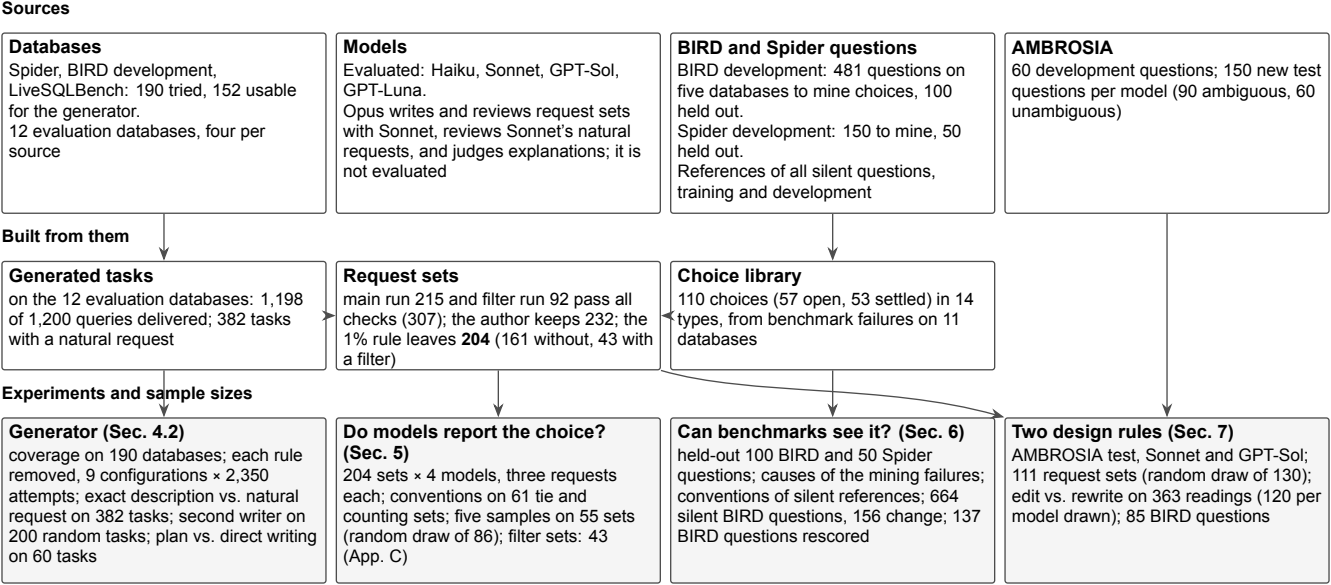
Sources
Databases
Spider, BIRD development, LiveSQLBench: 190 tried, 152 usable for the generator.
12 evaluation databases, four per source
Models
Evaluated: Haiku, Sonnet, GPT-Sol, GPT-Luna.
Opus writes and reviews request sets with Sonnet, reviews Sonnet's natural requests, and judges explanations; it is not evaluated
BIRD and Spider questions
BIRD development: 481 questions on five databases to mine choices, 100 held out.
Spider development: 150 to mine, 50 held out.
References of all silent questions, training and development
AMBROSIA
60 development questions; 150 new test questions per model (90 ambiguous, 60 unambiguous)
Built from them
Generated tasks
on the 12 evaluation databases: 1,198 of 1,200 queries delivered; 382 tasks with a natural request
Request sets
main run 215 and filter run 92 pass all checks (307); the author keeps 232; the 1% rule leaves 204 (161 without, 43 with a filter)
Choice library
110 choices (57 open, 53 settled) in 14 types, from benchmark failures on 11 databases
Experiments and sample sizes
Generator (Sec. 4.2)
coverage on 190 databases; each rule removed, 9 configurations × 2,350 attempts; exact description vs. natural request on 382 tasks; second writer on 200 random tasks; plan vs. direct writing on 60 tasks
Do models report the choice? (Sec. 5)
204 sets × 4 models, three requests each; conventions on 61 tie and counting sets; five samples on 55 sets (random draw of 86); filter sets: 43 (App. C)
Can benchmarks see it? (Sec. 6)
held-out 100 BIRD and 50 Spider questions; causes of the mining failures; conventions of silent references; 664 silent BIRD questions, 156 change; 137 BIRD questions rescored
Two design rules (Sec. 7)
AMBROSIA test, Sonnet and GPT-Sol; 111 request sets (random draw of 130); edit vs. rewrite on 363 readings (120 per model drawn); 85 BIRD questions