Text-to-SQL Models Make Choices Without Saying So [Experiment, Analysis & Benchmark]: Supplementary Material

Yichen Xu
University of California, Berkeley
Berkeley, CA, USA
yichen_xu@berkeley.edu

Shipei Lin
School of Data Science
The Chinese University of Hong Kong, Shenzhen
Shenzhen, China
225045017@link.cuhk.edu.cn

Chenhao Ma
School of Data Science
The Chinese University of Hong Kong, Shenzhen
Shenzhen, China
machenhao@cuhk.edu.cn

This supplement complements the paper; it does not repeat it. Section, figure, table and equation numbers without an S refer to the paper. The data and code behind every number are in the reproducibility package of the repository. The last two sections list every library choice and every accepted request set in full, so that each can be checked by hand.

A QUERY GENERATION DETAILS

Algorithm S1 lists the construction of a request set. Figure S1 shows, for each feature, the share of 152 databases whose schema allows it and the share where a generated query with it was accepted by both model reviewers; and Table S1 the queries delivered on the 12 evaluation databases.

Formal statement. Let $\mathcal{L}(K_D)$ be the queries that the catalogue K_D of database D and the type rules allow, $\delta(p)$ a query's arithmetic and step depth, and $O(p)$ its operations. For target depths d and required operations R the generator samples from

$$\mathcal{P}(D, d, R) = \{p \in \mathcal{L}(K_D) : \delta(p) = d,\ R \subseteq O(p),\ I_D(p) = 1\}, \quad (1)$$

where $I_D(p) = 1$ means that all data checks pass. It keeps a partial query z only if some completion can still reach the target,

$$\mathrm{can}(z, d, R) = \mathbf{1}\{\exists (d, U) \in \mathrm{Reach}(z) : R \subseteq U\}, \quad (2)$$

where $\mathrm{Reach}(z)$ lists the depths and operation sets still possible. The look-ahead is exact for the rules, and no query is drawn twice. For a request set F with specific requests q_A, q_B and open request $q_\circ$, let $H_{ij}^{(r)} = 1$ when reviewer r accepts request i for reading j, $f_r = 1$ when r finds no changed or missing requirement, and $c(F) = 1$ when all code checks pass. The set is kept when

$$\mathrm{keep}(F) = c(F) \wedge \bigwedge_r \left[f_r \wedge H^{(r)}_{\{A,B\},\{A,B\}} = I_2 \right] \wedge \bigvee_r \left[H^{(r)}_{\circ A} \wedge H^{(r)}_{\circ B} \right]. \quad (3)$$

B LIBRARY COMPOSITION

Table S2 gives the labels and counts in the fixed example library for every type. "Open" includes both unresolved and contested records. "Settled" denotes failure records the reviewers label as settled by the wording; this label is not the word-list check used for the defaults experiment. Failures that differ only in output format (51 with two accepted readings, 10 with one) are not choices of meaning and are not in the table. Possible reference errors (27) and unplaced cases (four) are also kept apart.

Algorithm S1 Build a request set around one calculation choice

Require: Databases D, target depths d, required operations R
Ensure: Accepted request sets and reasons for rejected sets

for each requested task **do**
  Sample $p_A \in \mathcal{P}(D, d, R)$ with the look-ahead of Eq. (2)
  If no query is delivered, record the shortfall and continue
  Describe p_A; write and check its request
  Propose edits $p_A \to p_B$ that express a different choice
  Execute and describe both queries; require different answers
  Repair failed proposals once; review both readings
end for
Select one choice per task at most, balancing types and databases
for each selected choice **do**
  List what stays fixed and what changes
  Write requests specifying A, specifying B, and leaving it open
  Check requirements and all six request–reading pairs
  Revise a failing set once and recheck
  Keep F if keep(F); otherwise record the reason
end for

Table S1: Program construction on 12 databases, with no model calls after catalogue preparation. Every delivered SQL is distinct, has the requested depth, and passes both integrity checks.

Source	DBs	Delivered	Constructions	Seconds
Spider	4	398/400	289	127
BIRD	4	400/400	270	274
LiveSQLBench	4	400/400	309	37
Total	12	1198/1200	868	437

Source-level times are rounded individually; the total is rounded from the unrounded sum. Separate task review delivers 384/400 tasks, and language generation supplies 382/384 requests.

Table S2 counts the 204 accepted request sets by type, 43 of them with a choice in a filter (Appendix C). Over the 215 sets whose choice a model proposes, aggregation level contributes 63, ties 44, missing values 40, counting unit 35, calculation 15, relationship role nine, and time scope two. The remaining seven belong to

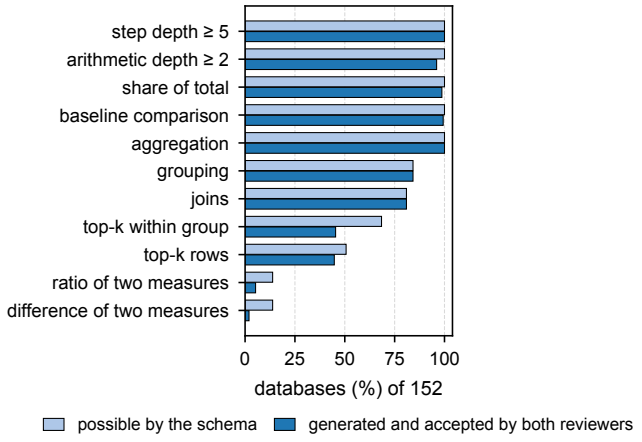


Figure S1: Features the generator can build, as a share of 152 databases: possible by the schema, and generated and accepted.

Table S2: All choice types (Figure 4 of the paper shows the major ones). Library: choices of meaning in benchmark failures, open (the wording fits both readings) or settled (it selects one). References: BIRD and Spider questions silent on the choice, by the reading their reference SQL takes: the usual one, or the other with its share (%), after the less common readings were reviewed; –: not readable from the SQL. Sets: the 204 accepted request sets, of the types whose other reading at least 1% of silent references take. The library's value-representation type is split by its SQL edit: comparison strictness is a strict cutoff, any other edit value matching.

		Library		References		
Choice type	The choice	Open	Settled	Usual	Other (%)	Sets
Aggregation level	Rows or group averages?	6	2	244	15 (5.8)	61
Calculation	Which formula is meant?	1	1	613	1 (0.2)	0
Quantifier	"Than any": one or all?	3	0	8	7 (46.7)	0
Condition linkage	Conditions on one row?	1	0	260	0 (0.0)	0
Counting unit	Count rows or entities?	7	27	2,521	237 (8.6)	30
Missing values	Keep or skip missing?	8	3	1,630	26 (1.6)	29
Negation	Some non-match, or none?	0	1	272	3 (1.1)	1
Population scope	Keep unmatched rows?	0	1	12,462	5 (0.0)	0
Referent column	Which column is meant?	19	6	–	–	0
Relationship role	Which link to follow?	2	4	15	0 (0.0)	0
Strict cutoff	Strict or inclusive?	1	0	309	5 (1.6)	42
Ties and selection	One tied row, or all?	7	3	2,826	397 (12.3)	41
Time scope	Period? Ends included?	1	2	254	1 (0.4)	0
Value matching	Exact or loose match?	1	3	9,330	22 (0.2)	0
Total		57	53			204

the six additional labels: running-total interpretation (two), active-status filtering, comparison baseline, baseline population, baseline definition, and row-count versus quantity measurement (one each). These counts reflect the generation and selection procedure. For Table 1, AmbiQT's column and table ambiguity and AMBROSIA's vagueness map to *which column*, AMBROSIA's scope to *any vs. all*, and attachment to *how conditions link*.

The References columns of Table S2 read each reference query with an SQL parser and the schema's keys and column types, without a model, and count a question only when its wording has no cue for the choice. Ties, counting unit, missing values, population scope, "between" and value matching follow Section 6.1 of the paper; calculation is a question that says percent, answered times 100 or as a fraction. *Quantifier*: a question that says "than any", answered by comparing with every value (`> MAX`, `< MIN`, `ALL`) or with at least one (`> MIN`, `< MAX`, `ANY`). *Negation*: a negated question answered with no matching record (`NOT IN`, `NOT EXISTS` or `EXCEPT` over another table) or with some non-matching record (<> on the many side of a join). *Aggregation level*: an entity ranked by a numeric column on the many side of a join, by one row's value or by a sum or average per group. *Condition linkage*: two or more conditions on different columns of the many side of a join, on one row or on separate rows (the table joined twice by the same key, or an `INTERSECT`). *Relationship role*: two tables that share a parent and are also connected by a link table, joined through the parent or through the link table. Which column a phrase means depends on the words, not the SQL, so the referent column has no count. Each rule reads one SQL form of its choice, so it detects only part of the choices mined from failures: applied to the 85 library records of the twelve types, the rules find the choice in 50, including all ties and quantifier records, 27 of 34 counting records and 4 of 6 relationship records, but none of the aggregation, time scope or value matching records, whose forms (grouping by a non-key attribute, an average used as a threshold) the rules do not read. The counts are therefore a lower bound on how often the choices arise. Every reference that takes the less common reading was then read against fixed criteria (the SQL takes that reading, the wording leaves the choice open, and the other reading could change the answer), and every verdict is listed for the author's review (`human_supervision/review-reference-minority.md`). Of 1,259 such references, 540 are dropped, most often because the wording settles the choice ("list the students with the highest", "how many different"), the join cannot repeat the counted entity, or a `LIKE` only escapes an apostrophe. Only this side was reviewed, so the shares in the table are conservative. The same drops apply to Section 6.1 of the paper.

C REQUEST SETS WITH FILTER CONDITIONS

The filter sets of Section 2.4 of the paper come from one optional rule of the generator. It adds a single condition to the first stage of a query, on the rows of that stage's table, and builds the second reading at the same time. The five forms are: a date range, inclusive (>= and <=) or exclusive (> and <), with ends at the 30th and 70th percentile of a date column; a cutoff at a frequent value of a measure, `>= v` or `> v`; "without" a category, no record of the entity has it (`NOT IN`) or some record lacks it (`IN` with <>); two categories, on one record or on any records of the entity; and a per-entity count, over entities with matching rows (inner join) or over all entities (outer join, zero counts kept). A condition is used only if it keeps 20–80% of the stage's rows, and the generator prefers forms it has used less often. The seed is 20261016; 161 of 180 requested queries are delivered, 121 of them have readings with different answers, and 92 request sets pass all checks.

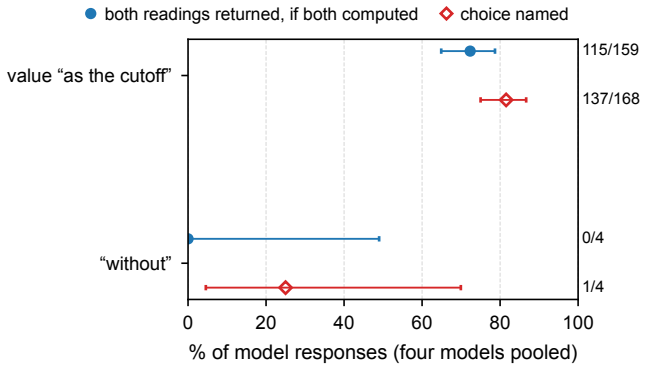


Figure S2: Filter request sets (43), four models pooled, by the choice in the filter and the wording of the open request (Wilson 95% intervals).

Table S3 gives the results per model on the 43 accepted filter sets, and Figure S2 by choice type. The author rejects all 17 date-range sets. Their open requests say “between,” which SQL’s `BETWEEN` includes, and so do 254 of the 255 BIRD and Spider references whose question says “between” with no other cue; the wording therefore settles the ends. The one exception is a Spider reference that applies strict bounds to the start and end columns of a staff role. A BIRD reference whose evidence says `BETWEEN 121 AND 130` but whose query excludes both ends is not counted, since its evidence settles the ends. Four further references use strict bounds moved one step out, or a negated range, and select the inclusive range. What models raise follows the wording of the open request. Forty of the 42 cutoff requests give a value “as the cutoff” or “threshold,” with no comparison word, and the other two say “upwards of”; models return both the strict and the inclusive reading in 115 of the 159 cases where they compute both.

The settled request is never called open. GPT-Luna returns both readings least often, as on the other sets. Of the 49 rejected sets, 17 are the date ranges, 15 are cutoffs whose wording (“passes”, “clears”, “on the high side of”) already makes the comparison strict, five are “without” requests that already mean no such record, and one is a per-entity average whose request restricts the population. The other 11 test per-entity counts with or without unmatched entities (eight) and conditions on one or any records (three), choices whose other reading real references almost never take (Figure 4 of the paper).

D CALCULATION CHOICES WITH REAL CONSEQUENCES

Table S4 surveys documented cases, across six domains, in which one open or mistaken choice of the kind our library records changed a consequential number or decision. None is a text-to-SQL failure. Each entry was checked against the original paper or official report; we left out cases whose primary source we could not reach, such as a bank’s value-at-risk spreadsheet and a classic study of graduate admissions.

Table S3: Filter request sets (43 accepted). Top, per model: both specific requests right (with exploration); among those, open requests answered with both readings; open requests that name the choice (Opus judge); specific requests right (single call); missed readings the model computes when asked, of all missed.

Figure S2 gives the results by choice type.

Model	Both right	Both returned	Named	Right	Missed, computable
Haiku	42	33	40	40	9 / 10
Sonnet	41	36	38	42	5 / 6
GPT-Sol	42	32	35	40	10 / 10
GPT-Luna	38	14	25	41	24 / 29

E FURTHER DETAILS

Seeds, versions and reruns. The grammar is `compose-v3.4`. Task requests, Spider splitting, request set selection, explanation allocation, and bootstrap resampling use seed 20261003; BIRD holdout selection uses 20261001; the paired wording comparison uses 20261004. Coverage uses seed 1 and depth-control runs seeds 1–3. Python was 3.11.5 and SQLite 3.41.2. During the experiments, the Claude command-line client changed from 2.1.287 to 2.1.289 and the Codex client from 0.147 to 0.160. Temporary usage-limit outputs were removed and retried rather than treated as SQL answers. A final query that failed under concurrent execution was rerun alone: ten such answers then matched their reference (nine GPT-Luna, one GPT-Sol). Token counts are unavailable for the explanation runs, so latency is reported instead of cost.

Rewrite comparison. The comparison of the original specific request (a) with the other texts was fixed before the run. The comparisons with the edited request (c) were added after an interim look at half of the readings.

Response selection. The explanation examples are the first saved responses displaying each tested pattern in file order. The worked aggregation example uses the smallest case-identifier hash among eligible request sets. Library examples use the first eligible record per type at high confidence, or medium confidence if none is high, with fixed benchmark and task ordering. Generated examples outside library labels enumerate retained cases before selecting the four shown in the main text. These selection rules support traceability; illustrative examples are not frequency estimates.

REFERENCES

[1] Rotem Botvinik-Nezer, Felix Holzmeister, Colin F. Camerer, et al. 2020. Variability in the Analysis of a Single Neuroimaging Dataset by Many Teams. *Nature* 582 (2020), 84–88.
[2] Ann Furedi. 1999. The Public Health Implications of the 1995 ‘Pill Scare’. *Human Reproduction Update* 5, 6 (1999), 621–626.
[3] Thomas Herndon, Michael Ash, and Robert Pollin. 2014. Does High Public Debt Consistently Stifle Economic Growth? A Critique of Reinhart and Rogoff. *Cambridge Journal of Economics* 38, 2 (2014), 257–279.
[4] Jacques-Louis Lions et al. 1996. *ARIANE 5 Flight 501 Failure: Report by the Inquiry Board.* Technical Report. European Space Agency and Centre National d’Études Spatiales, Paris.
[5] Mars Climate Orbiter Mishap Investigation Board. 1999. *Mars Climate Orbiter Mishap Investigation Board Phase I Report.* Technical Report. National Aeronautics and Space Administration.
[6] Ziad Obermeyer, Brian Powers, Christine Vogeli, and Sendhil Mullainathan. 2019. Dissecting Racial Bias in an Algorithm Used to Manage the Health of Populations. *Science* 366, 6464 (2019), 447–453.

Table S4: Documented cases where one calculation choice changed a consequential number or decision, by domain, with the closest choice type in our library.

Domain	Case	Choice and consequence	Library type
Economics and policy	Public debt and growth [3]	Growth was averaged within each country before averaging across countries, so one year of New Zealand data weighed as much as decades of another country. With corrected data and weighting, high-debt countries grew 2.2% a year, not −0.1%.	aggregation level
Health care	Risk algorithm for extra care [6]	"Need" was computed as health-care cost. The algorithm looked accurate by standard measures, yet fixing the proxy would raise the share of Black patients receiving extra help from 17.7% to 46.5%.	which column a phrase means
Health care	Hospital mortality rates [7]	Four methods, given the same 2.5 million discharges, included 28–95% of them. 12 of 28 hospitals with higher-than-expected mortality by one method had lower-than-expected mortality by another.	population scope
Public health	1995 UK contraceptive pill warning [2]	A warning of higher thrombosis risk for two pill types was followed by pill use among under-16s falling from 40% to 27%, and by estimated NHS costs of £21 million for maternity care and £46 million for abortions. The author concludes that risk should be more carefully presented.	how a quantity is stated (not in our library)
Social science	Referees and skin tone [8]	29 teams analysed the same data. Their odds ratios ranged from 0.89 to 2.93; 20 teams found a significant effect and 9 did not, using 21 combinations of covariates.	several analysis choices
Science	Brain imaging [1]	70 teams analysed one dataset to test the same 9 hypotheses. No two chose the same workflow, and the results of the hypothesis tests varied substantially.	several analysis choices
Science	Gene names in spreadsheets [9]	Spreadsheet software converts gene names to dates and numbers by default; about one fifth of papers with supplementary gene lists in spreadsheets contain such conversions.	value representation
Engineering	Mars Climate Orbiter [5]	Thruster data in English units were used where the trajectory software expected metric units; the spacecraft was lost.	value representation
Engineering	Ariane 5 flight 501 [4]	A 64-bit floating-point value was converted to a 16-bit signed integer that could not hold it; the launcher broke up about 40 seconds after lift-off.	value representation

[7] David M. Shahian, Robert E. Wolf, Lisa I. Iezzoni, Leslie Kirle, and Sharon-Lise T. Normand. 2010. Variability in the Measurement of Hospital-Wide Mortality Rates. *New England Journal of Medicine* 363, 26 (2010), 2530–2539.

[8] Raphael Silberzahn, Eric L. Uhlmann, Daniel P. Martin, et al. 2018. Many Analysts, One Data Set: Making Transparent How Variations in Analytic Choices Affect Results. *Advances in Methods and Practices in Psychological Science* 1, 3 (2018), 337–356.

[9] Mark Ziemann, Yotam Eren, and Assam El-Osta. 2016. Gene Name Errors Are Widespread in the Scientific Literature. *Genome Biology* 17 (2016), 177.

F CATALOGUE OF LIBRARY CHOICES

Table S5: Every choice of meaning found in a benchmark failure (110: 57 open, 53 settled by the wording), by type, as counted in Table S2. Question: the benchmark question with its evidence. The two readings are written as the reviewers recorded them; each also has an SQL query in the package (`library/examples.jsonl`, by record id).

Source	Question	The choice	Reading A	Reading B
Aggregation level, open				
BIRD `thrombosis_prediction` batch2:b049	For all patient born in 1982, state if their albumin is within normal range. Evidence: Year(Birthday) = '1982'; albumin is within normal range refers to ALB between 3.5 and 5.5	Whether the normal/abnormal status is given once per laboratory record or once per patient born in 1982, since each patient can have many Laboratory rows	Output one normal/abnormal label for each lab test of each patient born in 1982. A patient appears once per Laboratory record.	Output one label per patient born in 1982. The patient is 'normal' only if their recorded albumin values are all between 3.5 and 5.5; otherwise the label is 'abnormal'.
BIRD `california_schools` batch3:b007	Which cities have the top 5 lowest enrollment number for students in grades 1 through 12? Evidence: K-12 refers to students in grades 1 through 12.	Whether enrollment is summed per city and cities are ranked by the total, or individual schools are ranked by their own K-12 enrollment and the cities of the 5 lowest schools are returned	Add up K-12 enrollment across all schools in each city, then return the 5 cities with the smallest totals.	Rank individual schools by their own K-12 enrollment and return the cities of the 5 schools with the lowest enrollment.
BIRD `california_schools` batch4:b007	How many students from the ages of 5 to 17 are enrolled at the State Special School school in Fremont for the 2014-2015 academic year? Evidence: State Special School means EdOpsCode = 'SSS'	Whether to give one enrollment figure per matching State Special School in Fremont, or add them into a single total, given that the question says 'the ... school' (singular) but two schools match the filters.	List the ages 5-17 enrollment of each State Special School in Fremont for 2014-2015 as its own row (40 and 335).	'How many students' asks for one number: the combined ages 5-17 enrollment across all State Special Schools in Fremont for 2014-2015 (375).
BIRD `financial` batch4:b021	Who are the account holder identification numbers whose who have transactions on the credit card with the amount is less than the average, in 1998? Evidence: Operation = 'VYBER KARTOU' refers to credit card withdrawal	Which transactions the 'average' amount threshold is computed over	The average is taken over all 1998 transactions of any type. A card withdrawal qualifies if its amount is below that average.	The average is taken over 1998 credit card withdrawals only (operation = 'VYBER KARTOU'). A card withdrawal qualifies if its amount is below the average of the other 1998 card withdrawals.
BIRD `thrombosis_prediction` batch5:b049	What number of patients with a degree of thrombosis level 2 and ANA pattern of only S, have a level of anti-Cardiolip in antibody (IgM) 20% higher than average? Evidence: thrombosis level 2 refers to Thrombosis = 2; ANA pattern of only S refers to ANA = 'S'; average anti-Cardiolip in antibody (IgM) refers to AVG(`aCL IgM`); calculation = MULTIPLY(AVG + AVG, 0.2)	Which population the average aCL IgM is computed over, which sets the 1.2x threshold	Compute the average aCL IgM over only the examination records with Thrombosis = 2 and ANA Pattern 'S'. Then count the records in that same subset whose aCL IgM is more than 20% above that subgroup average.	Compute the average aCL IgM over all examination records. Then count the records with Thrombosis = 2 and ANA Pattern 'S' whose aCL IgM is more than 20% above that overall average.
BIRD `california_schools` bird_pairs:b034	Which active district has the highest average score in Reading?	Whether to rank districts by one school's AvgScrRead value or by the average of AvgScrRead across each district's schools	Among active schools, find the school record with the highest average Reading score and return that school's district	Among active schools, group the records by district, average their average Reading scores, and return the district with the highest average
Aggregation level, settled by the wording				
BIRD `thrombosis_prediction` batch6:b043	For all female patient with total protein (TP) beyond the normal index, what is the deviation of their TP idex from the normal. Evidence: female refers to SEX = 'F'; total protein (TP) beyond the normal index refers to TP > 8.5; deviation of TP index from normal refers to SUBTRACT(TP, 8.5)	Whether to report one deviation per lab record or one combined deviation per female patient, given that a patient can have several Laboratory rows with TP > 8.5	Return TP - 8.5 for every lab measurement where a female patient's TP is above 8.5. A patient with several high readings gets one row per reading.	'Their deviation' is a per-patient quantity: for each female patient with high TP, return the average of TP - 8.5 over that patient's high readings, so each patient gets one row.
BIRD `toxicology` batch7:b005	How many molecules without sulphur element is not having double bond? Evidence: double bond refers to bond_type = ' = '; bond_type ! = ' = '; sulphur refers to element = 's'	Where the 'without sulphur' and 'no double bond' conditions apply: to individual atom/bond rows in the join, or to the whole molecule as an exclusion	Count molecules that have at least one atom that is not sulphur and at least one bond that is not a double bond. Each condition is checked on single rows of the atom–bond join.	Count molecules that contain no sulphur atom at all and no double bond at all. Each condition excludes the whole molecule.
Calculation, open				
BIRD `financial` batch2:b023r	Provide the IDs and age of the client with high level credit card, which is eligible for loans. Evidence: the credit card is high-level refers to card.type = 'gold'; eligible for loans refers to disp.type = 'OWNER'	How to compute the client's age from birth_date: subtract the birth year from the current year, or count completed years of life.	Age is the current calendar year minus the client's birth year, ignoring whether the birthday has happened yet this year	Age is the number of full years the client has lived as of today. It is the current year minus the birth year, minus one more if the client's birthday has not come yet this year.
Calculation, settled by the wording				
BIRD `thrombosis_prediction` batch2:b053	For all patients who are older than 55 years old, what is the percentage of female who has abnormal prothrombin time (PT)? Evidence: older than 55 years old = SUBTRACT(year(current_timestamp), year(Birthday)) > 55; abnormal prothrombin time (PT) refers to PT > = 14; percentage = DIVIDE(SUM(PT > = 14 AND SEX = 'F'), SUM(PT > = 14)) * 100; female refers to sex = 'F';	Which population the percentage is taken over (the denominator)	Among all lab records of patients older than 55, the share that belong to a female patient and show an abnormal PT (PT >= 14). The denominator is every joined lab record for these patients.	Among the lab records of patients older than 55 that show an abnormal PT (PT >= 14), the share that belong to female patients. This matches the evidence formula SUM(PT>=14 AND SEX='F') / SUM(PT>=14).
Condition linkage, open				
BIRD `toxicology` batch7:b018	Calculate the percentage of carcinogenic molecules with triple bonded Hidrogen atoms. Evidence: hydrogen refers to element = 'h'; label = '+' mean molecules are carcinogenic; triple bond refers to bond_type = '#'; percentage = DIVIDE(SUM(label = '+'), COUNT(molecule_id)) * 100.0 where element = 'h' AND bond_type = '#';	How a hydrogen atom is tied to a triple bond: either the hydrogen is one of the two atoms in the '#' bond (linked through `connected`), or the hydrogen and some '#' bond just appear in the same molecule (linked only through molecule_id).	Count molecules that have at least one hydrogen atom and also at least one triple bond anywhere in the molecule. The hydrogen does not have to be part of that bond. The result is the share of these molecules that are carcinogenic.	Count molecules where a hydrogen atom is itself one of the two atoms in a triple bond. The result is the share of these molecules that are carcinogenic.
Counting unit, open				

Source	Question	The choice	Reading A	Reading B
BIRD `thrombosis_prediction` batch2:b052	What is the average age of the male patient with high cholesterol? Evidence: average age = DIVIDE(SUM(SUBTRACT(YEAR(NOW()), YEAR(birthday))), COUNT(ID)); male patient refers to sex = 'M'; high cholesterol refers to `T-CHO` > = 250;	Whether the average is taken once per patient or once per matching lab record, since a patient can have several Laboratory rows with T-CHO >= 250	Average the age over every joined lab record with T-CHO >= 250 for male patients, so a patient with several high-cholesterol tests is counted once for each of those tests.	Average the age over the distinct male patients who have at least one lab record with T-CHO >= 250, so each patient is counted once.
BIRD `toxicology` batch3:b033	What is the percentage of carcinogenic molecules in triple type bonds? Evidence: label = '+' mean molecules are carcinogenic; triple bond refers to bond_type = '#'; percentage = DIVIDE(SUM(bond_type = '#') * 100, COUNT(bond_id)) as percent where label = '+'	Whether the percentage counts molecules that contain a triple bond or counts the triple bonds themselves.	Of the distinct molecules that have at least one triple bond, the percentage that are carcinogenic (label = '+').	Of all triple bonds, the percentage that belong to carcinogenic molecules (label = '+'). This counts bond_id, as the evidence's COUNT(bond_id) suggests.
BIRD `thrombosis_prediction` batch4:b048	For the year that concluded on December 31, 1998, how many male patients on average were tested in the lab each month? Evidence: the year that concluded on December 31, 1998 refers to Date BETWEEN '1998-01-01' AND '1998-12-31'; male refers to SEX = 'M'; calculation = DIVIDE(COUNT(ID), 12)	What unit to count: every male lab-test record in 1998, or the distinct male patients tested in each month	Count all lab-test records for male patients dated in 1998 and divide by 12, so a patient tested several times is counted once per test	For each month of 1998, count the distinct male patients tested that month, add up the monthly counts and divide by 12, so a patient is counted at most once per month
BIRD `thrombosis_prediction` batch4:b050	What is the ratio of male to female patients among all those with abnormal uric acid counts? Evidence: male refers to SEX = 'M'; female refers to SEX = 'F'; abnormal uric acid refers to UA < = '8.0' where SEX = 'M', UA < = '6.5' where SEX = 'F'; calculation = DIVIDE(SUM(UA <= '8.0' and SEX = 'M'), SUM(UA <= '6.5 and SEX = 'F'))	Counting unit for the male and female counts: lab records (each test row that meets the abnormal-UA condition) or distinct patients	Count lab records. Every Laboratory row where a male's UA <= 8.0 or a female's UA <= 6.5 adds 1, so a patient with several abnormal tests is counted several times. The ratio is the male record count divided by the female record count.	Count distinct patients. Each male patient with at least one UA <= 8.0 record counts once, and so does each female patient with at least one UA <= 6.5 record. The ratio is the number of such male patients divided by the number of such female patients.
BIRD `toxicology` batch5:b019	What is the least common element of all carcinogenic molecules? Evidence: label = '+' mean molecules are carcinogenic	What unit 'least common' is counted in: the number of carcinogenic molecules that contain the element, or the number of atoms of that element across all carcinogenic molecules	The least common element is the one that appears in the fewest carcinogenic molecules (count of distinct molecules containing it)	The least common element is the one with the fewest atoms in total across all carcinogenic molecules
BIRD `thrombosis_prediction` bird_pairs:b031	What is the percentage of female patient had total protein not within the normal range? Evidence: female refers to sex = 'F'; total protein not within the normal range refers to TP < '6.0' or TP > '8.5'; calculation = DIVIDE((ID where sex = 'F' and TP < '6.0' or TP > '8.5'), COUNT(ID)) * 100	What to count: lab test records or distinct female patients (a patient with several out-of-range TP results counts once or many times)	Among all lab records of female patients, the percentage of records whose total protein is below 6.0 or above 8.5	Among female patients who have lab records, the percentage of distinct patients with at least one total protein result below 6.0 or above 8.5
Spider `student_transcripts_tracking` spider_dev:spider_dev_0520	How many different degrees are offered?	What counts as one degree: a distinct degree name (degree_summary_name) or each degree program row in Degree_Programs	Count the distinct degree names (e.g., Bachelor, Master, PhD) across all degree programs	Count every degree program offered, treating each row in Degree_Programs as a separate degree
Counting unit, settled by the wording				
BIRD `thrombosis_prediction` batch1:b052	How many male patients are are with creatinine index out of the normal range? Evidence: creatinine (CRE) out of the normal range refers to CRE > = 1.5; Male refers to Sex = 'M'	Whether to count distinct male patients or the joined lab records (one row per out-of-range CRE test) for male patients	Count every lab test record with CRE >= 1.5 that belongs to a male patient. A patient with several abnormal tests is counted once per test.	Count distinct male patients who have at least one lab test with CRE >= 1.5. Each patient is counted once.
BIRD `thrombosis_prediction` batch1:b058	How many patients diagnosed with SLE have a normal white blood cell level? Evidence: diagnosed with SLE refers to Diagnosis = 'SLE'; normal white blood cell level refers to WBC between 3.5 and 9.0;	Counting unit: whether to count every matching lab record from the Patient–Laboratory join, or each distinct SLE patient who has at least one normal WBC reading.	Count every lab test record in which an SLE patient's WBC is between 3.5 and 9.0. A patient with several normal readings is counted once per reading.	Count each SLE patient once if they have at least one lab test with WBC between 3.5 and 9.0.
BIRD `thrombosis_prediction` batch2:b054	Among the patients with a normal Ig G level, how many of them have symptoms? Evidence: normal Ig G level refers to IGG > 900 and IGG < 2000; have symptoms refers to Symptoms IS NOT NULL;	Whether to count distinct patients or rows of the Patient×Laboratory×Examination join, where a patient with several qualifying lab tests or exams appears many times.	Count every combination of a lab record with normal IgG and an examination record with symptoms for the same patient. A patient is counted once per such lab/exam pair.	Count each patient once if they have at least one lab record with normal IgG and at least one examination record with symptoms.
BIRD `thrombosis_prediction` batch2:b056	For the patient with an abnormal alkaliphophatase level, how many of them are admitted to the hospital? Evidence: abnormal alkaliphophatase level refers to ALP > = 300; admitted to the hospital refers to Admission = '+';	Whether to count patient-lab join rows (one per lab record with ALP >= 300) or distinct patients	Count every lab record with ALP >= 300 that belongs to an admitted patient, so a patient with several abnormal tests is counted several times	Count the distinct admitted patients who have at least one lab record with ALP >= 300
BIRD `thrombosis_prediction` batch2:b057	Among the patients followed at the outpatient clinic, how many of them have a normal level of alkaliphophatase? Evidence: followed at the outpatient clinic refers to Admission = '-'; normal level of alkaliphophatase refers to ALP < 300;	Counting unit: whether to count every joined lab record with ALP < 300, or each outpatient patient only once.	Count every laboratory record with ALP below 300 that belongs to an outpatient (Admission = '-'). A patient with several such tests is counted once per test.	Count distinct outpatient patients who have at least one laboratory record with ALP below 300. Each patient is counted once.
BIRD `thrombosis_prediction` batch2:b059	For the patients whose total cholesterol is higher than normal, how many of them have a negative measure of degree of coagulation? Evidence: total cholesterol is higher than normal refers to `T-CHO` > = 250; negative measure of degree of coagulation refers to KCT = '-' ;	Counting unit: the gold counts rows from the three-way join (one per matching Laboratory × Examination record pair), not distinct patients.	Count every combination of a lab record with total cholesterol >= 250 and an examination record with KCT = '-' for the same patient, so a patient with several qualifying records is counted more than once.	Count distinct patients who have at least one lab record with total cholesterol >= 250 and at least one examination record with KCT = '-'.
BIRD `thrombosis_prediction` batch5:b055	How many patients who were female got white blood cells that were below 3.5? Evidence: female refers to SEX = 'F'; white blood cells that were below 3.5 refers to WBC < 3.5	Counting unit: whether to count each lab record (joined row) where a female patient's WBC was below 3.5, or each distinct female patient with at least one such record	Count the lab test records (patient-visit rows) of female patients in which WBC was below 3.5, so a patient with several low readings is counted once per reading	Count distinct female patients who had at least one lab record with WBC below 3.5
BIRD `toxicology` batch6:b014	Among all chemical compounds identified in the database, what percent of compounds form a triple-bond. Evidence: triple bond refers to bond_type = '#';	What the percentage counts: individual bonds or compounds (molecules) that contain a triple bond.	Of all bonds in the database, the percentage that are triple bonds (bond_type = '#').	Of all compounds (molecules), the percentage that contain at least one triple bond.

Source	Question	The choice	Reading A	Reading B
BIRD thrombosis_prediction batch6:b057	For patients with abnormal platelet level, state the number of patients with lower than normal range. How is it compare to the number of patients with higher than normal range? Evidence: abnormal platelet level refers to PLT <= 100 or PLT >= 400; platelet level lower than normal range refers to PLT < 100; calculation = SUBTRACT(SUM(PLT < 100), SUM(PLT > 400)); platelet level higher than normal range refers to PLT > 400;	Whether the counting unit is lab records (each PLT test row) or distinct patients	Count lab test records: the number of Laboratory rows with PLT <= 100 minus the number of rows with PLT >= 400, so a patient with several abnormal tests is counted once per test	Count distinct patients: the number of distinct patients with at least one PLT reading below the normal range minus the number of distinct patients with at least one reading above it
BIRD toxicology batch7:b001	Calculate the percentage of molecules containing carcinogenic compounds that element is hydrogen. Evidence: hydrogen refers to element = 'h'; label = '+' mean molecules are carcinogenic; percentage = DIVIDE(SUM(label = '+' and element = 'h'), COUNT(molecule_id)) * 100.0	What gets counted and what the base is: atom rows over all joined atoms, or distinct molecules containing hydrogen	Count atom rows. Take the atoms that are hydrogen and belong to a carcinogenic molecule, and divide by all atom rows in the atom–molecule join. This follows the evidence formula literally.	Count distinct molecules. Of the molecules that contain at least one hydrogen atom, find the percentage that are carcinogenic.
BIRD toxicology batch7:b004	How many molecules have a double bond type and among these molecule, how many are labeled as carcinogenic compound? Evidence: double bond refers to bond_type = ' = '; label = '+' mean molecules are carcinogenic;	Unit counted for the carcinogenic number: double-bond rows belonging to carcinogenic molecules, or distinct carcinogenic molecules that have a double bond	Count the molecules with a double bond, then count the double bonds whose molecule is labeled '+'. This gives one tally per bond row, not per molecule.	Count the molecules with a double bond, then count how many of those distinct molecules are labeled '+' (carcinogenic).
BIRD thrombosis_prediction batch7:b038	How many patients with a normal Ig A level came to the hospital after 1990/1/1? Evidence: normal Ig A level refers to IGA > 80 AND IGA < 500; came to the hospital after 1990/1/1 refers to YEAR(`First Date`) > = 1990;	Whether to count distinct patients or every joined laboratory record. A patient with several normal IgA lab results can be counted more than once.	Count every laboratory record that has an IgA value between 80 and 500 and belongs to a patient whose first hospital visit was in the required year range. A patient is counted once per qualifying lab test.	Count each patient once if they have at least one lab record with a normal IgA value and their first hospital visit was in the required year range.
BIRD thrombosis_prediction batch7:b040	How many patients with a abnormal C-reactive protein don't have their data recorded? Evidence: abnormal C-reactive protein refers to CRP ='+'; don't have data recorded refers to Description IS NULL;	Whether to count distinct patients or the joined Patient–Laboratory rows (lab records with CRP = '+') for patients whose Description is NULL	Count every lab record with an abnormal CRP (CRP = '+') that belongs to a patient with no recorded Description. A patient with several such lab tests is counted once per test.	Count each distinct patient once if they have at least one lab record with CRP = '+' and their Description is NULL.
BIRD thrombosis_prediction batch7:b048	Among the patients with normal anti-SM, how many of them does not have thrombosis? Evidence: normal anti-SM refers to SM IN('-', '+-'); SM = 'negative' means '-'; SM = '0' means '+-'; SM = '1' means '+'; does not have thrombosis refers to Thrombosis = 0;	Whether to count joined Examination×Laboratory rows or distinct patients. (The agent's other difference, using the literal values '-' and '+-' instead of 'negative' and '0', is a plain mistake: the evidence states that mapping explicitly. That is why its runs returned 0.)	Count every joined row: each pair of a lab record with normal anti-SM and an examination record with Thrombosis = 0 counts once, so a patient with several such records is counted several times.	Count each patient once: the number of distinct patients who have a lab record with normal anti-SM ('negative' or '0') and an examination record with Thrombosis = 0.
BIRD thrombosis_prediction batch7:b049	For the patients with an abnormal anti-SM, please list the IDs of the three youngest ones. Evidence: abnormal anti-SM refers to SM NOT IN ('negative', '0'); youngest refers to MAX(Birthday);	Whether the three slots in LIMIT 3 are filled by abnormal lab records, where one patient can appear more than once, or by distinct patients.	Take every lab record with an abnormal anti-SM value. Sort these records by the patient's birthday, newest first, and return the patient ID from the top three records. A young patient with several abnormal lab records can fill more than one of the three slots.	Find the distinct patients who have at least one abnormal anti-SM lab record. Return the IDs of the three youngest, each patient listed once.
BIRD thrombosis_prediction batch7:b059	Of the patients with an normal level of IGG, how many of them admitted to the hospital? Evidence: normal level of IGG refers to 900 < IGG < 2000; admitted to the hospital refers to Admission = '+';	Whether the count is taken over distinct patients or over the joined patient–laboratory rows (lab records with normal IgG)	Count every laboratory record that has a normal IgG value (900 < IGG < 2000) and belongs to an admitted patient. A patient with several qualifying lab tests is counted once per test.	Count the distinct admitted patients who have at least one laboratory record with a normal IgG value. Each patient is counted once.
BIRD thrombosis_prediction batch8:b000	What is the percentage of patient who has a abnormal level of glutamic oxaloacetic transaminase level, yet he or she is diagnosed with SLE? Evidence: abnormal level of glutamic oxaloacetic transaminase refers to GOT > = 60; percentage = MULTIPLY(DIVIDE(COUNT(ID WHERE GOT > = 60 AND Diagnosis = 'SLE'), COUNT(ID WHERE GOT > = 60)), 1.0);	What to count in the ratio: each abnormal-GOT lab record (the gold's join rows) or each distinct patient. The gold's `ELSE 0` makes its numerator count every row, so it always returns 1. A real share is only possible if that branch becomes NULL.	Count lab records, not patients: every Laboratory row with GOT >= 60 joined to its patient is counted once. In practice the numerator also counts all of those rows (ELSE 0 is never NULL), and integer division gives 1.	Take the distinct patients who have at least one lab result with GOT >= 60. Report the percentage of them whose Patient.Diagnosis is SLE.
BIRD thrombosis_prediction batch8:b001	How many male patients have their glutamic oxaloacetic transaminase in the normal range? Evidence: male refers to Sex = 'M'; glutamic oxaloacetic transaminase in the normal range refers to GOT < 60;	Whether to count lab records or distinct patients among male patients with GOT < 60	Count every laboratory record from a male patient where GOT is below 60, so a patient with several normal GOT tests is counted once per test.	Count distinct male patients who have at least one laboratory record with GOT below 60.
BIRD thrombosis_prediction batch8:b002	Please list the top three patients' birthdays with the highest glutamic pylvic transaminase in the normal range. Evidence: highest glutamic pylvic transaminase in the normal range refers to MAX(GPT < 60);	Whether 'top three patients' counts the top three laboratory records (one patient can show up more than once) or the top three distinct patients (each patient counted once, by their highest GPT below 60)	Take every lab test with GPT below 60, sort the tests by GPT from highest to lowest, and return the patient's birthday for each of the first three tests. The same patient or birthday can appear more than once if they had several high tests.	Count each patient once. Use their highest GPT test below 60, sort patients by that value, and return the birthdays of the top three patients. A patient with several tests at the same high value, such as the repeated '1943-06-02', appears only once.
BIRD thrombosis_prediction batch8:b003	For the patients with the normal glutamic pylvic transaminase level, how many of them are male? Evidence: normal glutamic pylvic transaminase level refers to GOT < 60; male refers to Sex = 'M';	Whether to count lab records (rows of joined Laboratory entries) or distinct patients who have a normal-level lab record	Count every lab record with GOT < 60 that belongs to a male patient, so a patient with several qualifying tests is counted once per test	Count distinct male patients who have at least one lab record with GOT < 60
BIRD thrombosis_prediction batch8:b007	How many male patients have a normal level of both albumin and total protein? Evidence: male refers to Sex = 'M'; normal level of both albumin and total protein refers to ALB > 3.5 and ALB < 5.5 AND TP between 6.0 and 8.5;	Whether to count each patient once (distinct patient IDs) or count every qualifying lab record row from the Patient–Laboratory join	Count every laboratory record of a male patient where albumin and total protein are both in the normal range, so a patient with several qualifying tests is counted several times	Count the distinct male patients who have at least one laboratory record with albumin and total protein both in the normal range
BIRD thrombosis_prediction batch8:b009	Among the patients whose total cholesterol is within the normal range, how many of them have a P pattern observed in the sheet of ANA examination? Evidence: total cholesterol is within the normal range refers to `T-CHO` < 250; P pattern observed in the sheet of ANA examination refers to ANA Pattern = 'P';	Counting unit: whether to count distinct patients or every row produced by joining Patient, Laboratory and Examination (one row per qualifying lab record paired with each qualifying exam record).	Count every pairing of a lab record with T-CHO < 250 and an examination record with ANA Pattern = 'P' for the same patient. A patient with several such records is counted several times.	Count each distinct patient once if they have at least one lab record with T-CHO < 250 and at least one examination with ANA Pattern = 'P'.

Source	Question	The choice	Reading A	Reading B
BIRD `thrombosis_prediction` batch8:b010	Among the patients with the normal level of triglyceride, how many of them have other symptoms observed? Evidence: normal level of triglyceride refers to TG < 200; have other symptoms refers to Symptoms is not null;	Whether to count distinct patients, or the rows produced by joining each patient's Examination records to their Laboratory records	Count every joined Examination×Laboratory row where the lab record has TG < 200 and the examination has non-null Symptoms. A patient is counted once for each matching pair of lab record and examination record.	Count each patient once if they have at least one lab record with TG < 200 and at least one examination with non-null Symptoms.
BIRD `thrombosis_prediction` batch8:b013	For the patients with a normal range of creatinine phosphokinase, how many of them have a positive measure of degree of coagulation? Evidence: normal range of creatinine phosphokinase refers to CPK < 250; positive measure of degree of coagulation refers to KCT = '+' or RVVT = '+' or LAC = '+' ;	Whether to count distinct patients or every joined Laboratory×Examination row that matches the filters	Count every matching row in the Patient–Laboratory–Examination join, so a patient with several qualifying lab tests or examinations is counted once per combination	Count distinct patients who have a CPK < 250 lab result and a positive KCT, RVVT or LAC examination
BIRD `thrombosis_prediction` batch8:b014	How many patients accepted to the hospital have a normal level of white blood cells? Evidence: accepted to the hospital refers to Admission = '+'; normal level of white blood cells refers to WBC between 3.5 and 9.0;	Whether to count distinct patients or the joined patient-lab-test rows (lab records) that meet the WBC condition	Count every laboratory record of an admitted patient whose WBC is between 3.5 and 9.0 inclusive, so a patient with several qualifying tests is counted once per test.	Count distinct admitted patients who have at least one laboratory record with WBC between 3.5 and 9.0.
BIRD `financial` bird_pairs:b023	After making a credit card withdrawal, how many account/s with monthly issuance has a negative balance? Evidence: Negative balance means balance < 0; Operation = 'VYBER KARTOU' stands for credit card withdraw. Frequency = 'POPLATEK MESICNE' stands for monthly issurance	Whether to count distinct accounts or count qualifying transaction rows (one per negative-balance credit card withdrawal) for monthly-issuance accounts.	Count each credit card withdrawal transaction on a monthly-issuance account that left a negative balance, so an account with several such withdrawals is counted more than once.	Count the distinct monthly-issuance accounts that had at least one credit card withdrawal leaving a negative balance.
BIRD `thrombosis_prediction` bird_pairs:b036	Among the patients with a normal blood glucose, how many of them don't have thrombosis? Evidence: normal blood glucose refers to GLU < 180; don't have thrombosis refers to Thrombosis = 0;	Counting unit: whether to count every row from the Patient–Laboratory–Examination join, or each distinct patient once	Count every joined combination of a lab record with GLU < 180 and an examination record with Thrombosis = 0. A patient with several such lab tests or examinations is counted once per combination.	Count distinct patients who have at least one lab record with GLU < 180 and at least one examination record with Thrombosis = 0. Each patient counts once.
Missing values, open				
BIRD `california_schools` batch2:b002	What is the average score in writing for the schools that were opened after 1991 or closed before 2000? List the school names along with the score. Also, list the communication number of the schools if there is any. Evidence: Communication number refers to phone number.	Whether to keep schools that match the open/close date condition but have no writing score (no SAT record, or AvgScrWrite is NULL), or list only schools that actually have a writing score.	List every school that opened after 1991 or closed before 2000, with its writing score and phone number. Keep a school even when it has no writing score and show the score as empty.	List only the schools that opened after 1991 or closed before 2000 and actually have a writing score. The question asks for the score, so schools without one are left out.
BIRD `california_schools` batch3:b001	For the school with the highest average score in Reading in the SAT test, what is its FRPM count for students aged 5-17?	How to handle the top SAT reading school when it has no matching FRPM record: skip it and take the highest-scoring school that has an FRPM record, or keep it and report its FRPM count as missing.	Look only at schools that have both SAT scores and an FRPM record. Among those, take the one with the highest average reading score and return its FRPM count for ages 5-17.	Take the school with the highest average SAT reading score of all schools, whether or not it has an FRPM record, and return its FRPM count for ages 5-17. That count is NULL if the school has no FRPM record.
BIRD `california_schools` batch3:b004	What is the phone number of the school that has the highest average score in Math?	How to handle the top AvgScrMath row in satscores when it has no matching record in schools: drop it and take the best-scoring school that does have a record, or keep it as the answer even though its phone is missing.	Only consider schools that have a record in the schools table, and return the phone number of the one with the highest average Math score among them.	Take the school with the highest average Math score overall in satscores, and return its phone number. Since that school has no schools record, the phone number is missing (NULL).
BIRD `superhero` batch5:b030	Which superhero has the same eyes, hair and skin colour? Indicate the publisher of the superhero. Evidence: which superhero refers to superhero_name; the same eyes, hair and skin colour refers to hair_colour_id = skin_colour_id AND hair_colour_id = eye_colour_id; publisher refers to publisher_name;	Whether superheroes with no matching publisher row should be kept (inner join vs. left join to publisher)	List only superheroes whose eye, hair and skin colour ids match and who have a publisher; heroes with no matching publisher row are dropped.	List every superhero whose eye, hair and skin colour ids match, and show the publisher if there is one; heroes without a matching publisher still appear with a NULL publisher.
BIRD `superhero` batch6:b023	How many superheroes didn't have any publisher? Evidence: didn't have any publisher refers to publisher.id = 1;	Whether superheroes with a NULL publisher_id count as having no publisher, along with those linked to publisher.id = 1 (the placeholder publisher record).	Count only superheroes whose publisher_id points to the publisher record with id 1. Superheroes with a missing (NULL) publisher_id are left out.	Count superheroes with no publisher in either form: those linked to the publisher record with id 1, and those whose publisher_id is missing (NULL).
Spider `cre_Doc_Template_Mgt` spider_dev:spider_dev_0309	Show all template ids and number of documents using each template.	Whether templates that no document uses should appear in the output with a count of 0, or only templates that at least one document uses should be listed	List only the template ids that appear in Documents (templates actually in use), each with its count of documents	List every template id in the Templates table with its document count, so templates with no documents show a count of 0
Spider `world_1` spider_dev:spider_dev_0799	What are the different continents and the total popuation and average life expectancy corresponding to each, for continents that have an average life expectancy less than 72?	Whether countries with a missing (NULL) life expectancy still count toward a continent's total population	Total population sums every country on the continent, including countries with no recorded life expectancy. The average life expectancy covers only the countries that have a value.	Use only countries that have a recorded life expectancy, so countries with no life expectancy are left out of both the total population and the average.
Spider `network_1` spider_dev:spider_dev_0882	Show the student IDs and numbers of friends corresponding to each.	Whether students who have no rows in Friend should be listed with a count of 0, or left out.	List only the students who appear in the Friend table, each with how many friend rows they have, so every student shown has at least one friend.	List every student in Highschooler with their number of friends, and show 0 for students who have no friends.
Missing values, settled by the wording				
BIRD `california_schools` batch3:b009	List the names of virtual schools that are among the top 5 in their respective counties based on average reading scores. Evidence: Exclusively virtual refers to Virtual = 'F'; respective counties means PARTITION BY County	Whether exclusively virtual schools with no average reading score (NULL AvgScrRead) can count as top-5 in their county	Rank every exclusively virtual school in each county by average reading score, including schools with no reading score. Those schools sort last but are still listed when their county has five or fewer such schools.	Only exclusively virtual schools that have an average reading score are ranked. Schools with no reading score cannot be among the top 5.

Source	Question	The choice	Reading A	Reading B
BIRD `superhero` batch7:b028	Among all superheroes in Marvel Comics, identify the percentage of 'good' superheroes. Evidence: Marvel Comics refers to publisher_name = 'Marvel Comics'; percentage = MULTIPLY(DIVIDE(SUM(alignment = 'Good' WHERE publisher_name = 'Marvel Comics'), COUNT(publisher_name = 'Marvel Comics')), 100.0); good superheroes refers to alignment = 'Good';	Whether Marvel superheroes with no matching alignment row (NULL or unmatched alignment_id) count in the denominator of the percentage	Count only Marvel superheroes that have a recorded alignment. The percentage is the share of 'Good' heroes among those with a known alignment, and heroes with no alignment are dropped.	Count all Marvel superheroes, including those with no recorded alignment. The percentage is the share of 'Good' heroes among every Marvel hero, and heroes with no alignment count as not good.
BIRD `california_schools` bird_pairs:b042	What is the complete address of the school with the lowest excellence rate? Indicate the Street, City, Zip and State. Evidence: Execellence Rate = NumGE1500 / NumTstTakr; complete address has Street, City, State, Zip code	Whether schools with an undefined excellence rate count as candidates for the lowest rate. The rate is undefined when NumGE1500 is NULL or NumTstTakr is 0. Under SQLite's NULLs-first ascending sort, such a school ends up ranked lowest.	Sort every joined school by NumGE1500 / NumTstTakr in ascending order and take the first row. Schools whose rate is missing or undefined are kept, and because SQLite sorts NULLs first, one of them is returned as having the 'lowest' rate (the Fresno school at 1111 Van Ness Avenue).	Only consider schools where the excellence rate can actually be computed, meaning NumGE1500 is not NULL and there is at least one test taker. Return the address of the school with the smallest real rate (222 North Virgil Avenue, Los Angeles).
Negation, settled by the wording				
BIRD `toxicology` batch3:b032	What percentage of carcinogenic-type molecules does not contain fluorine? Evidence: label = '+' mean molecules are carcinogenic; contain fluorine refers to element = 'f'; percentage = DIVIDE(SUM(element = 'f') * 100, COUNT(molecule_id)) where label = '+'; Should consider the distinct atoms when counting;	What it means for a carcinogenic molecule to 'not contain fluorine': the molecule has at least one atom that is not fluorine, or the molecule has no fluorine atom at all.	A carcinogenic molecule counts if at least one of its atoms is not fluorine. The result is the share of '+' molecules that have some non-fluorine atom.	A carcinogenic molecule counts only if none of its atoms is fluorine. The result is the share of '+' molecules with no fluorine atom: all '+' molecules minus those with any fluorine atom, divided by all '+' molecules.
Population scope, settled by the wording				
BIRD `thrombosis_prediction` batch7:b037	How many patients with an Ig G higher than normal? Evidence: Ig G higher than normal refers to IGG >= 2000; Should consider DISTINCT in the final result;	Whether the patient count only includes patients who also have records in the Examination table (because of the inner join to Examination), or every patient with any lab IgG result of at least 2000	Count distinct patients who have a lab IgG of at least 2000 and who also have at least one examination record	Count distinct patients who have any lab IgG of at least 2000, whether or not they have an examination record
Quantifier, open				
Spider `world_1` spider_dev:spider_dev_0772	What are the countries that have greater surface area than any country in Europe?	Whether 'greater than any country in Europe' means larger than at least one European country (compare to the smallest European area) or larger than every European country (compare to the largest European area)	Countries whose surface area is greater than at least one European country, meaning greater than the smallest European surface area	Countries whose surface area is greater than every European country, meaning greater than the largest European surface area
Spider `world_1` spider_dev:spider_dev_0774	What are the African countries that have a population less than any country in Asia?	Whether 'less than any country in Asia' means less than at least one Asian country (compare to the Asian maximum) or less than every Asian country (compare to the Asian minimum)	African countries whose population is smaller than at least one Asian country's population, i.e., smaller than the largest Asian population	African countries whose population is smaller than every Asian country's population, i.e., smaller than the smallest Asian population
Spider `world_1` spider_dev:spider_dev_0777	What are the Asian countries which have a population larger than that of any country in Africa?	Whether 'larger than that of any country in Africa' means larger than at least one African country (exceeds the smallest African population) or larger than every African country (exceeds the largest African population)	Asian countries whose population is larger than the population of at least one African country, meaning it exceeds the smallest African country's population	Asian countries whose population is larger than the population of every African country, meaning it exceeds the largest African country's population
Referent column, open				
BIRD `financial` batch2:b015	How many clients opened their accounts in Jesenik branch were women? Evidence: A2 has region names; Woman and female share the same meaning; female refers to gender = 'F'	Which district decides that a client 'opened their accounts in Jesenik branch': the client's own district_id, or the district_id of the account(s) the client is linked to through disp.	Count female clients whose own home or registered district (client.district_id) is Jesenik.	Count distinct female clients linked through disp to at least one account whose branch (account.district_id) is Jesenik.
BIRD `financial` batch2:b020	List the top nine districts, by descending order, from the highest to the lowest, the number of female account holders. Evidence: A2 refers to districts; Female refers to gender = 'F'	Which district a female account holder is counted under: the client's own district (client.district_id) or the district of the account's branch (account.district_id, reached through disp)	For each district where female clients are registered (client.district_id), count those clients and list the nine districts with the most.	For each district where the accounts are held (the branch location in account.district_id), count the distinct female clients linked to those accounts through disp, and list the nine districts with the most.
BIRD `financial` batch2:b022	What proportion of customers who have accounts at the Prague branch are female? Evidence: Female refers to gender = 'F'; Percentage of female clients in Prague branch = count[female clients with accounts in Prague branch / count(clients with accounts in Prague branch)] * 100%; A3 may contain information about Prague	Which district_id decides that a customer is at the Prague branch: the client's own district_id (client.district_id) or the district_id of the account they hold (account.district_id, reached through disp)	Customers at the Prague branch are the clients whose own district is in the Prague region (client.district_id -> district.A3 = 'Prague'). The result is the share of those clients who are female.	Customers at the Prague branch are the clients linked through disp to an account whose branch district is in the Prague region (account.district_id -> district.A3 = 'Prague'). Each client is counted once, and the result is the share of those clients who are female.
BIRD `california_schools` batch3:b006	Name schools in Riverside which the average of average math score for SAT is grater than 400, what is the funding type of these schools? Evidence: Average of average math = sum(average math scores) / count(schools).	Which geographic unit the place name 'Riverside' refers to: the school district or the county.	Schools whose district name starts with 'Riverside', such as Riverside Unified. Average math score is computed per school and funding type, and schools above 400 are returned with their charter funding type.	Schools located in Riverside County. Average math score is computed per school and funding type, and schools above 400 are returned with their charter funding type.
BIRD `financial` batch3:b017	State different accounts who have account opening date before 1997 and own an amount of money greater than 3000USD	Which trans column stands for the 'amount of money greater than 3000 USD' the account owns: the transaction amount (trans.amount) or the account balance after a transaction (trans.balance)	Accounts opened before 1997 that have at least one transaction whose amount is over 3000 USD	Accounts opened before 1997 whose balance was over 3000 USD after at least one transaction
BIRD `thrombosis_prediction` batch3:b059	What are the symptoms observed by the youngest patient to ever did a medical examination? Identify their diagnosis. Evidence: The larger the birthday value, the younger the person is, and vice versa; symptoms observed refers to the symptoms is not NULL	Which table's Diagnosis column supplies 'their diagnosis': the patient-level diagnosis or the diagnosis recorded on the examination record	Report the diagnosis stored on the patient's record (Patient.Diagnosis) next to the symptoms from their examination.	Report the diagnosis recorded on the same examination where the symptoms were seen (Examination.Diagnosis).

Source	Question	The choice	Reading A	Reading B
BIRD `california_schools` batch4:b002	Please specify all of the schools and their related mailing zip codes that are under Avetik Atoian's administration.	Which administrator slots count as being 'under Avetik Atoian's administration': only the primary administrator columns (AdmFName1/AdmLName1), or any of the three administrator slots.	A school is under Avetik Atoian's administration only if they are listed as the primary administrator (administrator 1).	A school is under Avetik Atoian's administration if they are listed in any of the three administrator slots (1, 2, or 3).
BIRD `thrombosis_prediction` batch4:b049	The oldest SJS patient's medical laboratory work was completed on what date, and what age was the patient when they initially arrived at the hospital? Evidence: The larger the birthday value, the younger the person is, and vice versa; 'SJS' refers to diagnosis; (SUBTRACT(year(`First Date`)), year(Birthday)); age of the patients when they initially arrived at the hospital refers to year(Birthday)	Which Diagnosis column decides whether a patient is an 'SJS patient': Patient.Diagnosis or Examination.Diagnosis.	An SJS patient is one whose Patient.Diagnosis is exactly 'SJS'. Among these patients, take the oldest (earliest Birthday) and return a Laboratory date and their age at first arrival.	An SJS patient is one with an examination record whose Examination.Diagnosis mentions SJS. Among these patients, take the oldest (earliest Birthday) and return a Laboratory date and their age at first arrival.
BIRD `thrombosis_prediction` batch4:b053	How old was the patient who had the highest hemoglobin count at the time of the examination, and what is the doctor's diagnosis? Evidence: How old the patient refers to SUBTRACT(year(`Examination Date`), year(Birthday)); the highest hemoglobin count refers to MAX(HGB)	Which table supplies the doctor's diagnosis: Patient.Diagnosis or Examination.Diagnosis (both columns are documented as 'disease names')	The doctor's diagnosis is the patient's overall diagnosis stored in the Patient record (Patient.Diagnosis).	The doctor's diagnosis is the one recorded at the patient's examination (Examination.Diagnosis), joined by patient ID. The question's wording 'at the time of the examination' makes this a reasonable choice.
BIRD `thrombosis_prediction` batch4:b056	What was the gender of the first AORTITIS diagnosed patient? Evidence: gender means SEX; 'AORTITIS' refers to Diagnosis;	Which table's Diagnosis field (and matching date) defines when a patient was 'diagnosed' with AORTITIS: the patient-level Diagnosis ordered by `First Date`, or the examination-level Diagnosis ordered by `Examination Date`.	Use the patient record. Find patients whose Patient.Diagnosis is AORTITIS, take the one who first came to the hospital (earliest `First Date`), and return that patient's sex.	Use the examination records. Find examinations whose Diagnosis mentions AORTITIS, take the earliest by `Examination Date`, and return the sex of that examination's patient.
BIRD `financial` batch5:b007	What year was account owner number 130 born?	Which ID the number 130 in 'account owner number 130' refers to: the account ID or the client ID	130 is the account's number. The question asks for the birth year of the client linked to account 130 through disp. The gold has no filter requiring that client to be the OWNER.	130 is the owner's own number. The question asks for the birth year of the client whose client_id is 130.
BIRD `thrombosis_prediction` batch7:b057	For the patients who have an abnormal level of anti-DNA, please list the diseases they are diagnosed with. Evidence: abnormal level of anti-DNA refers to DNA > = 8; diseases refers to Diagnosis;	Which table's Diagnosis column supplies the diseases: Patient.Diagnosis or Examination.Diagnosis	List the distinct diagnoses recorded on the patient record (Patient.Diagnosis) for patients with at least one lab test where anti-DNA >= 8.	List the distinct diagnoses recorded on the examination records (Examination.Diagnosis) for patients with at least one lab test where anti-DNA >= 8.
BIRD `thrombosis_prediction` batch8:b005	When is the latest patient's medical data recorded? This patient should have an abnormal level of lactate dehydrogenase. Evidence: latest patient refers to ID with MAX('First Date'); abnormal level of lactate dehydrogenase refers to LDH > = 500;	Which date column answers 'when is the medical data recorded' for the patient with the latest First Date who has LDH >= 500: Patient.`First Date` or the Laboratory.Date of that patient's abnormal-LDH test	Among patients with any lab record where LDH >= 500, take the one with the latest First Date and return that First Date (the date the patient first came to the hospital).	Among patients with any lab record where LDH >= 500, take the one with the latest First Date and return the date of that patient's most recent lab test with LDH >= 500 (Laboratory.Date).
BIRD `thrombosis_prediction` batch8:b006	For the patients who are diagnosed with SJS, how many of them have a normal level of total protein? Evidence: diagnosed with SJS refers to Diagnosis = 'SJS'; normal level of total protein refers to TP > 6.0 and TP < 8.5;	Which table's Diagnosis column identifies patients 'diagnosed with SJS': Patient.Diagnosis or Examination.Diagnosis. The evidence gives Diagnosis = 'SJS' without naming a table, and both columns are documented as 'disease names'.	A patient is diagnosed with SJS if their patient record (Patient.Diagnosis) says SJS.	A patient is diagnosed with SJS if one of their examination records (Examination.Diagnosis) says SJS.
BIRD `thrombosis_prediction` batch8:b011	What is the disease name of the patient who has the highest level of triglyceride within the normal range? Evidence: disease name referse to Diagnosis; highest level of triglyceride within the normal range refers to MAX(TG < 200);	Which table's Diagnosis column supplies the 'disease name': Examination.Diagnosis or Patient.Diagnosis	Report the Diagnosis recorded on the Examination record(s) of the patient whose lab TG is the highest value below 200.	Report the Diagnosis recorded on the Patient record of the patient whose lab TG is the highest value below 200.
BIRD `thrombosis_prediction` batch8:b015	Among the patients who have a normal platelet level, how many of them have other symptoms observed? Evidence: normal platelet level refers to PLT > 100 and PLT < 400; have other symptoms refers to Diagnosis is not null;	Which table's Diagnosis column the evidence's 'Diagnosis is not null' refers to: Patient.Diagnosis or Examination.Diagnosis. Both tables have a column with that name and the same description, 'disease names'.	Count the Laboratory records with a normal platelet level (100 < PLT < 400) whose patient has a non-null diagnosis in the Patient table.	Count the Laboratory records with a normal platelet level (100 < PLT < 400) whose patient has a non-null diagnosis in an Examination record.
BIRD `thrombosis_prediction` bird_pairs:b029	How many female patients were given an APS diagnosis? Evidence: female refers to SEX = 'F'; APS diagnosis refers to Diagnosis='APS'	Which table's Diagnosis column should hold 'APS': the Patient record's diagnosis or the diagnosis recorded on an Examination?	Count female patients whose patient-level diagnosis (Patient.Diagnosis) is 'APS'.	Count distinct female patients who had at least one examination with the diagnosis 'APS' (Examination.Diagnosis).
Spider `car_1` spider_dev:spider_dev_0103	Which distinct car models are the produced after 1980?	Which column counts as the 'car model' to return: the general model name in model_list.Model (for example 'plymouth') or the full car name in car_names.Make (for example 'plymouth reliant').	A car model is the general model name kept in model_list.Model. List the distinct model names that have at least one car built after 1980.	A car model is the specific full car name kept in car_names.Make. List the distinct full car names that have a car built after 1980.
Spider `car_1` spider_dev:spider_dev_0122	What are the makers and models?	Which column supplies 'makers': the integer foreign key model_list.Maker, or the maker name car_makers.Maker reached by joining on that key	List each model in model_list with its Maker value as stored there, which is the maker's integer id	List each model with the name of its maker, taken from car_makers.Maker by joining on model_list.Maker = car_makers.Id
Referent column, settled by the wording				
BIRD `financial` batch2:b021	How many of the account holders in South Bohemia still do not own credit cards? Evidence: A3 contains the region names; South Bohemia is one of region names.	What 'do not own credit cards' maps to: the gold treats it as a disposition whose type is not 'OWNER' (disp.type != 'OWNER'); the agent treats it as having no row in the card table	Count the dispositions of clients living in South Bohemia whose disposition type is not 'OWNER' (users/disponents), treating the non-owner role as 'not owning a card'	Count the account holders in South Bohemia who have no credit card, meaning no card row is linked to their disposition
BIRD `financial` batch4:b020	Which accounts placed orders for household payment in Pisek? Evidence: k_symbol = 'SIPO' refers to household payment	Which table's k_symbol = 'SIPO' counts as placing an order for household payment: the transaction records in `trans` or the payment orders in `order`.	An account qualifies if it has at least one transaction (trans row) marked as a household payment (k_symbol = 'SIPO') and its branch is in the Pisek district.	An account qualifies if it has at least one payment order (`order` row) whose purpose is household payment (k_symbol = 'SIPO') and its branch is in the Pisek district.

Source	Question	The choice	Reading A	Reading B
BIRD `thrombosis_prediction` batch5:b054	What proportion of patients who had signs of thrombocytopenia had SLE diagnosed? Evidence: thrombocytopenia' refers to symptoms; 'SLE' refers to diagnosis; calculation = DIVIDE(SUM(DIAGNOSIS LIKE '%ITP%'), SUM(DIAGNOSIS LIKE '%SLE%')) MULTIPLY 100	How to identify patients with 'signs of thrombocytopenia': by the Examination.Symptoms text, or by a low platelet count (PLT < 100) in the Laboratory table	Patients have thrombocytopenia when their examination record lists 'thrombocytopenia' in the Symptoms column. This matches the evidence, which says thrombocytopenia refers to symptoms.	Patients have thrombocytopenia when a lab test shows a platelet count below the normal range (PLT < 100), which is the clinical sign of thrombocytopenia.
BIRD `thrombosis_prediction` batch6:b054	List each patient's ID and blood glucose (GLU) index that were within normal range for patient's whose data was first recorded in 1991. Evidence: blood glucose (GLU) index that were within normal range refers to GLU < 180; data that was first recorded in 1991 refers to year(Description) = 1991;	Which Patient date column defines the year the patient's data was 'first recorded' (to be filtered to 1991)	Patients whose First Date (the date they first came to the hospital) falls in 1991, listed with each of their GLU lab values below 180	Patients whose Description date (documented as the first date the patient's data was recorded, and named in the evidence as year(Description) = 1991) falls in 1991, listed with each of their GLU lab values below 180
BIRD `thrombosis_prediction` batch7:b052	How many patients with a normal anti-SSA came to the hospital before 2000? Evidence: normal anti-SSA refers to SSA IN('-','+-'); came to the hospital before 2000 refers to YEAR(`First Date`) < 2000; Should compute the number of distinct ones	Which date column decides 'came to the hospital before 2000': the lab test date (Laboratory.Date) or the patient's first visit date (Patient.`First Date`)	Count distinct patients who have a lab record with a normal anti-SSA result and whose lab test date falls before 2000.	Count distinct patients who have a lab record with a normal anti-SSA result and whose first visit to the hospital (Patient.`First Date`) was before 2000. This is what the evidence says.
Spider `car_1` spider_dev:spider_dev_0121	What are all the makers and models?	Which column should be returned for 'makers': the numeric maker foreign key stored in model_list.Maker, or the maker name from car_makers?	For every row in the model list, return the maker as recorded in that table (the numeric maker id in model_list.Maker) alongside the model name.	For every model in the model list, return the maker's name (car_makers.Maker, looked up through the maker id) alongside the model name.
Relationship role, open				
BIRD `toxicology` batch7:b015	Which molecule consisted of Sulphur atom with double bond? Evidence: sulphur refers to element - 's'; double bond refers to bond_type = ' = ';	Does the sulfur atom itself have to be part of the double bond, or is it enough that the molecule has both a sulfur atom and a double bond somewhere?	Return any molecule that has at least one sulfur atom and at least one double bond. The two do not need to be connected to each other.	Return molecules where a sulfur atom is one of the two atoms in a double bond.
Spider `network_1` spider_dev:spider_dev_0917	What is the lowest grade of students who do not have any friends?	Whether a student who appears in Friend only as friend_id (someone else lists them as a friend) counts as having a friend	A student has no friends if they never appear as student_id in the Friend table. Being listed as someone else's friend_id does not count.	A student has no friends only if they appear nowhere in the Friend table, as neither student_id nor friend_id. Friendship is treated as two-way.
Relationship role, settled by the wording				
BIRD `toxicology` batch2:b027	How many bond id have element iodine? Evidence: iodine refers to element = 'i'	How an iodine atom is tied to a bond: by the bond being in the same molecule as the atom (join through molecule_id), or by the atom being one of the bond's two endpoints (join through the connected table).	For each iodine atom, count every bond in the molecule that contains that atom. Nothing is de-duplicated, so a bond is counted once for each iodine atom in its molecule.	Count the distinct bonds that have an iodine atom as one of their two endpoints, as recorded in the connected table.
BIRD `toxicology` batch3:b027	What elements are in a double type bond? Evidence: double type bond refers to bond_type = '=';	How an atom must relate to a double bond to count: being an actual endpoint of that bond (linked through connected.bond_id), or just belonging to a molecule that has at least one double bond (linked through molecule_id).	List the distinct elements of every atom that has a connection and sits in a molecule containing at least one double bond, whether or not the atom is part of a double bond itself.	List the distinct elements of atoms that are actually one of the two atoms joined by a double bond.
BIRD `toxicology` batch7:b017	List all the elements with double bond, consisted in molecule TR024. Evidence: double bond refers to bond_type = '='; element = 'cl' means Chlorine; element = 'c' means Carbon; element = 'h' means Hydrogen; element = 'o' means Oxygen, element = 's' means Sulfur; element = 'n' means Nitrogen, element = 'p' means Phosphorus, element = 'na' means Sodium, element = 'br' means Bromine, element = 'f' means Fluorine; element = 'i' means Iodine; element = 'sn' means Tin; element = 'pb' means Lead; element = 'te' means Tellurium; element = 'ca' means Calcium	How an element is linked to the double bond: through any shared molecule (molecule_id join), or by the atom actually taking part in the double bond (through the connected table)	List the element of every atom in molecule TR024, given that TR024 has at least one double bond. Each atom row is repeated once per double bond in the molecule.	List the elements of the atoms in TR024 that are themselves an endpoint of a double bond ('=').
BIRD `toxicology` batch7:b023	What is the total number of molecules with double bonded oxygen? Evidence: oxygen refers to element = 'o'; double bond refers to bond_type = ' = ';	Whether the oxygen atom has to be one of the atoms in the double bond, or only has to be in a molecule that contains a double bond somewhere	Count the molecules that contain at least one oxygen atom and at least one double bond, with no requirement that the oxygen is part of that bond	Count the molecules where an oxygen atom is itself one end of a double bond (linked through the connected table, at either endpoint)
Strict cutoff, open				
BIRD `thrombosis_prediction` batch6:b041	Provide the ID and age of patient with lactate dehydrogenase (LDH) between 100-300 index above the normal range. Evidence: age refers to SUBTRACT(year(current_timestamp), year(Birthday)); lactate dehydrogenase (LDH) between 100-300 index above the normal range refers to LDH between 600 and 800;	Whether the LDH range endpoints 600 and 800 are included (strict vs. inclusive bounds)	Patients with any lab record where LDH is strictly greater than 600 and strictly less than 800	Patients with any lab record where LDH is between 600 and 800 inclusive (600 <= LDH <= 800)
Ties and selection, open				
BIRD `superhero` batch1:b042	Which hero was the fastest? Evidence: which hero refers to superhero_name; fastest refers to MAX(attribute_value) WHERE attribute_name = 'Speed';	How to handle the 40 heroes tied at the maximum Speed value: return only one of them or return all of them	Return a single hero: the first row after sorting by Speed value from highest to lowest (LIMIT 1). When several heroes tie, the one returned is arbitrary.	Return every hero whose Speed value equals the maximum Speed value, including all tied heroes.
BIRD `superhero` batch1:b043r	What is the race of the superhero with maximum attribute value? Evidence: maximum attribute value refers to MAX(attribute_value);	How to handle ties at the overall maximum attribute_value: return the race from one arbitrary top row (LIMIT 1), or return the race of every row that has the maximum value.	Find the single top row by attribute_value (one arbitrary hero among any ties) and return that hero's race.	Keep every hero_attribute row whose attribute_value equals the overall maximum, and return the race of the hero on each of those rows, not just one arbitrarily chosen row.
BIRD `thrombosis_prediction` batch1:b056	For the patient who has the highest Ig A within the normal range, what is his or her diagnosis? Evidence: highest Ig A within the normal range refers to MAX(IGA BETWEEN 80 AND 500);	How to handle several lab records tied at the maximum in-range IgA value (IGA = 500): return one row or every tied patient's diagnosis.	Return a single diagnosis, taken from one lab record with the highest IgA between 80 and 500. When several records tie, whichever one sorts first wins.	Return the diagnoses of all patients whose lab record matches the highest IgA between 80 and 500, so every tied patient is included.

Source	Question	The choice	Reading A	Reading B
BIRD `california_schools` batch2:b009	Which county reported the most number of school closure in the 1980s with school wonership code belonging to Youth Authority Facilities (CEA)? Evidence: Youth Authority Facilities (CEA) refers to SOC = 11; 1980s = years between 1980 and 1989	How to handle counties tied for the highest number of closures: return one arbitrary county, or return every tied county.	Return a single county: sort counties by their number of closed SOC 11 schools in 1980–1989 and keep only the top row, even though several counties tie.	Return every county that has the highest number of closed SOC 11 schools in 1980–1989, so all tied counties are listed.
BIRD `thrombosis_prediction` batch6:b050	Who is the oldest patient with the highest total cholesterol (T-CHO). State the patient ID and T-CHO index. Evidence: oldest patient refers to MIN(birthday); highest total cholesterol refers to MAX(T-CHO);	Which criterion is ranked first: highest T-CHO, with age only breaking ties, or oldest patient, with T-CHO taken as that patient's highest value	Find the highest T-CHO value across all lab records. If several patients share that value, pick the oldest one (earliest birthday). Return that patient's ID and the T-CHO value.	Find the oldest patient, meaning the one with the earliest known birthday who has lab records. Return that patient's ID and their highest T-CHO value.
BIRD `superhero` bird_pairs:b008	Who is the dumbest superhero? Evidence: the dumbest superhero refers to MIN(attribute_value) where attribute_name = 'Intelligence'	How to handle ties at the minimum Intelligence value: return one hero or every hero tied at the minimum	Return a single hero with the lowest Intelligence value, picking one arbitrarily if several are tied	Return every hero whose Intelligence value equals the minimum Intelligence value
BIRD `superhero` bird_pairs:b015r	Give the hero ID of superhero with the lowest attribute value. Evidence: lowest attribute value refers to MIN(attribute_value);	How to handle ties at the minimum attribute_value: return every tied hero, or return just one hero as the singular 'the hero ID' suggests.	Return the hero IDs of all hero_attribute rows whose attribute_value equals the overall minimum, so every tied hero is included (duplicates are possible when a hero has more than one attribute at the minimum).	Return exactly one hero ID whose attribute_value equals the overall minimum. When several heroes tie, any one of them may be returned.
Ties and selection, settled by the wording				
BIRD `financial` bird_pairs:b010	List out the accounts who have the earliest trading date in 1995 ?	How to handle ties at the earliest 1995 transaction date: return one arbitrary account, or every account that traded on that date	Return a single account: the one on the first transaction row when 1995 transactions are sorted by date (ties broken arbitrarily).	Return every distinct account that had a transaction on the earliest transaction date in 1995.
Spider `world_1` spider_dev:spider_dev_0741	Give the name of the nation that uses the greatest amount of languages.	Several countries tie for the most languages (12). The question does not say which of them to return, so the answer depends on how the tie is broken.	Return a single country with the most languages. Count rows grouped by country name and keep whichever tied country LIMIT 1 happens to pick first under that grouping ('United States').	Return a single country with the most languages. Count rows grouped by country code and keep whichever tied country LIMIT 1 happens to pick first under that grouping ('Canada').
Spider `dog_kennels` spider_dev:spider_dev_0955	Which dogs are of the rarest breed? Show their names and treatment dates.	How to handle ties when choosing the 'rarest breed', meaning several breeds share the minimum dog count	The rarest breed is a single breed: rank breeds by dog count and take the one with the fewest dogs. If several breeds tie, keep only one of them, chosen arbitrarily. Show the names and treatment dates of the treated dogs of that one breed.	Every breed whose dog count equals the minimum counts as rarest. Show the names and treatment dates of the treated dogs of all those tied breeds.
Time scope, open				
BIRD `financial` batch1:b012	For the client who applied the biggest loan, what was his/her first amount of transaction after opened the account?	Whether a transaction on the account-opening date counts as 'after opened the account' (inclusive vs. strictly later boundary)	Take the account with the biggest loan and return the amount of its earliest transaction, including any transaction made on the day the account was opened.	Take the account with the biggest loan and return the amount of its earliest transaction made strictly after the opening date, so transactions on the opening day are excluded.
Time scope, settled by the wording				
BIRD `financial` batch1:b015	How much is the average amount in credit card made by account holders in a month, in year 2021? Evidence: Operation = 'VYBER KARTOU' refers to credit card withdrawn	Which year to filter card-withdrawal transactions on: the gold uses 1998, while the question text says 2021.	Average amount of card withdrawal transactions (operation = 'VYBER KARTOU') made in 1998 by accounts that have a card. The year 1998 does not appear in the question or evidence; the question's '2021' seems to be a typo for a year the data actually covers.	Take the question literally: average amount of card withdrawal transactions made in 2021. The data only covers 1993-1998, so this returns NULL.
BIRD `thrombosis_prediction` batch2:b051	For all patients with normal uric acid (UA), what is the average UA index based on their latest laboratory examination result? Evidence: uric acid (UA) with normal range refers to UA < 8.0 and SEX = 'M' or UA < 6.5 and SEX = 'F'; average UA index refers to AVG(UA)	Which date counts as the 'latest laboratory examination': the single latest date in the whole Laboratory table, or each patient's own latest test date	Use only lab records dated on the latest date found anywhere in the Laboratory table. Because of how the gold's WHERE clause groups AND and OR, this date filter only applies to the male branch. Then average UA over the records with normal UA.	For each patient, take that patient's own most recent lab record. Keep it if its UA is in the normal range for the patient's sex, then average UA across those patients.
Value matching, open				
BIRD `financial` batch5:b011	How much, in total, did client number 617 pay for all of the transactions in 1998?	Which 1998 transactions count as the client having 'paid': every transaction on the client's accounts, or only outgoing ones (debits).	Add up the amounts of all 1998 transactions on the accounts linked to client 617, whatever their direction (credits and debits alike).	Add up only the outgoing 1998 transactions (type 'VYDAJ', i.e. withdrawals/debits) on the accounts linked to client 617, since money paid by the client is money leaving the account.
Value matching, settled by the wording				
BIRD `thrombosis_prediction` batch7:b035	Please list the disease names of the patients that have a proteinuria level higher than normal. Evidence: disease names refers to Diagnosis; proteinuria level higher than normal refers to `U-PRO` > = 30;	Whether the threshold `U-PRO >= 30` compares the TEXT column `U-PRO` as text (string order) or as numbers	Keep laboratory records where the U-PRO text value sorts at or after '30' in string order, which is what the literal evidence `U-PRO >= 30` does on a TEXT column in SQLite, and list the diagnoses of those patients	Keep laboratory records where the U-PRO value, read as a number, is at least 30 (proteinuria above the normal range 0 < N < 30), and list the diagnoses of those patients
BIRD `thrombosis_prediction` batch7:b056	Among the patients who has a normal level of anti-centromere and a normal level of anti-SSB, how many of them are male? Evidence: normal level of anti-centromere refers to CENTROMEA IN('-', '+-'); normal level of anti-SSB refers to SSB IN('-', '+-'); male refers to Sex = 'M'; Should consider DISTINCT in the final result;	Which stored codes count as a 'normal' anti-centromere and anti-SSB result: the stored values 'negative'/'0', or the literal symbols '-'/'+-' given in the evidence	A result is normal when the column holds the values the database actually uses for a negative result, 'negative' or '0'. Count the distinct male patients who have a lab record where both CENTROMEA and SSB are normal in this sense.	A result is normal when the column holds exactly the symbols the evidence lists, '-' or '+-'. Count the distinct male patients who have a lab record where both CENTROMEA and SSB hold one of those symbols.
Spider `car_1` spider_dev:spider_dev_0133	Which model saves the most gasoline? That is to say, have the maximum miles per gallon.	Whether the TEXT-typed MPG column is ranked by its stored string value or by its numeric value when picking the maximum	Rank cars by MPG as stored text (alphabetical order, so a value like '9' outranks '46.6') and return the model of the top-ranked car	Rank cars by MPG as a number (the true highest miles per gallon, 46.6) and return the model of that car

G CATALOGUE OF ACCEPTED REQUEST SETS

Table S6: Every accepted request set (204, 43 of them with a filter condition), by choice type. Open request: the text the models were given (each set also has one specific request per reading). Results: how the two reference answers differ on the database. Review: the author's reason for accepting the set. Models (Haiku, Sonnet, GPT-Sol, GPT-Luna), for the open request in one call with explanation: B both readings returned; N one reading, the choice named; S one reading, silent, although the model computes both when asked; – one reading, and the model does not compute both. Full sets with SQL, schema and grading: `benchmark/request_sets.jsonl`, by set id.

Set	Open request	Readings	Results	Review	Models
Aggregation level					
`7c4a89c45b alien`	Which RFI statuses (`rfistat`) have an above-average modulation index (`modindex`)? Ignore missing index values, and treat signals with no status as their own group. For each such status, show the status, its average modulation index, how far above average it is, and how it differs from the mean of the qualifying statuses' averages. Leave all values unrounded.	**A:** A status counts as above average if its mean modulation index exceeds the simple mean of the per-status means. **B:** A status counts as above average if its mean modulation index exceeds the mean modulation index over all signals.	A: 2 rows, B: 2 rows; first row differs; largest value change: 0.003134 (A) vs 0.00293 (B), 6.5%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	`-BBS`
`9bb5ac76bd alien`	Linking signals to their telescopes and observatories (signals that don't match are dropped), find the average modulation index and average phase shift for each telescope and weather profile (`weathprofile`). Keep the telescope/weather groups whose modulation index is above the overall average modulation index. Then, per weather profile, show the average modulation index and average phase shift of those groups, and how far above the overall average they are on average. Don't round anything.	**A:** The baseline is the average of the group-level averages, with every telescope/weather group counting equally. **B:** The baseline is the average modulation index over all linked signals, with every signal counting equally.	A: 3 rows, B: 3 rows; first row differs; largest value change: 185.7 (A) vs 184.6 (B), 0.59%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	`BBB-`
`cde68b2348 alien`	For observatory stations whose average telescope bandwidth usage (`bandusagepct`) is above average, what is their mean bandwidth usage, and how far above average are they on average? Ignore missing usage values, and give both figures unrounded in a single row.	**A:** Compare each station's average bandwidth usage with the unweighted mean of the station averages. **B:** Compare each station's average bandwidth usage with the mean bandwidth usage over all individual telescopes.	A: 1 row, B: 1 row; first row differs; largest value change: 24.9 (A) vs 25.03 (B), 0.52%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	`SBBN`
`d31fb48995 alien`	For each telescope in the signals table (`telescref`), work out its average pulse rate (`pulsepersec`), ignoring missing values, and keep the telescopes whose average is above the overall average pulse rate. Among those, keep only the ones whose average is also above the average of the kept telescopes' averages. In one row, give these final telescopes' mean average pulse rate, how far on average they sit above the overall average, and how far on average they sit above that second average. Don't round anything.	**A:** The overall baseline pulse rate is the average of the per-telescope average pulse rates, with each telescope weighted equally. **B:** The overall baseline pulse rate is the average pulse rate over all signal records, with each signal weighted equally.	A: 1 row, B: 1 row; first row differs; largest value change: 1.213 (A) vs 1.208 (B), 0.44%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	`SBSS`
`e7b2c8e815 alien`	Which observatory stations (`observstation`) have an average telescope bandwidth usage (`bandusagepct`) above the overall average bandwidth usage? Ignore missing values. For each one, show the station, its average bandwidth usage, and how much it exceeds the overall average. Return the top 10 by that margin, largest first, without rounding.	**A:** Each station's average bandwidth usage is compared against the mean of the station averages, with every station weighted equally. **B:** Each station's average bandwidth usage is compared against the mean bandwidth usage over all individual telescopes.	A: 10 rows, B: 10 rows; first row differs; largest value change: 49.42 (A) vs 49.54 (B), 0.26%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	`SBSS`
`1a73aa1ae0 archeology`	For environment records that match a site, show each site access status (`entrystat`) with its average ambient temperature (`ambictemp`) and average image count (`imgcount`). Also show how far that image-count average sits above or below the average image count overall (status average minus overall). Ignore missing values and don't round.	**A:** Compare each status's average image count with the mean of the status-level averages. **B:** Compare each status's average image count with the average image count over all matched environment records.	A: 3 rows, B: 3 rows; first row differs; largest value change: -11.54 (A) vs -10.94 (B), 5.2%	Overall can refer to matched records or the displayed status groups. B is more natural; A requires equal status weighting.	`SSSS`
`1d3fc42a2c archeology`	Match each spatial record (`scanspatial`) to the project whose `arcregistry` equals its `arcref`, and drop records with no matching project. What is each project's bounding-box X-to-Z ratio (`boxx` to `boxz`), ignoring missing values? Which 10 projects beat the simple average ratio of the projects in their own funding source (`projects.fundflux`) by the widest margin? Show the project ID, funding source, ratio and margin, sorted by margin from largest to smallest (missing last), with no rounding.	**A:** A project's ratio is the mean of its per-record boxx/boxz ratios. **B:** A project's ratio is its total boxx divided by its total boxz.	A: 10 rows, B: 10 rows; first row differs; largest value change: 27.87 (A) vs 27.96 (B), 0.29%	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	`NBNN`
`284b059491 archeology`	In the feature scan data, which 10 site artifact counts (`articount`, identified by site code `zoneref`) stand furthest above the average in their share of all artifacts? Each count's share is that count divided by the overall artifact total across all of `scanfeatures` (return nothing if the overall total is missing), and the average share is the mean of those shares (again return nothing if that average is missing). For each, give the site code, the artifact count, its share, and its share minus the average share. Ignore missing counts, and sort from the highest difference down with missing values last.	**A:** Total the artifact counts per site, then compute and rank each site's share against the average site share. **B:** Treat each feature record separately, then compute and rank each record's share against the average record share.	A: 10 rows, B: 10 rows; first row differs; 18 row keys in only one answer; largest value change: 0.002942 (A) vs 0.001054 (B), 64%	Site artifact counts suggests site totals, while each count and the named field suggest individual records. The intended row grain is unclear.	`BB--`

Set	Open request	Readings	Results	Review	Models
30b46b4254 archeology	Looking only at scans that match a site, what's the GB-per-minute rate of file size (gbsize) to scan duration (spanmin) for each site risk evaluation status (riskeval), treating scans at sites with no recorded status as a group of their own? Use the plain average of those group rates as a benchmark, with each group counting once. For the groups strictly above the benchmark (none if there's no benchmark), give me one row: their average rate and their average margin above the benchmark, unrounded.	**A:** Each risk group's rate is the mean of the individual scans' size-to-duration ratios. **B:** Each risk group's rate is the group's total file size divided by its total scan duration.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	SBBB
6a53a32ddd archeology	Using the spatial records (scanspatial) whose arcref matches a project's arcregistry and dropping any that don't match, look at each funding source (projects.fundflux). For each one, give the average relative difference between boxy and boxx, computed per record as (boxy − boxx) ÷ boxx, and give its area-to-volume ratio (aream2 to volm3). Ignore missing values. Which funding sources have a ratio above the simple average of that ratio across all funding sources? Show the funding source, the average Y-vs-X difference, the ratio, and how far the ratio is above that average, with no rounding.	**A:** A funding source's area-to-volume figure is the mean of each spatial record's area ÷ volume. **B:** A funding source's area-to-volume figure is its summed area divided by its summed volume.	A: 2 rows, B: 2 rows; first row differs; largest value change: 7.456 (A) vs 0.5305 (B), 93%	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	BBBS
95f2f15de7 archeology	For scans matched to their site (skip any whose site code isn't in sites), break things down by the site's health and safety status (healtheval). For each status, show the total number of scans (scancount) and the GB per minute of scanning, based on gbsize and spanmin, ignoring missing values. Only include statuses whose GB per minute is strictly above the average of that figure across statuses (none if that average is missing). Show the status, the total scans, the GB per minute, and how much it exceeds that average. No rounding.	**A:** Compute GB per minute for each scan, then average those ratios within each health/safety status. **B:** Divide each status's total gbsize by its total spanmin.	A: 2 rows, B: 2 rows; first row differs; 2 row keys in only one answer; largest value change: 0.02076 (A) vs 0.005421 (B), 74%	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	-BBB
d9ab5276bc archeology	Using the feature records (scanfeatures) and ignoring missing detected-feature counts (traitcount), which site codes' feature counts make up an above-average share of the overall feature total? For the top 10 by margin above the average share, show the site code, the feature count, the share, and how far the share exceeds the average share, largest margin first. Don't round.	**A:** Sum detected features per site, then compare each site's share with the average site share. **B:** Treat each feature record separately, and compare its share with the average record share.	A: 10 rows, B: 10 rows; first row differs; 20 row keys in only one answer; no shared value changes	Site codes' feature counts suggests totals per site, while using feature records and showing the count can imply individual records. The output grain is not explicit.	-S--
ef237a4024 archeology	Using only scans that match a project, what's each project's size-per-minute rate (gbsize per spanmin), grouped by project reference (scans.arcref) and funding source (fundflux)? Take the plain average of those project rates as a benchmark, with each project counting once. Show the 10 projects furthest above it, keeping only those strictly above, with project reference, funding source, rate and the gap above the benchmark, biggest gap first. Leave values unrounded.	**A:** A project's rate is the mean of its per-scan gbsize/spanmin ratios. **B:** A project's rate is its total gbsize divided by its total spanmin.	A: 10 rows, B: 10 rows; first row differs; largest value change: 4.376 (A) vs 4.382 (B), 0.15%	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	BBBN
f12d27cb47 archeology	Using mesh records (scanmesh) that match a site (sites.zoneregistry = scanmesh.zoneref), find each site's average number of mesh faces (facetfaces, ignoring missing values) along with its access status (sites.entrystat). Within each access status, take the simple average of its sites' averages, and focus on the sites strictly above their status's average; sites with no recorded access status are left out of this comparison. In a single row, give the average face count of those sites, how far on average they sit above the overall average number of mesh faces across all the matched data (sites without an access status included; if there is no such overall average, no sites count), and how far on average they sit above their status's average. No rounding.	**A:** The overall baseline is the average of the per-site average face counts, with each site weighted equally. **B:** The overall baseline is the average face count over every matched mesh record, with each record weighted equally.	A: 1 row, B: 1 row; first row differs; largest value change: 2.466e+06 (A) vs 2.48e+06 (B), 0.55%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	SBSS
fecab5d065 baseball_1	For players who received award votes (votes matched to players on player_id), grouped by death_country with a blank country as its own group: how well do players from each death country do on average in award voting, measured by the share of available points (points_won ÷ points_max on each vote, skipping any that can't be computed)? Also give each country's figure divided by the mean of all the country figures. List country, average share and relative ratio, unrounded.	**A:** Within each death country, every award-vote row is weighted equally when averaging points_won/points_max. **B:** Within each death country, each player's mean points_won/points_max is computed first, and the players are then weighted equally.	A: 8 rows, B: 8 rows; first row differs; largest value change: 0.1334 (A) vs 0.07359 (B), 45%	Per-vote points shares are specified, but country averaging can weight votes or first average within players. How well players do makes equal player weighting plausible.	-SBN
07335c94bb chinook_1	For each album (tracks with no album id count as one group), what's its bytes-per-millisecond rate (from Bytes and Milliseconds, ignoring missing sizes) and how many tracks does it have? Take the mean of those album rates across all albums. For the albums that come in above that mean, give me one row with their average rate, their average track count, and how far above the mean their rates are on average. No rounding.	**A:** Find each track's Bytes/Milliseconds ratio, then average those ratios within each album to get the album rate. **B:** Divide each album's total Bytes by its total Milliseconds to get the album rate. This makes longer tracks count for more.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	-BBS
1486e0d33f chinook_1	For each media type ID, show its average track unit price and how that compares to the overall average track price, as a ratio (its average divided by the overall average). Leave the numbers unrounded. List the top 10 by that ratio, highest first, with any missing ratios last.	**A:** Divide each media type's average track price by the mean of the per-media-type average prices (each media type weighted equally). **B:** Divide each media type's average track price by the average price over all tracks (each track weighted equally).	A: 5 rows, B: 5 rows; first row differs; largest value change: 0.8326 (A) vs 0.9421 (B), 12%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	SSSS

Set	Open request	Readings	Results	Review	Models
`1adc54e3cf chinook_1`	Which media types (by name, counting only tracks that have a media type) have a bytes-per-millisecond rate higher than the simple average of that rate across all media types? For each, show the media type name, its rate, and how much it exceeds that average, unrounded.	**A:** Each media type's rate is the mean of its tracks' individual Bytes/Milliseconds ratios; list the media types whose rate is above the unweighted mean of all media types' rates, with the rate and how far above it is. **B:** Each media type's rate is its total Bytes divided by its total Milliseconds; list the media types whose rate is above the unweighted mean of all media types' rates, with the rate and how far above it is.	A: 1 row, B: 1 row; first row differs; largest value change: 181.7 (A) vs 179.5 (B), 1.2%	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	`-BBN`
`57706f39b2 chinook_1`	For each media type, looking at all tracks, work out the average track unit price and its bytes per millisecond. Keep the media types whose bytes per millisecond is above the simple average across media types. Give one row with: the average of their average unit prices, the average of their bytes per millisecond, and on average how far their rate is above that cross-media-type average (rate minus average). No rounding.	**A:** Each media type's byte rate is the mean of its tracks' individual Bytes/Milliseconds ratios. **B:** Each media type's byte rate is its tracks' total Bytes divided by their total Milliseconds.	A: 1 row, B: 1 row; first row differs; largest value change: 181.7 (A) vs 179.5 (B), 1.2%	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	`NBBN`
`5a20154d84 chinook_1`	Looking only at tracks that have a matching media type, how many bytes per millisecond (`Bytes` vs `Milliseconds`) does each media type name come out at? Also show that figure divided by the simple average of the figures of all media types (each media type counted once). Leave the numbers unrounded.	**A:** Each media type's rate is the mean of its tracks' own Bytes/Milliseconds ratios. **B:** Each media type's rate is its total Bytes divided by its total Milliseconds, so long tracks count for more.	A: 5 rows, B: 5 rows; first row differs; largest value change: 35.16 (A) vs 33.57 (B), 4.5%	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	`-BBN`
`5e80a4dead chinook_1`	Looking only at tracks that have a matching media type, what is each media type's bytes-per-millisecond rate (`Bytes` vs `Milliseconds`), and how many tracks does it have? Show only the media types whose rate is above the simple average of all media types' rates (each type counted once). For those, give the name, the rate, the track count, and how far the rate is above that average. Don't round anything.	**A:** Each media type's rate is the mean of its tracks' own Bytes/Milliseconds ratios. **B:** Each media type's rate is its total Bytes divided by its total Milliseconds.	A: 1 row, B: 1 row; first row differs; largest value change: 181.7 (A) vs 179.5 (B), 1.2%	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	`BBBB`
`6f949de77e chinook_1`	Consider only the media types (by name, counting only tracks that have a media type) whose bytes-per-millisecond rate is above the simple average of that rate across all media types. In one unrounded row, what are their average rate, their average number of tracks, and their average margin above that overall average?	**A:** Per media type, rate = the mean of its tracks' Bytes/Milliseconds ratios; among media types above the unweighted mean rate, report the average rate, average track count, and average excess. **B:** Per media type, rate = its total Bytes / its total Milliseconds; among media types above the unweighted mean rate, report the average rate, average track count, and average excess.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	`-BBN`
`fc09eb5c8c chinook_1`	Which 10 albums have the highest bytes-per-millisecond rate compared with the average across all albums? Group tracks by `AlbumId`, with tracks that have no album id counted together as one group. For each album, show the album id, the average track `UnitPrice`, the album's average bytes per millisecond, and how far it is above the all-album average. Ignore missing values, don't round, and sort by that difference, largest first.	**A:** Each album's rate is the mean of the per-track ratios Bytes/Milliseconds; compare it with the average of these rates across albums and show the top 10 albums by how far above the average they are. **B:** Each album's rate is the sum of its tracks' Bytes divided by the sum of their Milliseconds; compare it with the average of these rates across albums and show the top 10 albums by how far above the average they are.	A: 10 rows, B: 10 rows; first row differs; largest value change: 161.4 (A) vs 161.8 (B), 0.26%	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	`-BBS`
`43277f1d9e crypto`	Using fee records matched to their orders (the order's `recordvault` equals the fee's `orderslink`, compared exactly as text; unmatched fees dropped), which trading pairs (`mktnote`, with records lacking a pair treated as one group) have an average net fee rate (`feerate` minus `rebrate`) strictly above the overall average net fee rate? Give me one row for those pairs: the mean of their average net fee rate, and the mean amount by which they exceed the overall average. If the overall average is missing, no pairs count. Skip missing values in the averages (missing if there's nothing to average) and don't round anything.	**A:** The benchmark is the average of the per-trading-pair average net fee rates, so each pair counts equally. **B:** The benchmark is the average net fee rate over all matched fee records pooled together, so each fee record counts equally.	A: 1 row, B: 1 row; first row differs; largest value change: 3.76e-05 (A) vs 3.726e-05 (B), 0.9%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	`SSSS`
`43cd7485bf crypto`	For fees tied to orders (`fees.orderslink` = `orders.recordvault`), which order channels (`orders.orderbase`) have an above-average fee rate (`feerate`, ignoring missing values)? Give me one row: those channels' mean average fee rate and, on average, how far above average they are, unrounded.	**A:** Compare each channel's average fee rate to the simple mean of all channel averages. **B:** Compare each channel's average fee rate to the overall average fee rate across all linked fee records.	A: 1 row, B: 1 row; first row differs; largest value change: 1.808e-05 (A) vs 1.812e-05 (B), 0.22%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	`SBBB`
`50dc01d833 crypto`	Using the order executions, find each order's (`ordersmark`) average unfilled quantity (`remaincount`), ignoring missing values; executions with no order reference count together as one order. Which orders are above the overall average unfilled quantity, and by how much? Among those orders, keep only the ones that are also above the group's own average, and note how far above it they are. Return one row with three values, without rounding: the average of these orders' averages, their average excess over the overall average, and their average excess over the second average.	**A:** The first benchmark is the mean of the per-order average unfilled quantities. **B:** The first benchmark is the mean unfilled quantity over all execution records.	A: 1 row, B: 1 row; first row differs; largest value change: 3.984 (A) vs 3.984 (B), 0.01%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	`SBBS`

Set	Open request	Readings	Results	Review	Models
7d5e124519 crypto	Match order executions to their orders (orderexecutions.ordersmark = orders.recordvault; skip unmatched executions) and get each order's average filled quantity (fillcount, ignoring missing values) and its time in force (timespan). Which orders have an average fill strictly above their time in force's average fill? For each, show ordersmark, timespan, the order's average fill, how far above the time-in-force average it is, and its rank within its time in force by average fill (highest first, missing last, ties share a rank and the next rank skips). Keep only orders ranked 3 or better. No rounding.	**A:** The time-in-force benchmark averages the per-order average fills, with each order weighted equally. **B:** The time-in-force benchmark averages fill quantity directly over all execution records in that time in force.	A: 12 rows, B: 12 rows; first row differs; largest value change: 6.929 (A) vs 6.928 (B), 0.016%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	BBBS
9a9ca0e90a crypto	Looking at fees per order (orderslink, with fee records that have no order treated as one group), which orders have an average net fee rate (feerate minus rebrate) above the overall average net fee rate? For those orders, give me one row: the mean of their average total rebate (rebtotal), the mean of their average net fee rate, and the mean amount by which they exceed the overall average net fee rate. Skip missing values in the averages and don't round anything.	**A:** The benchmark is the average of the per-order average net fee rates, with each order weighted equally. **B:** The benchmark is the average net fee rate over all individual fee records, so orders with more fee records weigh more.	A: 1 row, B: 1 row; first row differs; largest value change: 0.000529 (A) vs 0.0005288 (B), 0.025%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	BBBS
9e5dda7b4a crypto	For each account scope, using fees tied to orders (fees.orderslink = orders.recordvault) and to the users who placed them (orders.userlink = users.userstamp), with unmatched rows dropped, show the scope, the average net fee rate (fee rate minus maker rebate rate), how many it was averaged over, and how far that average sits from the simple mean of the scope averages. Don't round.	**A:** Per account scope, average the net fee rate over every fee record and count fee records. **B:** Per account scope, collapse fee records to one net rate per order (averaging within the order), then average across orders and count orders.	A: 4 rows, B: 4 rows; first row differs; largest value change: -2.461e-06 (A) vs -2.327e-06 (B), 5.4%	How many it was averaged over does not identify fees versus orders. Fee-level averaging is more literal, but equal order weighting is a plausible alternative.	SSSS
fd1f07ae25 crypto	Linking fees to their orders (fees.orderslink = orders.recordvault, skipping fees without a matching order), for each orderslink and order type (orders.ordertune) combination, what are the average total rebate (rebtotal) and the average relative rebate-to-fee gap, (rebrate − feerate) / feerate? Show the combinations whose average rebate is above the overall average rebate, with orderslink, order type, average rebate, average relative gap, and how far above the overall average they are. List the top 10 by that margin, largest first, unrounded.	**A:** Compare each group's average rebate to the mean of the group averages. **B:** Compare each group's average rebate to the mean rebate across all matched fee records.	A: 10 rows, B: 10 rows; first row differs; largest value change: 199.1 (A) vs 199.1 (B), 0.003%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	SBSS
6c1b7822e6 customers_and_invoices	For each order in Order_Items, how does its average product_quantity (ignoring missing values) compare with the overall average quantity? Show the order id, the order's average quantity, and the order's average minus the overall average, with results only if the overall average exists. Sort by the difference from largest to smallest, with missing differences last, and keep the top 10.	**A:** The baseline is the average of the per-order average quantities. **B:** The baseline is the average quantity over all individual order items.	A: 8 rows, B: 8 rows; first row differs; largest value change: 0.2917 (A) vs -0.3333 (B), 1.9e+02%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	SBNS
b4231a9afb customers_and_invoices	For each invoice in the invoice line items, find its average product_quantity and its relative gap between product cost and total cost, i.e. (derived_product_cost − derived_total_cost) ÷ derived_total_cost. Keep only invoices whose average quantity is above the average of all the invoices' average quantities, and list the invoice number, its average quantity, that relative gap, and how far its average quantity is above that overall average. Sort by that difference from largest to smallest (missing last), show the top 10, and don't round. Ignore missing values in averages.	**A:** Each invoice's cost gap is the mean of the per-line ratios (product cost − total cost) ÷ total cost. **B:** Each invoice's cost gap is (summed product cost − summed total cost) ÷ summed total cost over its lines.	A: 4 rows, B: 4 rows; largest value change: 1.413 (A) vs 1.516 (B), 6.8%	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	BBBS
bd51f2bd3c customers_and_invoices	For each invoice in Invoice_Line_Items, show the invoice number, the invoice's ratio of product cost (derived_product_cost) to total cost (derived_total_cost), how many line items it has, and that ratio divided by the average ratio across all invoices. Ignore missing values throughout, and return rows only when that overall average exists. Sort by the relative figure from highest to lowest, with missing values last, show the top 10, and don't round anything.	**A:** Each invoice's cost ratio is the mean of its per-line product-cost ÷ total-cost ratios. **B:** Each invoice's cost ratio is its summed product cost ÷ its summed total cost.	A: 9 rows, B: 9 rows; first row differs; largest value change: 1.834 (A) vs 2.032 (B), 9.7%	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	BBBN
ed587ea981 customers_and_invoices	In Products, find each production_type_code's average unit_price (ignoring missing prices) and keep the types whose average is above the overall average price, if that average exists. Next, take the simple mean of the kept types' averages, and keep only the types also above that mean. Return one row with three values, each averaged over the remaining types: their type average price, how far their type average is above the overall average price, and how far it is above the mean of the kept types.	**A:** The overall baseline price is the mean of the per-type average prices. **B:** The overall baseline price is the mean unit price over all individual products.	A: 1 row, B: 1 row; first row differs; largest value change: 265.3 (A) vs 299.5 (B), 11%	Overall average does not specify product weighting versus equal weighting of type averages; B is the more usual default.	SSSS
fbb044ab6b customers_and_invoices	For each invoice, look at its average line product cost (derived_product_cost) and its ratio of total cost (derived_total_cost) to unit price (product_price), ignoring missing values. Take the invoices whose ratio is above the average ratio across all invoices. In a single row, what are their average product cost, their average ratio, and on average how far their ratio sits above that overall average? No rounding.	**A:** Each invoice's ratio is the mean of its line-level total-cost ÷ unit-price ratios. **B:** Each invoice's ratio is its summed total cost divided by its summed unit price.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	BBBB
1e89ccdc51 formula_1	Looking at race results by race name (races.name, matched to results by raceId), which race names have a total of points above the average total across race names? For those race names together, give me in one row their combined total points, their average laps completed, and the average amount by which their totals exceed the average total.	**A:** For race names with above-average total points, the laps figure is the plain mean of each race name's average laps. **B:** For race names with above-average total points, the laps figure is the mean laps over all their individual result rows, so race names with more results weigh more.	A: 1 row, B: 1 row; first row differs; largest value change: 44.51 (A) vs 42.83 (B), 3.8%	Their average laps can mean an equal-weight mean of race-name averages or pooled result-row laps. The requested combined total does not fix the weighting of this separate average.	BBBN

Set	Open request	Readings	Results	Review	Models
c759056a31 formula_1	For each circuit country, what are the average constructor standings points and wins (the season-to-date figures) at races held there? Also, how far is each country's average wins above or below the average of the country figures?	**A:** A country's averages are the mean of per-race averages, so each race counts equally. **B:** A country's averages are the plain mean over all of its standings records pooled together.	A: 31 rows, B: 31 rows; first row differs; largest value change: 0.039 (A) vs -0.03776 (B), 2e+02%	At races held there can support averaging races equally or pooling standings rows. No race-first aggregation or weighting rule is specified.	S−−S
fe2c3edae0 formula_1	For each circuit (`races.circuitId`, with its `circuits.country`, counting only races whose circuit is in the circuits table), what is the average round number (`races.round`) at which it hosts races, and how far is that above or below its country's average round? List circuit id, country, the circuit's average round and the difference (circuit minus country). Sort by the difference from largest to smallest (missing last), show the top 10, and don't round.	**A:** Compare each circuit's average round with the unweighted mean of its country's circuit averages. **B:** Compare each circuit's average round with the average round of all races in its country.	A: 10 rows, B: 10 rows; first row differs; 8 row keys in only one answer; largest value change: 4.079 (A) vs 2.211 (B), 46%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	SBBB
13e92662d1 insider	For each transaction that has sentiment data (matching `transref` to `transactionrecord.transreg`, dropping anything unmatched), along with its order type mix, show the transaction reference, the order type mix, its average short interest days to cover and average option volume ratio (ignoring missing values), and how its average short interest compares with the average short interest for its order type mix, as a ratio. Do not round. List the 10 highest ratios first, with missing ratios last.	**A:** The order-type baseline is the mean of the per-transaction averages, with each transaction weighted equally. **B:** The order-type baseline is the mean of short interest over the raw matched snapshot records of that order type.	A: 10 rows, B: 10 rows; first row differs; largest value change: 2.021 (A) vs 2.023 (B), 0.1%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	SBBB
56f991749c insider	Which broker reporting statuses (`brokrepstat`) have above-average model confidence (`modelconf`)? Treat cases with no status as their own status. For each one, show the status, its average model confidence, how far above average it is, and its average divided by the mean of the qualifying statuses' averages. Ignore missing values when averaging, return nothing if a reference value is missing, and don't round.	**A:** A status qualifies if its average model confidence exceeds the mean of all the status-level averages, with each status weighted equally. **B:** A status qualifies if its average model confidence exceeds the average model confidence taken over every individual compliance case.	A: 2 rows, B: 2 rows; first row differs; largest value change: 0.0001236 (A) vs 0.0001505 (B), 18%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	SSBS
86ce9ada6a insider	Link the sentiment data to transaction records on `transref` = `transreg`, dropping anything without a match and anything with no order type mix (`ordertypedist`), and look at institutional ownership (`instownpct`, ignoring missing values) against each transaction's order type mix. Take the cases whose ownership is strictly above the overall average ownership. In one unrounded row, give their average ownership, the average ratio of their ownership to the average ownership for their order type mix, and the average amount by which they exceed the overall average.	**A:** Each transaction is one unit, and its ownership is the average across the snapshots that point to it. **B:** Each sentiment snapshot is its own unit, even when several snapshots point to the same transaction.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Cases could be individual sentiment snapshots or transactions summarized over their snapshots. The transaction linkage does not define the unit being filtered and averaged.	SSBS
e59fa9440b insider	Looking at announcement timing (`eventannotm`, with a missing timing as its own group), which timings have an average institutional ownership strictly above average? Ignoring missing ownership values, give in one row the mean of those timings' average ownership and how far above average they are on average, without rounding.	**A:** Compare each timing group's average institutional ownership with the simple mean of the group averages. For the groups above it, report the mean of their averages and their mean excess. **B:** Compare each timing group's average institutional ownership with the overall per-snapshot average. For the groups above it, report the mean of their averages and their mean excess.	A: 1 row, B: 1 row; first row differs; largest value change: 0.841 (A) vs 0.8387 (B), 0.27%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	SBBB
06768b2b94 sakila_1	For each film rating, what are the average replacement cost and the average rental price per day of rental, and how does each rating's per-day price compare, as a ratio, to the plain average of that figure across the ratings (each rating counted once)? Leave the numbers unrounded.	**A:** For each rating, the per-day rental price is the average over its films of rental_rate / rental_duration. **B:** For each rating, the per-day rental price is AVG(rental_rate) / AVG(rental_duration) over its films.	A: 5 rows, B: 5 rows; first row differs; largest value change: 0.6685 (A) vs 0.5972 (B), 11%	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	SBSN
293277ae13 sakila_1	For each film rating, look at its average film length (ignoring films with no length) and its rental rate per day of rental. Which ratings have a per-day rate above the simple mean of the ratings' per-day rates, with each rating counting once? For just those ratings, give one row with their average length, their average per-day rate, and the average amount by which they exceed that mean. Don't round anything.	**A:** For each rating, work out each film's rental rate divided by its rental duration, then average those per-film values. **B:** For each rating, divide the average rental rate by the average rental duration.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	SBSS
6d0476914f sakila_1	Which film ratings have an average `length` (ignoring missing lengths) that is longer than average? Of those ratings, keep only the ones that also run longer than the mean of their own average lengths. Give one row with: the mean of the remaining ratings' average lengths, how far on average they sit above that first average, and how far on average they sit above the mean of the ratings that passed the first cut. No rounding.	**A:** Each rating's average length is compared with the unweighted mean of the per-rating averages. **B:** Each rating's average length is compared with the mean length over all individual films.	A: 1 row, B: 1 row; first row differs; largest value change: 5.366 (A) vs 5.172 (B), 3.6%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	SNSS
796562a668 sakila_1	For each film `rating`, what is the average rental price per day (`rental_rate` per day of `rental_duration`), ignoring missing values? Average those per-rating figures across ratings and keep the ratings strictly above that average (none if it is missing). Give me one row: the average per-day price of those ratings and how far above the across-rating average they are on average, ignoring missing values and unrounded.	**A:** Divide each film's rental_rate by its rental_duration, then average those ratios within each rating. **B:** Within each rating, average rental_rate and average rental_duration separately, then divide the first by the second.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	SSSS

Set	Open request	Readings	Results	Review	Models
8f7229f972 sakila_1	For each film rating, what is the ratio of rental price to replacement cost? Show the rating, that ratio, the number of films, and how far the ratio is above or below the plain average of the ratios across all ratings (the rating's ratio minus that average). Averages ignore missing values. Don't round.	**A:** The rating's ratio is the average of the per-film rental-price-to-replacement-cost ratios. **B:** The rating's ratio is the rating's average rental price divided by its average replacement cost.	A: 5 rows, B: 5 rows; first row differs; largest value change: 0.0007622 (A) vs -0.001679 (B), 1.5e+02%	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	SBBS
a2c14da395 sakila_1	For each film rating, look at its average replacement cost and its rental rate per day of rental duration. Which ratings have a per-day rate above the simple average across ratings? For those ratings, give one row showing their average replacement cost, their average per-day rate, and on average how far that rate is above the cross-rating average. Don't round.	**A:** Each rating's daily rental price is the mean of the per-film ratio rental_rate / rental_duration. **B:** Each rating's daily rental price is the rating's mean rental_rate divided by its mean rental_duration.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	SBNN
0d34b6c962 student_club	For each event in the budget table (budget.link_to_event), find its spend ratio, meaning how much was spent relative to the amount budgeted across its budget lines, ignoring missing values. Compare each event with the average spend ratio across all events: show the event id, its spend ratio, and its ratio divided by that average. Sort by that relative figure, highest first (missing values last), and give the top 10. No rounding.	**A:** Per event, average the spent/amount ratios of its budget lines; compare each event's value with the mean across events; top 10 by that relative figure. **B:** Per event, divide total spent by total budgeted; compare each event's value with the mean across events; top 10 by that relative figure.	A: 10 rows, B: 10 rows; first row differs; 2 row keys in only one answer; largest value change: 0.799 (A) vs 1.001 (B), 20%	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	BBSN
14054d8e5e student_club	For budget lines whose total expense cost (summing cost by link_to_budget) is above the average spend per budget line, what is their average total, how many expenses do they have altogether, and on average by how much do they exceed that average?	**A:** The threshold is the average total expense cost over budget lines that have expenses. **B:** The threshold is the average total expense cost over all budget lines, with lines that have no expenses counting as 0.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Average spend per budget line does not specify whether lines without expenses are included with zero cost or excluded. This is a baseline-population choice.	SBBN
37095ad966 student_club	Using budget lines matched to their event (drop lines with no matching event), first work out for each event the average remaining and the average overspend rate, (spent − amount) ÷ amount, across its lines. Then, for each event status, average those event-level figures. Show the status, its average remaining, its average overspend rate, and how its average remaining compares to the overall average remaining, as a ratio. Ignore missing values in the averages, and don't round anything.	**A:** The ratio column divides each status's average remaining by the mean of the status-level average remaining values, with each status weighted equally. **B:** The ratio column divides each status's average remaining by the mean of the event-level average remaining values over all events regardless of status, with each event weighted equally.	A: 3 rows, B: 3 rows; first row differs; largest value change: 1.895 (A) vs 1.904 (B), 0.47%	The request defines both event and status aggregates but does not specify which level supplies the overall baseline.	SNSS
bd67ddb5d9 student_club	For each event (budget.link_to_event), show the average spent and the average remaining across its budget lines. Also show how far its average remaining is above the overall average remaining (its average minus the overall average). Ignore missing values in the averages. List the 10 events with the largest difference, highest first, with missing values last. Don't round.	**A:** Each event's average remaining is compared with the mean of the per-event average remaining values, so every event carries equal weight. **B:** Each event's average remaining is compared with the mean of remaining over every budget line, so every budget line carries equal weight.	A: 10 rows, B: 10 rows; first row differs; largest value change: -10.35 (A) vs -15.22 (B), 32%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	SBBS
bf298fafe5 student_club	For each event status in the budget table, what's the average amount left over (amount budgeted minus amount spent, ignoring missing values), and how does it compare with the overall average? Show the status, its average, and its difference from the overall average, unrounded, and return nothing if there is no overall average to compare against.	**A:** Compare each status's average leftover (amount − spent) with the unweighted mean of the per-status averages. **B:** Compare each status's average leftover (amount − spent) with the average leftover over all budget rows taken together.	A: 3 rows, B: 3 rows; first row differs; largest value change: 4.347 (A) vs 5.061 (B), 14%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	SSSS
df956358d7 student_club	Which events have an above-average budget amount? For each event (budget.link_to_event), using the average amount across its budget lines and ignoring missing amounts, show the event link, its average amount, how much it beats the overall average budget amount, and its average minus the mean among those above-average events. List the top 10 by that last gap, largest first.	**A:** The overall average budget amount is the average of the per-event average amounts, with each event weighted equally. **B:** The overall average budget amount is the average amount over all budget lines, so events with more lines weigh more.	A: 9 rows, B: 9 rows; first row differs; largest value change: 21.47 (A) vs 21.73 (B), 1.2%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	SBS-
2ab3b368e4 thrombosis_prediction	For each patient in Laboratory, show the patient ID, their average GOT, their average RBC, and their average GOT minus the overall average GOT. Use that one overall figure for every patient, and return nothing if it is missing. Ignore missing values. Sort by the difference, highest first with missing differences last, return the top 10, and don't round.	**A:** Compare each patient's average GOT with the mean of the per-patient averages. **B:** Compare each patient's average GOT with the mean GOT taken over all lab records.	A: 10 rows, B: 10 rows; first row differs; largest value change: 43.66 (A) vs 41.6 (B), 4.7%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	SBSS
4a2d376afa thrombosis_prediction	Among patients whose average hemoglobin (HGB) across their lab records (ignoring missing values) is above the overall average hemoglobin, list the patient ID, their average HGB, how far it is above the overall average, and their average minus the average of these above-average patients' averages. Show the top 10 by that last value, highest first, unrounded.	**A:** The overall HGB average is an average of patient averages, so each patient counts once. **B:** The overall HGB average is taken over every lab record directly, so patients with more tests weigh more.	A: 10 rows, B: 10 rows; first row differs; largest value change: 1.652 (A) vs 1.758 (B), 6%	The comparison baseline does not specify equal weighting of group averages versus pooling the underlying records. Both are plausible; pooled records give more weight to groups with more observations.	SBBS

Set	Open request	Readings	Results	Review	Models
568a31721e thrombosis_prediction	Looking at each patient's hemoglobin-to-hematocrit (HGB/HCT) ratio from their lab records (ignoring missing values), take the patients whose ratio is above the average across all patients. Give one summary row for them: the mean of their average PIC, the mean of their ratios, and the mean amount by which their ratios exceed that cross-patient average.	**A:** Each patient's HGB/HCT ratio is the mean of the ratios computed for each of their lab records. **B:** Each patient's HGB/HCT ratio is their mean HGB divided by their mean HCT.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The patient-level ratio is not defined: averaging per-record ratios and dividing patient-level means are both plausible. Missing-value pairing also needs clarification.	SNBN
7fd142e2d9 thrombosis_prediction	By patient sex, what are the average aCL IgG and ANA values in the examination data (only examinations linked to a patient, missing values ignored)? Also give each sex's average ANA as a ratio of the mean of the per-sex ANA averages, unrounded.	**A:** Average aCL IgG and ANA over all matched examination records within each sex. Each examination counts once. **B:** Average each patient's examinations into per-patient values first. Then average those values within each sex, so each patient counts once.	A: 2 rows, B: 2 rows; first row differs; largest value change: 3.281 (A) vs 3.334 (B), 1.6%	The examination-data wording favors equal examination weighting; the patient-sex framing also permits equal patient weighting. No intermediate patient aggregation is specified.	SSSS
c94de1efe9 thrombosis_prediction	Match laboratory records to patients on ID (lab records with no matching patient are dropped) and group by patient ID and sex. For each patient, find their average white blood cell count (WBC) and their average hemoglobin-to-hematocrit ratio (HGB/HCT), ignoring missing values. Keep only patients whose ratio is greater than the mean of that ratio across all patients (missing values ignored). Show the patient ID, sex, average WBC, the ratio, and the ratio minus that mean. Sort by that gap, largest first with missing values last, return only the first 10 rows, and don't round anything.	**A:** For each patient, compute HGB/HCT on each lab record and then average those ratios. **B:** For each patient, divide their average HGB by their average HCT.	A: 10 rows, B: 10 rows; first row differs; 2 row keys in only one answer; largest value change: 0.01235 (A) vs 0.01192 (B), 3.5%	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	SNSB
de5cec30a7 thrombosis_prediction	Linking examinations to patients by ID (dropping examinations with no matching patient), how far is patients' aCL IgM above or below the overall mean? Within each sex (exact text match, so a blank sex is its own group and a missing sex matches nothing), what is the average of those deviations? Show the top 10 entries whose deviation is above their sex's average deviation, ranked by how far above it they are, largest first (missing last): patient ID, SEX, aCL IgM, its deviation from the overall mean, and how far that deviation exceeds the sex's average deviation. All averages ignore missing values; don't round.	**A:** Average aCL IgM per patient (ID, sex) first, then compute deviations from the mean of those patient averages, compare them with sex-group average deviations, and rank the patients. **B:** Use each examination's own aCL IgM value, compute deviations from the mean over all examinations, compare them with sex-group average deviations, and rank the examination rows.	A: 10 rows, B: 10 rows; first row differs; largest value change: 1.68 (A) vs 1.681 (B), 0.059%	Patients' values suggests one aggregate per patient, while top entries and the raw measurement field suggest examination rows. The ranking unit is not explicitly fixed.	SSSS
ef4d2acbf1 thrombosis_prediction	Work out each patient's hemoglobin-to-hematocrit (HGB/HCT) ratio based on their lab tests. Keep the patients whose ratio is above the average across all patients, noting how far above they are. Then, within that group, keep those who are also above the group's own average, noting how far above that they are. Give one row with the averages, over these final patients, of the ratio, the amount above the overall average, and the amount above the group average, unrounded.	**A:** Each patient's ratio is the average of their per-test HGB/HCT ratios; then apply the two above-average filters and average the results. **B:** Each patient's ratio is their mean HGB divided by their mean HCT; then apply the two above-average filters and average the results.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The within-group ratio is not operationally defined: averaging record-level ratios and dividing aggregated numerator by aggregated denominator are both plausible. They weight records differently; missing-value pairing also needs a policy.	SNBN
Counting unit					
10cf8f1547 alien	For each power supply status (powerstatus) of our telescopes, with a missing status counted as its own group, what share of all telescopes does it account for? For the statuses whose share is above the average share across statuses, give one row with the total number of telescopes in them, their average share, and how far above the average share they are on average. Don't round anything.	**A:** Each power status's share of all telescopes is written as a fraction from 0 to 1. **B:** Each power status's share of all telescopes is written as a percentage from 0 to 100.	A: 1 row, B: 1 row; first row differs; largest value change: 0.3435 (A) vs 34.35 (B), 99%	Share can be expressed as a fraction or a percentage, and no scale is specified. This is an output-unit choice, not a different counting unit; label the chosen scale.	SSN-
11c73410e6 alien	Link each signal to its telescope (signals.telescref = telescopes.telescregistry), and drop signals with no matching telescope. For each equipment status (telescopes.equipstatus), what is the average pulse rate (signals.pulsepersec) of the telescopes in that status, based on the signals they detected and ignoring missing pulse rates? Keep the statuses whose average is above the mean of all the status averages, and note how far above it each one is. From those, keep only the statuses whose average is above the mean of the kept ones, again noting the excess. Give one row with three values: the average pulse rate of the final statuses, their average excess over the first mean, and their average excess over the second mean. Don't round.	**A:** Each equipment status's average pulse rate is taken over all its signals, each signal weighted equally. **B:** Each equipment status's average pulse rate is the mean of per-telescope average pulse rates, each telescope weighted equally.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Average pulse rate of telescopes suggests equal telescope weighting, while based on signals permits pooling signals. The within-status weighting is unspecified.	SSBS
55907567ca alien	For each observatory (observstation in the telescopes table), work out its share of all telescopes. Taking just the observatories whose share is above the average share across all observatories, give a single summary row with: their combined number of telescopes, their average share, and the average amount by which their share exceeds the average share. Don't round.	**A:** Each observatory's share of all telescopes is a fraction: count ÷ total. **B:** Each observatory's share of all telescopes is a percentage: 100 × count ÷ total.	A: 1 row, B: 1 row; first row differs; largest value change: 0.002173 (A) vs 0.2173 (B), 99%	Share can be expressed as a fraction or a percentage, and no scale is specified. This is an output-unit choice, not a different counting unit; label the chosen scale.	SNS-
6c21dc4586 alien	Match each telescope to its observatory on observstation, and drop telescopes without a match. For each moon phase (observatories.lunarstage; a missing phase is its own group), show four things. First, the phase. Second, the average bandwidth usage (telescopes.bandusagepct), averaging per station first and then across stations (missing values ignored). Third, how many there are with that phase. Fourth, what share of the overall total that is. Don't round.	**A:** For each moon phase, count the telescopes, and give each phase's share of the total number of telescopes. **B:** For each moon phase, count the observatory stations, and give each phase's share of the total number of stations.	A: 4 rows, B: 4 rows; first row differs; largest value change: 236 (A) vs 226 (B), 4.2%	There are with that phase can refer to telescopes or stations. Explicit station-first averaging makes B more plausible here than in a purely telescope-level request.	-BBB
9fe5c73bb8 alien	For each observatory (observstation in the telescopes table), count its telescopes and work out its share of all telescopes. Then, taking only the observatories whose share is above the average share across all observatories, give a single summary row: their average telescope count, their average share, and how far on average their share sits above the average share. Don't round anything.	**A:** An observatory's share of telescopes is a fraction from 0 to 1. **B:** An observatory's share of telescopes is a percentage from 0 to 100.	A: 1 row, B: 1 row; first row differs; largest value change: 0.002173 (A) vs 0.2173 (B), 99%	Share can be expressed as a fraction or a percentage, and no scale is specified. This is an output-unit choice, not a different counting unit; label the chosen scale.	SSS-

Set	Open request	Readings	Results	Review	Models
19b51206f0 archeology	Look at the operators (scans.crewref) whose average file size per minute, computed per scan as gbsize ÷ spanmin with missing values ignored, is above the simple average of that figure across all operators. In a single row, how many scans have these operators done on average, what is their average GB per minute, and by how much on average do they exceed the all-operator figure? Don't round.	**A:** An operator's number of scans is the total of the scancount field over their records. **B:** An operator's number of scans is how many scan records they have.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Number of scans may mean summed scancount or the number of scan records. The request defines the rate calculation but not how to count scans.	-NBS
20f81515dc archeology	Using mesh records (scanmesh) that match a site (sites.zoneregistry = scanmesh.zoneref), work out each site's average ratio of mesh resolution (facetresmm) to geometric accuracy (geomdeltamm), ignoring missing ratios, plus its mesh count, and note the site's insurance status (sites.insurstat). Within each insurance status, take the simple average of its sites' ratios (each site counting once), and pick out the sites strictly above their status's average; sites with no recorded insurance status are left out. For each insurance status with such sites, show the status, the above-average sites' average ratio, their count, and how far on average their ratios exceed the status average. No rounding or ordering is needed.	**A:** For each insurance status, the count column is the total number of mesh records belonging to the sites that are above that status's average. **B:** For each insurance status, the count column is the number of sites that are above that status's average.	A: 3 rows, B: 3 rows; first row differs; largest value change: 96 (A) vs 89 (B), 7.3%	The intermediate mesh count points to A, while the final above-average sites and their count points to B. The requested final count has an ambiguous antecedent.	SSSS
60a79fabd7 archeology	Matching scans to their sites (scans.zoneref = sites.zoneregistry, skipping unmatched scans) and grouping by the sites' health and safety status (sites.healtheval), take the statuses whose average file size per minute (gbsize ÷ spanmin, ignoring missing values) is above the average across all statuses (none if that average is missing). In one unrounded row, what is their average rate, what is their average count, and on average how far does their rate exceed the all-status average?	**A:** Each status's count is the number of scan records linked to sites with that status. **B:** Each status's count is the number of distinct sites with that status that have scans.	A: 1 row, B: 1 row; first row differs; largest value change: 325.5 (A) vs 306 (B), 6%	Count has no explicit noun: it could count matched scan records or sites represented by them. The opening scan-level framing favors A.	SSNS
8f823fb0da archeology	Using the QC records matched to their projects (scanqc.arcref = projects.arcregistry, dropping any without a match), how many does each funding source have in QC? For each funding source (fundflux, with a missing one as its own group), show that count, the count as a fraction of the total over all funding sources, and the count divided by the average count per funding source, unrounded.	**A:** For each funding source, count the QC records linked to its projects. Report the share of the overall total and the ratio to the average per funding source. **B:** For each funding source, count the distinct projects that have at least one QC record. Report the share of the overall total and the ratio to the average per funding source.	A: 4 rows, B: 4 rows; first row differs; largest value change: 268 (A) vs 250 (B), 6.7%	How many in QC can mean QC records or distinct projects undergoing QC. The join alone does not choose the counting unit.	-S--
afa4a85bda archeology	Using spatial data matched to projects (ignore any whose project isn't in projects), how much spatial survey coverage does each funding source (fundflux) have? Give each source's count and its share of the total across all sources. Only include funding sources whose count is above the average across sources. Show the source, the count, the share, and how far the count is above that average. No rounding.	**A:** Count the spatial records under each funding source, then compute shares and keep the sources whose count is above average. **B:** Count the distinct projects with spatial records under each funding source, then compute shares and keep the sources whose count is above average.	A: 2 rows, B: 2 rows; first row differs; largest value change: 3 (A) vs 6.25 (B), 52%	Survey coverage can mean spatial-record volume or distinct projects surveyed. No explicit counting entity resolves the choice.	-B--
00d645c60b baseball_1	How do manager award winners break down by batting hand (bats)? Use only awards that match a record in player, and treat a blank or missing hand as its own group. For each hand, show the hand, how many manager award winners it has, and that figure's share of the overall total (no rounding).	**A:** For each batting hand, count the manager award records won by people with that hand, and give each hand's share of all matched award records. **B:** For each batting hand, count the distinct managers with that hand who won at least one award, and give each hand's share of all such managers.	A: 3 rows, B: 3 rows; first row differs; largest value change: 122 (A) vs 55 (B), 55%	Award winners can mean award-winning occurrences or distinct people with at least one award. B is the more literal people count; repeated wins are not explicitly addressed.	BBBB
1490d5931d baseball_1	Break fielding down by batting hand: match fielding to players on player_id and give the fielding count for each batting hand (player.bats), with blank or missing hands as their own groups. Give each hand's share of the total. Keep only the hands whose count is above the average across hands. For each, show the batting hand, the count, the share, and how far the count is above that average, unrounded.	**A:** Count every fielding row (one per player, year, stint and position) for each batting hand. **B:** Count the distinct players with fielding records for each batting hand.	A: 2 rows, B: 2 rows; first row differs; largest value change: 5356 (A) vs 167.8 (B), 97%	Fielding count could measure fielding records or players represented in fielding. A is favored by the data-level wording, but no counting unit is explicitly named.	SSBN
2f5c7dea6e baseball_1	Tally manager award winners by throwing hand (throws, with missing as its own group; awards matched to players on player_id, unmatched ones dropped) and work out each hand's share of the overall tally. For the hands whose share is above the average share, give one row with their combined tally, their average share, and how far above the average share they sit on average. Unrounded.	**A:** Each throwing-hand group is counted by its manager_award rows. **B:** Each throwing-hand group is counted by its distinct managers who won awards.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Award winners can mean award-winning occurrences or distinct people with at least one award. B is the more literal people count; repeated wins are not explicitly addressed.	-SB-
47e074c369 baseball_1	Grouping by the state players died in (death_state, with missing as its own group; Hall of Fame records matched to players on player_id), how much Hall of Fame representation does each state have? For the states above the average across states, give one row with their average count and how far above that average they are on average. Unrounded.	**A:** Each death state's count is its number of Hall of Fame ballot records. **B:** Each death state's count is its number of distinct players appearing in the Hall of Fame records.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Hall of Fame representation can mean ballot appearances or distinct represented players. Neither reading implies induction; the request leaves the counting unit open.	NBNN
6782deeb9c baseball_1	How do player award winners break down by batting hand (bats), counting awards matched to a known player and treating a blank or missing hand as its own group? For each hand, give the tally and its share of the overall total. Keep only the hands whose share is above the simple average share across hands, and show the hand, the tally, the share, and how much the share exceeds that average (no rounding).	**A:** For each batting hand, count the award records won by players with that hand. Compute each hand's share of all matched award records. Keep the hands whose share is above the average share across hands, and report hand, count, share and the excess over the average. **B:** For each batting hand, count the distinct players with that hand who won at least one award. Compute each hand's share of all such players. Keep the hands whose share is above the average share across hands, and report hand, count, share and the excess over the average.	A: 2 rows, B: 2 rows; first row differs; largest value change: 1924 (A) vs 376 (B), 80%	Award winners suggests distinct people, whereas counting awards suggests award records. The wording supports both counting units.	SSS-

Set	Open request	Readings	Results	Review	Models
696a17f661 baseball_1	Using player_college, work out each college's player count (by college_id, with records lacking a college_id treated as one group). Then find the colleges whose player count is strictly above the average across all colleges. Return a single row with: (1) those colleges' average player count; (2) the average of each one's player count as a fraction of the total of all colleges' player counts; and (3) the average amount by which they exceed the all-college average. Do not round.	**A:** Each college's size is the number of player_college records it has (player-years). **B:** Each college's size is the number of distinct players it has in player_college.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Player count favors distinct players, while counting player_college records would count repeated player-years. B is stronger, but record count is a plausible data-summary reading.	-BBN
7d77e789e0 baseball_1	Using player_award matched to player (dropping awards with no matching player), tally the award winners by throwing hand (player.throws). Treat a blank or missing hand as its own group. Work out each hand's share of the total across all hands. Keep only the hands whose share is above the average share. For those hands, give a single row with three values: the total of their tallies, their average share, and the average amount by which their share exceeds the average share. Don't round anything.	**A:** For each throwing hand, count award records. Then keep the hands whose share of the total is above the average share, and summarise them. **B:** For each throwing hand, count distinct award-winning players. Then keep the hands whose share of the total is above the average share, and summarise them.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Award winners can mean award-winning occurrences or distinct people with at least one award. B is the more literal people count; repeated wins are not explicitly addressed.	BBBS
a83a9b0bf3 baseball_1	Using the college attendance data matched to players by player_id (leaving out anything with no matching player), tally college attendance for each player birth_country, along with each country's share of the overall tally. For the countries above the average tally across all countries, give me one row: their combined tally, their average share, and on average how far above the average tally they are. No rounding.	**A:** Per birth country, count the player_college records matched to a player. Shares, the above-average filter and the summary are all built on those record counts. **B:** Per birth country, count the distinct players who have college records. Shares, the above-average filter and the summary are all built on those player counts.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	College attendance can mean attendance records or distinct players who attended college. Repeat attendance is not explicitly addressed.	-SB-
f447a6c8c2 baseball_1	For each college in player_college that matches a college record (unmatched rows dropped), find how many players it has, along with its state. Then for each state, with a blank or missing state as its own group, report four values: the state; the average of its colleges' player counts; the average of each college's count relative to the average count across all colleges; and the average of each college's share of the total across all colleges. Don't round anything.	**A:** Each college's count is its number of player_college records (player-years). Those counts are then summarised by state. **B:** Each college's count is its number of distinct players. Those counts are then summarised by state.	A: 49 rows, B: 49 rows; first row differs; largest value change: 15 (A) vs 5 (B), 67%	Player count favors distinct players, while counting player_college records would count repeated player-years. B is stronger, but record count is a plausible data-summary reading.	BSS-
3cc388876c crypto	Break down executed orders by channel (orders.orderbase, with a missing channel as its own group), using only executions that match an order on ordersmark = recordvault: for each channel, show the average filled quantity (fillcount, ignoring missing values), the count, and the channel's share of the total, unrounded.	**A:** Per channel, count execution rows and give each channel's share of all execution rows. **B:** Per channel, count distinct executed orders and give each channel's share of the total of those order counts.	A: 4 rows, B: 4 rows; first row differs; largest value change: 247 (A) vs 246 (B), 0.4%	Executed orders favors distinct orders, while the execution-level average and join favor counting execution rows. Count is not explicitly tied to one unit.	SB--
65eee834c9 crypto	Matching risk/margin records to their orders (ordervault = recordvault) and ignoring unmatched ones, count the risk/margin-matched orders on each order side (dealedge). For the sides whose count is above the average across sides, give one row with the total of their counts and the average amount by which they exceed that average, without rounding.	**A:** Count the joined risk/margin records for each order side. For the sides above the average count, return the total of their counts and their average excess over that average. **B:** Count the distinct orders with risk/margin records for each order side. For the sides above the average count, return the total of their counts and their average excess over that average.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Matching records suggests counting joined rows, while matched orders suggests distinct orders. B is the stronger entity-level reading.	SBB-
00d65781cc customers_and_invoices	How much of each product color was invoiced? Products with no color count as their own group. Work out each color's share of the total invoiced. Look at the colors whose invoiced amount is above the average per color. For those colors, give one row with their average invoiced amount, their average share, and how far above the per-color average they are on average. Don't round anything.	**A:** A color's invoiced amount is the number of invoice line items for products of that color. **B:** A color's invoiced amount is the total product_quantity on invoice lines for products of that color.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	How much was invoiced does not specify line frequency versus total units. B better matches quantity wording, but A is a plausible invoice-line volume measure.	-NN-
a63897eb14 customers_and_invoices	Size up each customer state/province's (state_county_province) account base, letting customers with no state form their own group. Work out each state's share of the total and the plain average of those shares across states. For the states whose share is strictly above that average, give me one unrounded row: their combined account base, their average share, and how far, on average, their share sits above the average share.	**A:** Each state's size is the number of accounts its customers hold. **B:** Each state's size is the number of different customers there who hold accounts.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Account base normally means accounts, but can also describe the customer base holding them. A is favored; clarify whether multiple accounts per customer count separately.	----
f27aed7e53 formula_1	Using driver standings linked to their races by race id (leaving out standings with no matching race), for each race name (races.name) get the average season-to-date championship points (driverStandings.points, ignoring missing values) and how big its standings field is. Then focus on the race names whose field size is strictly above the average field size across all race names. In one row, give the average of their average points, their average field size, and how far above the overall average their field size is on average. No rounding.	**A:** Per race name, count every driver standings record. Compare those counts with their average across names and summarise the names that are above average. **B:** Per race name, count distinct drivers in the standings. Compare those counts with their average across names and summarise the names that are above average.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Standings field size can mean accumulated driver entries or distinct drivers represented across races sharing a name. The time and deduplication scope is unspecified.	SBSS
3a20179886 insider	Using the advanced behaviour data (only where the trade and its trader can be found), work out for each trader type how many analysed trade records it has and its share of the overall total. Then, for the trader types whose number is strictly above the average across trader types, give in one row their combined number, their average share, and on average how far above that average number they sit. Don't round.	**A:** Count analysis records per trader type and compute each type's share of the total. For the types above the average count, report the summed count, the average share and the average excess over the average count. **B:** Count distinct analysed transactions per trader type and compute each type's share of the total. For the types above the average count, report the summed count, the average share and the average excess over the average count.	A: 1 row, B: 1 row; first row differs; largest value change: 17 (A) vs 17.25 (B), 1.4%	Analysed trade records can mean analysis records or distinct transactions analysed. The advanced-behaviour-data framing favors A, but the entity counted remains ambiguous.	-S--

Set	Open request	Readings	Results	Review	Models
753e097c22 insider	Using the transaction records matched to their traders (trdref = tradereg, exact text match, dropping records with no matching trader), how big is each trader type (tradekind) in this data? Traders with no type form their own group. Which trader types have a count strictly above the plain average of the per-type counts? For each one, show the trader type, its count, that count divided by the total of all the per-type counts, and the count minus the average. Don't round anything, and return nothing if the average is missing.	**A:** Each trader type's size is the number of transaction records belonging to traders of that type. **B:** Each trader type's size is the number of different traders of that type who have transaction records.	A: 1 row, B: 1 row; first row differs; largest value change: 17.25 (A) vs 15.5 (B), 10%	How big is each trader type can refer to its transaction activity or the number of active traders. The request does not name the counted unit.	-B--
989c45ebd2 insider	Taking the advanced behaviour analyses linked to their transaction (translink = transreg) and on to that transaction's trader (trdref = tradereg), using exact text matches and dropping anything that doesn't match: how much behaviour analysis does each trader type (tradekind) have? Treat a missing type as its own group. Compute each type's share as its count divided by the total of the counts across all types. Which trader types have a share strictly above the plain average of the per-type shares? For each one, show the type, its count, its share, and its share minus the average share. Don't round anything, and return nothing if the average is missing.	**A:** Each type's count is the number of behaviour-analysis records tied to its traders. **B:** Each type's count is the number of different transactions with behaviour analysis tied to its traders.	A: 1 row, B: 1 row; first row differs; largest value change: 0.017 (A) vs 0.01727 (B), 1.5%	Amount of behaviour analysis can measure analysis records or coverage of distinct transactions. A is more direct, but the counting unit is not explicit.	----
e6f24d2438 insider	For each trader type, how many analysed trade records are there in the advanced behaviour data (only where the trade and its trader can be found), and what share of the overall total does each type make up? Keep the share unrounded.	**A:** For each trader type, count the advanced behaviour analysis records that join to a transaction and a trader. Report that count and its share of the total count over all trader types. **B:** For each trader type, count the distinct transactions that have at least one advanced behaviour analysis and a matching trader. Report that count and its share of the total count over all trader types.	A: 4 rows, B: 4 rows; first row differs; largest value change: 242 (A) vs 241 (B), 0.41%	Analysed trade records can mean analysis records or distinct analysed transactions. A is favored by records in the advanced behaviour data.	-S--
ee905af7d9 insider	Taking the advanced behaviour analyses that tie to a transaction (via translink → transreg), which order type mixes (ordertypedist, with a missing mix treated as its own group) make up an above-average share? For each mix, find how many there are and its share of the overall total. Compare that share with the plain average share across mixes, and list only the mixes above that average. Show the mix, its count, its share, and how far it sits above the average share, all unrounded.	**A:** Each mix's count is the number of behaviour-analysis records linked to transactions with that mix, and shares are computed from those counts. **B:** Each mix's count is the number of distinct linked transactions with that mix, and shares are computed from those counts.	A: 2 rows, B: 2 rows; first row differs; largest value change: 0.005667 (A) vs 0.005005 (B), 12%	How many there are may refer to the analyses or the transactions whose mix is used. A is favored by taking the analyses, but the counting unit remains implicit.	-S--
fc16192583 sakila_1	For each actor, show their actor ID, how many films they're credited in, that count as a fraction of the total, and how far that fraction is above or below the average fraction across actors. Return the top 10 by that difference, highest first, unrounded.	**A:** Each actor's film count is divided by the total number of casting credits in film_actor, which gives the actor's share of all credits. **B:** Each actor's film count is divided by the number of distinct films that have a cast in film_actor, which gives the fraction of those films the actor appears in.	A: 10 rows, B: 10 rows; first row differs; largest value change: 0.001591 (A) vs 0.008716 (B), 82%	The total could mean total actor-film credits or distinct films; both are meaningful denominators, though only A produces shares summing to one. Historical-column discrepancy: this row says no, but the prior completed 39-row review says yes; this judgment follows that review.	-N--
Missing values					
19bf5ec0c7 baseball_1	Matching appearances to players by player id (drop appearances with no matching player), how many appearance records are there for each batting hand? For each batting hand, show the count, its share of the total, and the count divided by the average count per batting hand. Leave the numbers unrounded; any order is fine.	**A:** Players whose batting hand is recorded as an empty string are kept as a separate batting-hand group. That group adds to the total and to the average group size. **B:** Players whose batting hand is recorded as an empty string are left out before grouping, so only non-blank batting hands form groups, the total and the average.	A: 4 rows, B: 3 rows; first row differs; 1 row keys in only one answer; largest value change: 2.523 (A) vs 1.929 (B), 24%	Dropping unmatched entities does not resolve an empty attribute on a matched entity. A separate unknown group and exclusion of unknown values are both plausible.	-NSN
568dff8074 baseball_1	Counting player_college records under the state of their college (matching on exactly equal college_id; records whose college isn't found don't count toward any state; a blank or missing state is its own group), how many records does each state have, and what share of the player-college records is that? Don't round.	**A:** Each state's matched count is divided by the sum of the matched counts across all states. **B:** Each state's matched count is divided by the total number of player_college records, whether or not they match a college.	A: 49 rows, B: 49 rows; first row differs; largest value change: 0.03662 (A) vs 0.0366 (B), 0.058%	Unmatched records are explicitly excluded from state counts, but the denominator says player-college records without restricting it to matches. A gives shares of classified records; B gives shares of all records.	SBSS
69f4c578b6 baseball_1	For each player in the postseason batting data (matched to player on player_id), take their average postseason doubles per record, ignoring missing values, along with their batting hand. Divide it by the overall average of these per-player averages to get a relative index. Then show how far each player's index sits above the average index for their batting hand. Give the top 10 players by that gap, largest first (missing last), with player id, batting hand, average doubles, index, and the gap, unrounded.	**A:** An empty-string batting hand is a real group value: those players stay in every step and are compared with the other empty-hand players. **B:** An empty batting hand means unknown: those players (and NULL-hand players) are removed before anything is computed.	A: 10 rows, B: 10 rows; first row differs; largest value change: 12.8 (A) vs 12.79 (B), 0.039%	Unknown sex or batting hand can be retained as a comparison group or omitted. The request does not explicitly choose a missing-key policy; A follows literal grouping more closely.	-SS-
f74c2bad9d baseball_1	For players who received award votes (matched to player by player_id), what is the average points_won for each death country (ignoring missing points)? Show each death country, its average, and how far it is from the simple average of those per-country averages, without rounding.	**A:** Average the points won per death country, with a blank death country kept as its own group. Compare each group's average with the plain average of all group averages, the blank group included. **B:** Average the points won per death country, leaving out players with no recorded death country. Compare each country's average with the plain average of the remaining countries' averages.	A: 8 rows, B: 7 rows; first row differs; 1 row keys in only one answer; largest value change: 0.2828 (A) vs 3.15 (B), 91%	The request does not specify whether a missing grouping key forms an unknown bucket or is excluded. Both are reasonable reporting policies, though literal grouping favors A.	NN-N

Set	Open request	Readings	Results	Review	Models
`0f4fb81d6a chinook_1`	Count invoices by billing state, and work out each billing state's share of the invoices (its count over the sum of all the counts) and the average invoice count per billing state. Keep the billing states with more invoices than that average. Give one row with: the total number of invoices across them, their average share, and on average how many invoices above the average count they are (count minus average). No rounding.	**A:** Invoices with no billing state are treated as one more group alongside the real states. **B:** Invoices with no billing state are dropped before grouping.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	No missing-state policy is given. Unknown billing states can form a reporting bucket or be excluded from a breakdown of known states.	`SBBN`
`70ea90e305 chinook_1`	Break employees down by their `ReportsTo` value. For each group, show the `ReportsTo` value, how many employees it has, that count as a share of the total across groups, and how far its share is above or below the average share. Don't round. Sort by that difference from highest to lowest and show the top 10.	**A:** The employees with no manager make up a group of their own that takes part in every calculation. **B:** The employees with no manager are left out, so only real managers form groups.	A: 4 rows, B: 3 rows; first row differs; 1 row keys in only one answer; largest value change: 0 (A) vs -0.04762 (B), 1e+02%	The request does not specify whether a missing grouping key forms an unknown bucket or is excluded. Both are reasonable reporting policies, though literal grouping favors A.	`BB-N`
`b61f4a52bd crypto`	For each order in `orderexecutions.ordersmark` (executions with no order reference grouped together), what's the average size (`fillcount`) of the fills it got? Keep the orders whose average beats the average across all orders, and note by how much. Then, among those, keep the ones that also beat that group's own average, and note that margin too. Give me one summary row with three values: the average per-order fill size of the final orders, their average margin over the first benchmark, and their average margin over the second. Leave everything unrounded.	**A:** Zero-quantity executions count as fills of 0 in each order's average fill quantity. **B:** Zero-quantity executions are not fills and are left out of each order's average fill quantity.	A: 1 row, B: 1 row; first row differs; largest value change: 4.024 (A) vs 4.021 (B), 0.06%	Fills it got can mean all execution quantities, including zero, or only actual positive fills. The semantic use of fills makes B plausible; clarify the zero-fill rule.	`SSSS`
`2553116e1d customers_and_ invoices`	Looking at customers' order counts, what share of all orders from customers of the same gender does each customer account for? Customers with no gender recorded are left out. Which customers have a share above the average share across customers? Show the customer id, gender, number of orders, share, and how far the share is above the average. Sort by that gap from largest to smallest, with missing values last, and show the top 10. Don't round anything.	**A:** Only customers who placed at least one order are counted, both in the shares and in the average share. **B:** Every customer is counted. Customers with no orders have 0 orders and a share of 0, and they are included in the average share.	A: 4 rows, B: 8 rows; first row differs; 4 row keys in only one answer; largest value change: 0.0404 (A) vs 0.08889 (B), 55%	The request excludes missing gender but does not say whether customers with zero orders participate in the average.	`-B--`
`69c4bd1cc2 customers_and_ invoices`	For each order, count its invoices and work out its share of all invoices (its count divided by the total). For the orders whose share is strictly above the average share, give one unrounded row with: their total number of invoices, their average share, and the average amount by which their share exceeds the average share.	**A:** Only orders that have at least one invoice are included. **B:** All orders in Orders are included, and orders with no invoice count as 0.	A: 1 row, B: 1 row; first row differs; largest value change: 0.07778 (A) vs 0.1111 (B), 30%	The population for the entity-level averages is not explicit: observed entities only or all eligible entities with zero records. Including zeros changes the baseline, so the population needs clarification.	`-BSN`
`bc39324d95 customers_and_ invoices`	For each product, get the average `product_quantity` it was ordered in, how many order items it has, and its share of all order items. Average those three figures over the products in each product type (by type description), then give a single row with the average of each across all product types. Ignore missing values and don't round.	**A:** Only products that have order items are included at the per-product level before averaging by type and then overall. **B:** All products with a category are included. Never-ordered products count 0 order items, have a missing average quantity and a share of 0.	A: 1 row, B: 1 row; first row differs; largest value change: 0.08889 (A) vs 0.06417 (B), 28%	The population for the entity-level averages is not explicit: observed entities only or all eligible entities with zero records. Including zeros changes the baseline, so the population needs clarification.	`BBS-`
`c7b0dd6493 customers_and_ invoices`	Using the financial transaction counts by account, show the 10 accounts whose count is highest relative to the average count per account (count ÷ average), highest first. For each, give the account id, the town/city (`town_city`) of the customer who holds it, its number of transactions, its share of its town's transactions, and that ratio to the average. Leave out accounts whose holder isn't on file or has no town, and don't round.	**A:** Only accounts that have at least one transaction are considered. **B:** All accounts are considered, and those without transactions count as 0.	A: 8 rows, B: 10 rows; first row differs; 2 row keys in only one answer; largest value change: 2.133 (A) vs 4 (B), 47%	The population for the entity-level averages is not explicit: observed entities only or all eligible entities with zero records. Including zeros changes the baseline, so the population needs clarification.	`-NN-`
`bbce981d7b financial`	For each account, work out its share of all orders: its order count divided by the total number of orders (`order`). For the accounts whose share is above the average share, give one row with the total number of orders they have, their average share, and on average how far above the average share they are. Do not round.	**A:** Shares and their average cover only accounts that placed at least one order. **B:** Shares and their average cover all accounts, with zero-order accounts given share 0.	A: 1 row, B: 1 row; first row differs; largest value change: 0.0001418 (A) vs 0.0001856 (B), 24%	The population for the entity-level averages is not explicit: observed entities only or all eligible entities with zero records. Including zeros changes the baseline, so the population needs clarification.	`BBS-`
`28ea354f9a formula_1`	How many constructor standings records does each constructor have, and how does that compare with the average per constructor? Show the 10 constructors furthest above the average. Give the constructor id, its nationality, its number of standings records, and the difference from the average, largest first.	**A:** Only constructors that appear in the standings are counted. The average count per constructor is computed over them. **B:** Every constructor is counted, and constructors with no standings records get a count of 0. The average count per constructor is computed over all constructors.	A: 10 rows, B: 10 rows; first row differs; largest value change: 238.7 (A) vs 257.8 (B), 7.4%	The population for the entity-level averages is not explicit: observed entities only or all eligible entities with zero records. Including zeros changes the baseline, so the population needs clarification.	`NBBN`
`af91f25689 formula_1`	Which 10 drivers have the most qualifying entries compared with the average driver? For each, show the driver ID, how many qualifying records they have, and that number divided by the average number of qualifying records per driver (unrounded), highest ratio first.	**A:** The average number of qualifying records per driver is computed only over drivers who appear in the qualifying table. **B:** The average number of qualifying records per driver is computed over every driver in the drivers table, with drivers who never qualified counting as zero.	A: 10 rows, B: 10 rows; first row differs; largest value change: 4.112 (A) vs 22.63 (B), 82%	The population for the entity-level averages is not explicit: observed entities only or all eligible entities with zero records. Including zeros changes the baseline, so the population needs clarification.	`SBBB`

Set	Open request	Readings	Results	Review	Models
5ae33e9bfe sakila_1	How many actors does each film have, and what is the average of those counts (ignoring missing values)? Keep the films with strictly more actors than that average and note how far above it each one is. Among those films, average their actor counts again, keep only the ones strictly above that, and note how far above it they are too. If either average is missing, keep no films. Give me one row: the total actor count of the films that are left, their average excess over the first average, and their average excess over the second average, ignoring missing values and without rounding.	**A:** Per-film actor counts and every average are taken only over films that have at least one row in film_actor. **B:** Per-film actor counts and every average are taken over all films in the film table, with films that have no actor rows counted as 0.	A: 1 row, B: 1 row; first row differs; largest value change: 3.559 (A) vs 3.575 (B), 0.46%	The population for the entity-level averages is not explicit: observed entities only or all eligible entities with zero records. Including zeros changes the baseline, so the population needs clarification.	BBBS
975469f62b sakila_1	Which 10 films have the most inventory copies compared with the average number of copies per film? For each, show the film id, its rating, its number of copies, how far that is from the average copies per film for its rating, and how far it is from the overall average copies per film (averaged over films, not over the rating averages; averages ignore missing values). Sort by the overall difference, largest first. Ties can go in any order, and nothing is rounded.	**A:** Average copies per film (per rating and overall) are computed only over films that have at least one copy in inventory. Films with no copies are excluded. **B:** Average copies per film (per rating and overall) are computed over every film in the catalogue. Films with no copies in inventory count as having 0 copies.	A: 10 rows, B: 10 rows; first row differs; largest value change: 2.951 (A) vs 3.237 (B), 8.8%	The population for the entity-level averages is not explicit: observed entities only or all eligible entities with zero records. Including zeros changes the baseline, so the population needs clarification.	SBBS
4eb2b4a71e student_club	Total up income amount by member (link_to_member) and work out each total's share of the sum of all the totals. For those whose share is strictly above the average share, give one row with: their combined amount, their average share, and the average amount by which their share is above the average share. No rounding.	**A:** Group income by linked member, with unlinked income as its own group, then compute shares, the average share, and summarise the above-average groups. **B:** Drop unlinked income, group the rest by linked member, then compute shares, the average share, and summarise the above-average groups.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The request does not specify whether a missing grouping key forms an unknown bucket or is excluded. Both are reasonable reporting policies, though literal grouping favors A.	SNB-
55be105b98 student_club	Group members by their major (link_to_major) and show each major id, how many members it has, and that count's share of the total. List the top 10 by share, highest first, with no rounding.	**A:** Members without a major form their own group, and they are included in the total used as the denominator. **B:** Members without a major are left out of both the groups and the total used as the denominator.	A: 10 rows, B: 10 rows; first row differs; 2 row keys in only one answer; largest value change: 0.09091 (A) vs 0.09375 (B), 3%	The request does not specify whether a missing grouping key forms an unknown bucket or is excluded. Both are reasonable reporting policies, though literal grouping favors A.	BBBN
7e75655320 student_club	Group members by their major (link_to_major). For each group, show the major id, how many members it has, that count's share of the total, and the count divided by the average group count. List the top 10 by that last ratio, highest first, with no rounding.	**A:** Members without a major are grouped together as their own group. That group counts toward both the total member count and the average group size. **B:** Members without a major are removed before grouping. Only real major ids form groups, and the total and the average use only those groups.	A: 10 rows, B: 10 rows; first row differs; 2 row keys in only one answer; largest value change: 0.09091 (A) vs 0.09375 (B), 3%	The request does not specify whether a missing grouping key forms an unknown bucket or is excluded. Both are reasonable reporting policies, though literal grouping favors A.	NBBN
be1e5ce229 student_club	For each location, what are the top 3 events by average budget overspend ratio, (spent − amount) / amount averaged over each event's budget lines? Rank events within each location from highest to lowest ratio, with events lacking a ratio last. Tied events share a rank and the next rank is skipped. Show the event id, location, average ratio and rank.	**A:** Events with a missing location are ranked among themselves as one location group and can appear in the result. **B:** Events with a missing location are left out of the ranking and the result.	A: 10 rows, B: 4 rows; first row differs; 6 row keys in only one answer; no shared value changes	The request does not specify whether a missing grouping key forms an unknown bucket or is excluded. Both are reasonable reporting policies, though literal grouping favors A.	SSSS
e7ceb527cb student_club	Group members by link_to_major and count the members in each group, then give each group its share of the total count. For the groups whose share is above the average share across groups, return one row with the total number of members in them, their average share, and on average how far their shares are above the average share. No rounding.	**A:** Members without a major are treated as one extra group in every step. **B:** Only members who have a major are considered.	A: 1 row, B: 1 row; first row differs; largest value change: 0.06667 (A) vs 0.06875 (B), 3%	The request does not specify whether a missing grouping key forms an unknown bucket or is excluded. Both are reasonable reporting policies, though literal grouping favors A.	BBSS
3e50c13e8f thrombosis_ prediction	In the Examination table, average the ANA titer (ANA) per ID, ignoring missing values. Which IDs have an average above the mean of all those per-ID averages? Show the ID, its average ANA, and how far above the mean it is. Return the top 10 by that gap, largest first (missing values last).	**A:** Examination records with no patient ID form their own group, which counts toward the overall mean and can be listed. **B:** Examination records with no patient ID are excluded before grouping and averaging.	A: 10 rows, B: 10 rows; first row differs; largest value change: 3734 (A) vs 3734 (B), 0.0082%	Per-ID grouping can retain a missing-ID bucket, but which IDs can reasonably imply known patient IDs only. No missing-ID rule is stated.	SSBS
40a0836cfa thrombosis_ prediction	Break the examination records down by patient ID. Which IDs have an above-average share of the examinations? Show the ID, its number of examinations, its share of the total, and how far above the average share it is, listing the top 10 by that margin from largest to smallest, unrounded.	**A:** Group examinations by patient ID, with the examinations lacking an ID forming one group that counts toward the total and the average share; list up to 10 above-average groups by excess share. **B:** Group only the examinations that have a patient ID; the total and the average share use only those; list up to 10 above-average groups by excess share.	A: 8 rows, B: 7 rows; first row differs; 1 row keys in only one answer; largest value change: 0.001172 (A) vs 0.001287 (B), 8.9%	The request does not specify whether a missing grouping key forms an unknown bucket or is excluded. Both are reasonable reporting policies, though literal grouping favors A.	SSNS
6514cf857d thrombosis_ prediction	Break down the lab test records that match a patient (by patient ID) by the patient's sex. Which sexes have more test records than the average per sex? For each, show the sex, its number of test records, its share of the total, and how far above the average it is, unrounded.	**A:** Group linked lab records by patient sex, with the empty-string sex value as its own group that counts toward the total and the average group count; report the groups above the average. **B:** Group linked lab records by patient sex after dropping records whose sex is an empty string; the total and the average group count use only the remaining groups; report the groups above the average.	A: 1 row, B: 1 row; first row differs; largest value change: 6730 (A) vs 4452 (B), 34%	Dropping unmatched entities does not resolve an empty attribute on a matched entity. A separate unknown group and exclusion of unknown values are both plausible.	-BS-

Set	Open request	Readings	Results	Review	Models
9c11ade5dd thrombosis_prediction	Matching lab tests to patients, average the CPK-to-GPT ratio for each SEX value, ignoring tests where the ratio is missing. Keep the sexes whose average is above the mean of those averages, then among them keep the ones above their own mean. In one unrounded row, give the average ratio of the sexes left, how far above the first mean they are on average, and how far above the second mean they are on average.	**A:** A blank (") SEX value is a sex group of its own. Its average ratio counts in the first mean and can survive the cutoffs. **B:** Records with a blank or missing SEX are excluded before grouping, so only the real sex values form groups.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Unknown sex or batting hand can be retained as a comparison group or omitted. The request does not explicitly choose a missing-key policy; A follows literal grouping more closely.	SSSS
aed8b6c38d thrombosis_prediction	Matching laboratory records to their patient on ID (lab records with no matching patient are dropped), work out each patient's (by ID and sex) average prothrombin time (PT) and average IGA − IGM difference per record, ignoring missing values. Also work out how each patient's IGA−IGM average compares to the mean across all patients, as their average divided by that mean. Then break this down by sex: for each sex, give the average PT, the average IGA−IGM difference, and the average ratio to the overall mean.	**A:** Blank sex is reported as its own group alongside the other sexes. **B:** Blank sex is treated as unknown and its group is not reported.	A: 3 rows, B: 2 rows; first row differs; 1 row keys in only one answer; no shared value changes	Unknown sex or batting hand can be retained as a comparison group or omitted. The request does not explicitly choose a missing-key policy; A follows literal grouping more closely.	SSSS
b3691a7c41 thrombosis_prediction	Count the Examination records for each ID and express each count as a share of the total. For the IDs whose share is above the average share, what are their average record count, their average share, and the average amount by which their share beats the average share? Give one unrounded row.	**A:** All examination records with no ID are lumped into one group that is treated like a patient. **B:** Examination records with no ID are excluded, so only real patient IDs form groups.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Counting records per ID can retain a missing-ID group, while IDs can imply only known identifiers. Missing-ID handling is unspecified.	-SNS
c535414107 thrombosis_prediction	Linking lab records to patients by ID (dropping unmatched lab records), first average each patient's (ID, SEX) ALB and their per-test ratio of UN to T-BIL. Then average those patient-level figures for each sex. Which sexes have an average ALB above the mean of the sex-level ALB averages? Show the sex, its average ALB, its average UN/T-BIL ratio, and how far its ALB average is above that mean. Averages ignore missing values (including ratios that can't be computed); don't round.	**A:** A blank sex is a group of its own and contributes to the mean of the sex-level ALB averages. **B:** A blank sex means unknown, so those patients are dropped before grouping and before the mean is computed.	A: 2 rows, B: 1 row; first row differs; 1 row keys in only one answer; largest value change: 0.2467 (A) vs 0.04688 (B), 81%	Dropping unmatched lab records does not resolve missing sex on matched patients. Retaining an unknown-sex group or excluding it are both plausible.	NSSS
de9929d7a2 thrombosis_prediction	Which ANA patterns (ANA Pattern) have an average aCL IgG (ignoring missing values) above the average of all the pattern averages? For each, show the pattern, its average aCL IgG, how far it is above that cross-pattern average, and how far it is from the average of these above-average patterns' averages.	**A:** Examinations with no recorded pattern are treated as their own pattern group. **B:** Examinations with no recorded pattern are excluded before grouping.	A: 5 rows, B: 5 rows; first row differs; 2 row keys in only one answer; largest value change: -1.643 (A) vs -0.1649 (B), 90%	The request does not specify whether a missing grouping key forms an unknown bucket or is excluded. Both are reasonable reporting policies, though literal grouping favors A.	SS-S
Negation, choice in a filter condition					
2e529bb193 customers_and_invoices	Ignoring transactions with no invoice number, which invoice numbers have transactions without the type 'Payment' (exact text)? Show each invoice number with the total number of transactions recorded for it. Sort by that count from highest to lowest and give the top 10.	**A:** Invoices with no Payment transaction at all. For each one, count its transactions and show the top 10 by count. **B:** Invoices with at least one transaction that is not a Payment. For each one, count all its transactions (Payments included) and show the top 10 by count.	A: 4 rows, B: 7 rows; first row differs; 3 row keys in only one answer; no shared value changes	Transactions without the type Payment can describe non-Payment transactions on an invoice or an invoice without any Payment transactions. B follows the transaction-level modifier more closely; specify whether any Payment disqualifies the invoice.	SSNS
Strict cutoff, choice in a filter condition					
4432585c37 alien	For telescopes with high bandwidth usage (bandusagepct), using 47.6 as the cutoff, give for each calibration status (calibrstatus) the average bandwidth usage and the number of telescopes.	**A:** A telescope with usage exactly at 47.6 is kept. **B:** Only usage strictly greater than 47.6 is kept.	A: 3 rows, B: 3 rows; first row differs; largest value change: 173 (A) vs 169 (B), 2.3%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	SBBN
57727b703a archeology	For the memory-heavy processing runs, using 60.8 GB of RAM (memusagegb) as the cutoff and counting only runs whose site code (zoneref) matches a site: first get, for each site code and access status (entrystat), the average RAM usage and the average GPU-to-CPU usage ratio (procgpu ÷ proccpu), ignoring missing values. Then, for each access status, give the access status, the average of those per-site RAM averages, and the average of those per-site GPU/CPU ratios, unrounded.	**A:** A processing record qualifies when its RAM usage is 60.8 GB or more, so 60.8 itself counts. **B:** A processing record qualifies only when its RAM usage is greater than 60.8 GB, so 60.8 itself does not count.	A: 3 rows, B: 3 rows; first row differs; largest value change: 96.1 (A) vs 96.57 (B), 0.49%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBB-
06934ae88f chinook_1	Using 3.98 as the cutoff on the invoice total (Invoice.Total) to keep only the higher-value invoices, what is their average total, as a single unrounded value? Ignore missing totals, and give a missing result if no invoices qualify.	**A:** Invoices whose total equals 3.98 are included in the average. **B:** Only invoices whose total is strictly greater than 3.98 are included in the average.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBS

Set	Open request	Readings	Results	Review	Models
76d3120235 chinook_1	Joining each track to its media type (dropping tracks with no matching media type) and using 10323804 bytes as the file-size threshold (Track.Bytes) to keep only the large files, for each media type name (MediaType.Name) give the media type name, the average list price per track (Track.UnitPrice), and the average of file size divided by track length (Track.Bytes ÷ Track.Milliseconds). Both averages should ignore missing values and be missing when there are none. Leave all values unrounded.	**A:** Tracks whose file size equals 10323804 bytes are kept along with larger ones. **B:** Only tracks whose file size is strictly greater than 10323804 bytes are kept.	A: 4 rows, B: 4 rows; first row differs; largest value change: 33.38 (A) vs 33.38 (B), 0.0054%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBB
8ff14b4a90 chinook_1	Considering only tracks that pass a file-size (Track.Bytes) cutoff of 10,323,804 bytes, work out for each media type (by name) the average bytes per millisecond (Bytes ÷ Milliseconds) of its tracks. What is the simple average of these per-media-type averages? Return just that single value, unrounded.	**A:** Tracks whose file size is exactly 10,323,804 bytes are included along with larger ones. **B:** Only tracks whose file size is strictly greater than 10,323,804 bytes are included.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Passes a cutoff can mean meeting a minimum qualifying threshold or exceeding the boundary. Both A and B are reasonable here; replace this phrase with at least or strictly greater than. This is a filter-boundary choice, not value representation.	BBSB
a42de66df9 chinook_1	Taking 10323804 bytes as the cutoff for a large track, what is the average UnitPrice of the large tracks and their average bytes per millisecond (Bytes divided by Milliseconds, as a decimal ratio)? Ignore missing values and don't round.	**A:** A track whose file size equals the threshold exactly is included. **B:** Only tracks whose file size is strictly above the threshold are included.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	NBB-
ecf2814a02 chinook_1	Considering only tracks that pass a file-size (Track.Bytes) cutoff of 10,323,804 bytes, work out for each media type (by name) the average bytes per millisecond (Bytes ÷ Milliseconds) of its tracks. Which media types have an average strictly higher than the simple average of all these per-media-type averages? Show the media type name, its average, and how far it is above that overall average (its average minus the overall average). Leave all values unrounded.	**A:** Tracks whose file size is exactly 10,323,804 bytes are included along with larger ones. **B:** Only tracks whose file size is strictly greater than 10,323,804 bytes are included.	A: 1 row, B: 1 row; first row differs; largest value change: 115.7 (A) vs 115.7 (B), 0.00039%	Passes a cutoff can mean meeting a minimum qualifying threshold or exceeding the boundary. Both A and B are reasonable here; replace this phrase with at least or strictly greater than. This is a filter-boundary choice, not value representation.	BBSN
f6a08cfdf0 chinook_1	For the large track files, using 10,323,804 bytes (Bytes) as the size cutoff, group the tracks by MediaTypeId. For each media type, work out the average bytes per millisecond (Bytes ÷ Milliseconds, ignoring missing values) and how many tracks there are. Then give me one summary row with (1) the average of those per-media-type bytes-per-millisecond averages and (2) the average of those per-media-type track counts, both unrounded.	**A:** Tracks whose file size is at least 10323804 bytes count, including any track of exactly that size. **B:** Only tracks whose file size is strictly greater than 10323804 bytes count; a track of exactly that size is excluded.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBB
f8b3cab9d0 chinook_1	Joining each track to its media type (dropping tracks with no matching media type) and using 10323804 bytes as the file-size threshold (Track.Bytes) to keep only the large files, for each media type name (MediaType.Name) give the media type name, the average list price per track (Track.UnitPrice), and the average of file size divided by track length (Track.Bytes ÷ Track.Milliseconds). Also give that last average divided by the plain average of it across all the media types. All averages should ignore missing values and be missing when there are none. Return no rows if the across-media-type average is missing. Leave all values unrounded.	**A:** Tracks whose file size equals 10323804 bytes are kept along with larger ones. **B:** Only tracks whose file size is strictly greater than 10323804 bytes are kept.	A: 4 rows, B: 4 rows; first row differs; largest value change: 33.38 (A) vs 33.38 (B), 0.0054%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBSS
2ba180393f crypto	Using only high-rate fee records, with 0.0009 on the fee rate (feerate) as the cutoff, match each fee to its order (fees.orderslink = orders.recordvault, exact text match; drop fees with no matching order). For each combination of orderslink and order side (dealedge), compute the average net fee rate (feerate minus rebrate, ignoring missing values) and the number of records. Then compute each side's average of those group averages. List the groups whose average net fee rate is strictly greater than their side's average. Show orderslink, side, the group's average net fee rate, its record count, and the group average minus the side average. Sort by that difference from highest to lowest (missing last), keep the top 10, and don't round anything.	**A:** Fee records with a fee rate of exactly 0.0009 are included along with those above it. **B:** Only fee records with a fee rate strictly above 0.0009 are included; exactly 0.0009 is excluded.	A: 10 rows, B: 10 rows; first row differs; largest value change: 0.0008818 (A) vs 0.0008242 (B), 6.5%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBB
4f1aa2827f crypto	For the high-rebate fee records, using 0.0005 on the maker rebate rate (rebrate) as the cutoff, what are the total rebate (the sum of rebtotal) and the average net fee rate (feerate minus rebrate)? Return one row with these two values in that order, unrounded, ignoring missing values (each value is empty if there is nothing to aggregate).	**A:** Fee records whose maker rebate rate equals 0.0005 exactly are included along with those above it. **B:** Only fee records whose maker rebate rate is strictly above 0.0005 are included; a rate of exactly 0.0005 is excluded.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBS

Set	Open request	Readings	Results	Review	Models
eb00b504f6 crypto	For fee records in the high fee-rate band, using 0.0009 as the feerate cutoff, what's the average ratio of maker rebate rate to fee rate (rebrate / feerate), ignoring missing values, and how many such records are there?	**A:** Fee records with a fee rate of exactly 0.0009 are included, along with those above it. **B:** Only fee records with a fee rate strictly greater than 0.0009 are included.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBN
f701cd1be1 crypto	Using only high-rate fee records, with 0.0009 on the fee rate (feerate) as the cutoff, match each fee to its order (fees.orderslink = orders.recordvault, exact text match; drop fees with no matching order). For each combination of orderslink and trading pair (mktnote), show orderslink, the trading pair, the average maker rebate rate (rebrate), and the total rebate (the sum of rebtotal), ignoring missing values. Sort by total rebate from highest to lowest (missing last), keep the top 10, and don't round.	**A:** Fee records with a fee rate of exactly 0.0009 are included along with those above it. **B:** Only fee records with a fee rate strictly above 0.0009 are included; exactly 0.0009 is excluded.	A: 10 rows, B: 10 rows; 2 row keys in only one answer; no shared value changes	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBN
252b8028e7 financial	Treating 2400 as the cutoff for a large transaction amount (trans.amount), which 10 accounts (trans.account_id) have the most large transactions? For each one, show the account ID, the average amount of its large transactions, and how many it has, sorted by that count from highest to lowest.	**A:** Only transactions with an amount of 2400 or more count, so an amount of exactly 2400 is included. **B:** Only transactions with an amount strictly greater than 2400 count, so an amount of exactly 2400 is excluded.	A: 10 rows, B: 10 rows; first row differs; 2 row keys in only one answer; largest value change: 368 (A) vs 366 (B), 0.54%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	NBBB
8a55588ad2 financial	Applying a cutoff of 2802.0 to the debited amount (order.amount), what is the average debited amount of the large standing orders, unrounded?	**A:** Orders of exactly 2802.0 are included. **B:** Only orders above 2802.0 are included.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBB
9149a88808 financial	Using an amount of 2802 as the cutoff for a large standing order (the order table), compute the average amount of the large orders for each payment purpose code (k_symbol), treating the blank purpose code as a group of its own. Then take the simple average of those per-purpose averages. Which purposes have an average above that overall figure? For each, list the k_symbol, its average amount, and how far it exceeds the overall figure (purpose average minus overall figure), without rounding.	**A:** Standing orders with an amount of 2802 or more are used in the per-purpose averages and in the overall average. **B:** Only standing orders with an amount strictly greater than 2802 are used in the per-purpose averages and in the overall average.	A: 3 rows, B: 3 rows; first row differs; largest value change: 38.03 (A) vs 29.64 (B), 22%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBS
a65deca8fc financial	Treating 2400 as the cutoff for a large transaction amount (trans.amount), break the large transactions down by operation (trans.operation), with transactions that have no operation as their own group. For each group, show the operation; the average signed amount, positive for credits (trans.type = 'PRIJEM') and negative for withdrawals ('VYDAJ'), with any other type left out of the average; and how many large transactions the group has in total, of any type.	**A:** Only transactions with an amount of 2400 or more count, so an amount of exactly 2400 is included. **B:** Only transactions with an amount strictly greater than 2400 count, so an amount of exactly 2400 is excluded.	A: 5 rows, B: 5 rows; first row differs; largest value change: 3632 (A) vs 3298 (B), 9.2%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBB
c4495596b4 financial	Treating an amount of 2802 as the cutoff for a large standing order (the order table), link each large order to its account and the account's district (dropping orders with no matching account or district), and compute the average order amount for each combination of account (order.account_id) and region (district.A3). Then take the simple average of all those group averages. Show the 10 groups whose average exceeds that overall figure by the most: give the account id, the region, the group's average amount, and the difference (group average minus overall figure), sorted by that difference from largest to smallest, without rounding.	**A:** Standing orders with an amount of 2802 or more are used in the per-account averages and in the overall average. **B:** Only standing orders with an amount strictly greater than 2802 are used in the per-account averages and in the overall average.	A: 10 rows, B: 10 rows; first row differs; largest value change: 8655 (A) vs 8651 (B), 0.052%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBB
c7d3136268 financial	Looking only at large transactions, with the cutoff at an amount (trans.amount) of 2400, work out two figures for each account (trans.account_id). The first is its average signed transaction amount, where credits ('PRIJEM') count as +amount, withdrawals ('VYDAJ') count as −amount, and other or missing types are left out. The second is its average balance after transaction (trans.balance). Then return a single row giving the average across accounts of each of these two figures, in that order, without rounding; accounts with no signed amount are skipped for the first figure.	**A:** Transactions of exactly 2400 are kept. **B:** Only transactions above 2400 are kept.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBN

Set	Open request	Readings	Results	Review	Models
cc6a148eaf financial	Looking only at large transactions, with a 2,400 cutoff on the amount, which 10 accounts have the highest average balance after those transactions? For each one, show the account_id, its statement frequency, the average signed amount (credits 'PRIJEM' positive, withdrawals 'VYDAJ' negative, other types ignored) and that average balance, highest first, unrounded.	**A:** Use only transactions with an amount of 2400 or more. For each account (with its frequency), give the average signed amount and the average balance, and return the 10 accounts with the highest average balance. **B:** Use only transactions with an amount strictly greater than 2400. For each account (with its frequency), give the average signed amount and the average balance, and return the 10 accounts with the highest average balance.	A: 10 rows, B: 10 rows; first row differs; 2 row keys in only one answer; largest value change: 1213 (A) vs 1300 (B), 6.7%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	-BB-
2b9a831682 formula_1	For race results where the driver started from a grid position upwards of 12: add up each driver's points and average their laps completed (laps), ignoring missing values. Across those drivers, what are the average points total and the average of their average laps? Give one unrounded row.	**A:** Use results where the driver started at grid position 12 or higher-numbered. Average the per-driver point totals and the per-driver average laps. **B:** Use results where the driver started at a grid position numbered higher than 12. Average the per-driver point totals and the per-driver average laps.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Upwards of can mean the stated value and higher or more than the stated value. Both inclusive A and strict B are reasonable; state the intended comparison explicitly. This is a filter-boundary choice, not value representation.	BBBN
544772024a formula_1	Looking only at back-of-the-field results, using 12 as the cutoff on classification order (positionOrder), and matching each result to its race: for each race (grouped by race ID and race name), show the race ID, the race name, and the average number of laps completed (laps, ignoring missing values). List the 10 races with the highest average first, with any race lacking an average placed last.	**A:** Results classified exactly 12th are included in the filtered set, along with all later classifications. **B:** Only results classified after 12th (13th onward) are included. Results classified exactly 12th are excluded.	A: 10 rows, B: 10 rows; first row differs; 2 row keys in only one answer; largest value change: 106.3 (A) vs 102 (B), 4%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBS
57277c47b4 formula_1	Among driver standings entries where the driver has upwards of 1 point so far that season, what is the average season-to-date win count? Ignore missing values and don't round.	**A:** Average the season-to-date win counts over standings entries whose season-to-date points are 1 or more. **B:** Average the season-to-date win counts over standings entries whose season-to-date points are more than 1.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Upwards of can mean the stated value and higher or more than the stated value. Both inclusive A and strict B are reasonable; state the intended comparison explicitly. This is a filter-boundary choice, not value representation.	BBBS
3bf6ac003c insider	Using 0.439 as the cutoff for similarity to illicit patterns (patsim), take the advanced behaviour records with high similarity and match each one to its transaction record(s), where transactionrecord.transreg equals advancedbehavior.translink; drop records with no match. For each order type mix (ordertypedist), treating a missing mix as its own group, show the order type mix and the number of matched records.	**A:** Pattern-similarity values equal to 0.439 pass the filter (at least 0.439). **B:** Only pattern-similarity values strictly greater than 0.439 pass the filter (more than 0.439).	A: 3 rows, B: 3 rows; first row differs; largest value change: 184 (A) vs 180 (B), 2.2%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBS
8ed79619ea insider	Using 0.149 as the cutoff, how many advanced behaviour records have high peer correlation (peercorr)? Give a single count.	**A:** Peer correlation values equal to 0.149 are counted (at least 0.149). **B:** Only peer correlation values strictly greater than 0.149 are counted (more than 0.149).	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BSBN
f892edf28a insider	Using 0.149 as the peer correlation (peercorr) cutoff, take the advanced behaviour records with high peer correlation and match each one to its transaction record(s), where transactionrecord.transreg equals advancedbehavior.translink; drop records with no match. For each combination of translink and order type mix (ordertypedist), count the matched records. For each combination, show the translink, the order type mix, the count, and the count as an unrounded fraction of the total count across all combinations with that same order type mix. Leave out combinations where the order type mix is missing. Sort by that fraction from highest to lowest and return the top 10.	**A:** Peer correlation values equal to 0.149 pass the filter (at least 0.149). **B:** Only peer correlation values strictly greater than 0.149 pass the filter (more than 0.149).	A: 10 rows, B: 10 rows; first row differs; largest value change: 0.007692 (A) vs 0.007812 (B), 1.5%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	-B--
0170ade8cb sakila_1	Treating 5 days as the cutoff for a long rental period, look only at the long-rental films (by rental_duration) and group them by rating. For each rating, show the rating, the average per film of rental_rate divided by rental_duration (ignoring missing values), the number of films, and that average divided by an overall benchmark. The benchmark is the plain unweighted mean of the per-rating averages across all ratings, ignoring any missing ones. If the benchmark itself is missing, return no rows. Leave all values unrounded.	**A:** Keep films whose rental duration is 5 days or more, so films with exactly 5 days count. **B:** Keep only films whose rental duration is strictly more than 5 days, so films with exactly 5 days are left out.	A: 5 rows, B: 5 rows; first row differs; largest value change: 100 (A) vs 60 (B), 40%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBB

Set	Open request	Readings	Results	Review	Models
15a3b1f80a sakila_1	Looking at the costlier films, with the `replacement_cost` threshold set at 20.99, what are the average running time (`length`) and the average ratio of `rental_rate` to `replacement_cost` for each `rating` (ignoring missing values)? Then summarize them as one row: the average of the per-rating lengths and the average of the per-rating ratios, in that order and unrounded.	**A:** Only films whose replacement cost is 20.99 or higher are used. Films costing exactly 20.99 are included. **B:** Only films whose replacement cost is strictly higher than 20.99 are used. Films costing exactly 20.99 are excluded.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBB
29a86460a7 sakila_1	Using 5 days (`film.rental_duration`) as the cutoff for a long allowed rental period and considering only films with a long one, group them by original language (`film.original_language_id`), treating films with no original language as one group of their own, and for each group work out the average rental rate (`rental_rate`) and the number of films. Then, across those groups, what is the average of the group average rental rates and the average of the group film counts? Treat each group equally, ignore missing values (missing if none remain), return the two figures in that order as a single row, and don't round.	**A:** Films with an allowed rental duration of 5 days or more are included; exactly 5 days counts. **B:** Only films with an allowed rental duration of strictly more than 5 days are included; exactly 5 days does not count.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBS
8c141562df sakila_1	Using 20.99 as the replacement-cost cutoff for films, group the qualifying films by `original_language_id`, with films that have no original language forming their own group. For each group, show the original language id, the average of `rental_rate` divided by `rental_duration` across its films (unrounded, ignoring missing values), and the number of films. Sort the groups from most films to fewest (any missing counts last) and show only the top 10.	**A:** Films with a replacement cost equal to 20.99 are included along with those above it. **B:** Only films with a replacement cost above 20.99 are included. A cost of exactly 20.99 is left out.	A: 1 row, B: 1 row; first row differs; largest value change: 486 (A) vs 429 (B), 12%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BNBN
9276ea1d18 sakila_1	Using 112 minutes (`film.length`) as the cutoff for a long film and considering only the long films, for each language (`film.language_id`) give the language id, the average replacement cost (`replacement_cost`), and the average running time, ignoring missing values in each average (missing if none remain). Sort by average running time from highest to lowest, with missing averages last, show the top 10, and don't round.	**A:** Films whose running time is 112 minutes or more are included; a film of exactly 112 minutes counts. **B:** Only films whose running time is strictly more than 112 minutes are included; a film of exactly 112 minutes does not count.	A: 1 row, B: 1 row; first row differs; largest value change: 148.2 (A) vs 149.1 (B), 0.61%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBB
bf7e064707 sakila_1	Applying a 4.99 cutoff on the payment amount, take the qualifying payments that can be linked to a rental and, through that rental, to an inventory item (leave out payments without such a match). For each rental and store (`inventory.store_id`), total the payment amounts and count the payments. Then for each store, show the store id, the average of those per-rental totals, and the average number of payments per rental (both unrounded, ignoring missing values).	**A:** Payments of exactly 4.99 are included along with larger ones. **B:** Only payments larger than 4.99 are included. Payments of exactly 4.99 are left out.	A: 2 rows, B: 2 rows; first row differs; largest value change: 6.205 (A) vs 7.389 (B), 16%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBS
c61660aba3 sakila_1	Using 20.99 as the replacement-cost cutoff for films, compute for each rating the average of `rental_rate` divided by `rental_duration` over the qualifying films (ignoring missing values). Then take the simple mean of these per-rating averages, ignoring missing ones. If that overall mean exists, list each rating whose average is strictly greater than it, showing the rating, its average, and how far its average is above the overall mean (all unrounded).	**A:** Films with a replacement cost equal to 20.99 are included along with those above it. **B:** Only films with a replacement cost above 20.99 are included. A cost of exactly 20.99 is left out.	A: 3 rows, B: 3 rows; first row differs; largest value change: 6.279e-05 (A) vs 0.006898 (B), 99%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	NBBN
e3e52d6207 sakila_1	Taking 20.99 as the cutoff for a film being expensive to replace (`replacement_cost`), look only at the expensive films and group them by `rating`. For each rating, show the rating, the average `length` (ignoring missing values), and the average per film of `rental_rate` divided by `rental_duration` (ignoring missing values). Leave all values unrounded.	**A:** Keep films whose replacement cost is 20.99 or more, so films costing exactly 20.99 count. **B:** Keep only films whose replacement cost is strictly more than 20.99, so films costing exactly 20.99 are left out.	A: 5 rows, B: 5 rows; first row differs; largest value change: 110.2 (A) vs 113 (B), 2.5%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBB
f0c4cb72e1 sakila_1	For the pricier-to-replace films, with 20.99 as the `replacement_cost` cutoff, what are the average `rental_duration` and the average `rental_rate`-to-`replacement_cost` ratio for each `rating` (ignoring missing values)? Then give me one row with the average of those per-rating durations and the average of those per-rating ratios, in that order and unrounded.	**A:** Only films whose replacement cost is 20.99 or higher are used. Films costing exactly 20.99 are included. **B:** Only films whose replacement cost is strictly higher than 20.99 are used. Films costing exactly 20.99 are excluded.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBS

Set	Open request	Readings	Results	Review	Models
f49dd87a11 sakila_1	Using 112 minutes (film.length) as the cutoff for a long film, take all the long films together as one group. What is their average running time, and what is the average of each film's replacement cost divided by its rental rate (replacement_cost / rental_rate, computed per film, then averaged)? Ignore missing values in each average (giving a missing value if none remain), return the two figures in that order as a single row, and don't round.	**A:** Films whose running time is 112 minutes or more are included; a film of exactly 112 minutes counts. **B:** Only films whose running time is strictly more than 112 minutes are included; a film of exactly 112 minutes does not count.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBBB
53d9deb09a student_club	Looking only at budget lines whose budgeted amount passes the 75-dollar cutoff, group them by category. For each category, compute the average remaining and the total of amount minus spent, ignoring missing values (each is blank if a category has no values). Then average those per-category totals across all categories. If that average exists, list the categories whose total is strictly greater than it. For each one, show the category, its average remaining, its total, and how far its total is above the overall average, without rounding.	**A:** Budget lines whose budgeted amount equals 75 are included along with those above 75. **B:** Only budget lines whose budgeted amount is strictly greater than 75 are included. Lines at exactly 75 are left out.	A: 1 row, B: 1 row; first row differs; largest value change: 738 (A) vs 843 (B), 12%	Passes a cutoff can mean meeting a minimum qualifying threshold or exceeding the boundary. Both A and B are reasonable here; replace this phrase with at least or strictly greater than. This is a filter-boundary choice, not value representation.	BBSS
d46bfabe90 student_club	Linking each budget line to its event (event.event_id = budget.link_to_event, exact text match; lines without a matching event are dropped), and using 6.0 as the cutoff on spent so that only the higher-spending budget lines are considered, show for each event status: the status, the average of amount minus spent (ignoring missing values; blank if there are none), and the number of such budget lines. Leave the results unrounded.	**A:** Keep budget lines whose spent is 6.0 or more, including lines at exactly 6.0. **B:** Keep only budget lines whose spent is strictly greater than 6.0. A line at exactly 6.0 is excluded.	A: 2 rows, B: 2 rows; first row differs; largest value change: 22 (A) vs 19 (B), 14%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BBSS
0cdfa9d90b thrombosis_prediction	For lab records that match a patient by ID (unmatched records dropped), screened with 180 as the total cholesterol (T-CHO) cutoff, break them down by patient sex (SEX), keeping each recorded value, including a blank one, as its own group. For each group, show the sex, the average anti-DNA level (Laboratory.DNA), and the average of albumin minus total protein (ALB − TP), ignoring missing values and leaving the average empty if there are none, unrounded.	**A:** Lab records with total cholesterol equal to 180 count as passing the filter (T-CHO at least 180). **B:** Only lab records with total cholesterol strictly above 180 pass the filter; records with exactly 180 are excluded.	A: 3 rows, B: 3 rows; largest value change: 25.42 (A) vs 26.22 (B), 3.1%	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	SBSS
33f486d239 thrombosis_prediction	Using 3.9 as the cutoff for elevated uric acid (UA), take only the laboratory records with elevated uric acid. For each patient, compute the average hemoglobin (HGB) and the average of APTT minus PT over those records. What are the averages of these two per-patient values across all patients? Give them as a single unrounded row and ignore missing values at every step.	**A:** Laboratory records with a uric acid value equal to 3.9 or greater are included. **B:** Only laboratory records with a uric acid value strictly greater than 3.9 are included.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	BSBS
8bafc1faa8 thrombosis_prediction	Taking individual lab records and using 6.5 as the white blood cell count (WBC) cutoff, group them by anti-centromere status (CENTROMEA), treating a missing status as its own group, and for each status work out the average AST minus ALT (GOT − GPT) and the average APTT minus PT (APTT − PT), skipping records where a difference can't be computed. Then give me a single row with two unrounded figures: the average of those per-status AST − ALT averages and the average of the per-status APTT − PT averages, with each status counting equally and missing per-status values ignored.	**A:** Laboratory records with WBC equal to 6.5 count as passing the filter (WBC at least 6.5). **B:** Only Laboratory records with WBC strictly above 6.5 pass the filter; records with exactly 6.5 are excluded.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	NSSB
e5b4a71fd9 thrombosis_prediction	Using 6.5 as the cutoff for a high white blood cell count (WBC), what are the average triglyceride level (TG) and the average WBC across the laboratory records with a high count? Give them as a single unrounded row and ignore missing values.	**A:** Records with a white blood cell count equal to 6.5 or greater are included. **B:** Only records with a white blood cell count strictly greater than 6.5 are included.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	The cutoff or threshold defines the boundary but does not explicitly say whether equality qualifies. Both at least the cutoff (A) and strictly above it (B) are reasonable; specify >= or >. This is a filter-boundary choice, not value representation.	NSSS
Ties and selection					
1c7e18597b alien	For each observatory station, count its telescopes and work out its share of all telescopes (its count divided by the total number of telescopes). Which are the top 10 observatories whose share is above the average share across observatories, ranked by how far above average they are? Show the observstation, the telescope count, the share, and the share minus the average share, unrounded, from the biggest gap to the smallest.	**A:** List observatories whose telescope share is above the average share, ordered by how far above average they are. Return exactly 10 rows, breaking any tie at the cutoff arbitrarily. **B:** List observatories whose telescope share is above the average share, ordered by how far above average they are. Return the top 10 plus every observatory tied with the 10th on that difference.	A: 10 rows, B: 32 rows; 22 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	----

Set	Open request	Readings	Results	Review	Models
65e2e1cfdd alien	What are the top 10 observatory stations (observstation) by share of all telescopes? For each one, show the station, how many telescopes it has, and its share of the total telescope count, without rounding, ordered from the largest share down.	**A:** Return exactly ten stations by share and break any tie at the tenth position arbitrarily. **B:** Return the top ten stations by share plus every other station tied with the tenth station's share.	A: 10 rows, B: 32 rows; 22 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	-NS-
da10959f8c alien	For each observatory (observstation in the telescopes table), count its telescopes and work out its share of all telescopes as a fraction (its count divided by the total across all observatories). How far is each observatory's share above or below the average share? Show the observatory, its telescope count, its share and that difference for the top 10 observatories by the difference, highest first.	**A:** Return exactly 10 observatories. Ties at the 10th place are broken arbitrarily. **B:** Return the top 10 observatories plus every observatory tied with the 10th place on the difference.	A: 10 rows, B: 32 rows; 22 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	SNS-
ee6b268047 alien	What are the top 10 telescopes by share of detected signals? For each telescope reference (telescref), show how many signals it detected and what fraction of all signals that is (unrounded), from the largest share down.	**A:** Rank telescopes by their share of all signals and return exactly 10 rows, choosing arbitrarily among telescopes tied at the 10th place. **B:** Rank telescopes by their share of all signals and return every telescope whose share is at least the 10th-highest share, including all ties at the cut.	A: 10 rows, B: 51 rows; 41 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	SSS-
4cabfc6bd3 archeology	In the QC data, tally the projects under each archival status (archstat), with a missing status treated as its own group. Work out each status's share as its tally divided by the sum of the tallies over all statuses. Return only the statuses whose share is above the average share across statuses. Show the status, the tally, the share, and how far the share is above the average share, with no rounding.	**A:** For each archival status, count the QC records. Each share is that count divided by the sum of the counts. Return the statuses whose share is above the average share, with the status, count, share and share minus average. **B:** For each archival status, count the distinct projects with QC records of that status. Each share is that count divided by the sum of these counts. Return the statuses whose share is above the average share, with the status, count, share and share minus average.	A: 2 rows, B: 2 rows; first row differs; largest value change: 0.004667 (A) vs 0.003401 (B), 27%	Projects favors distinct project IDs, while tallying the QC data can mean QC records. B is stronger, but the record-versus-project unit needs clarification. This is not about ties.	NB--
6739210c1d archeology	Match each conservation record to its site (zoneref = site code) and group by the site's insurance status (insurstat), with a missing status treated as its own group. For each status, count the conserved sites. Show the status, the count, and its share of the sum of the counts over all statuses, with no rounding.	**A:** For each site insurance status, count the conservation records at matching sites. Each share is that count divided by the sum of the counts. Show the status, count and share. **B:** For each site insurance status, count the distinct sites with conservation records. Each share is that count divided by the sum of these counts. Show the status, count and share.	A: 3 rows, B: 3 rows; first row differs; largest value change: 319 (A) vs 299 (B), 6.3%	Conserved sites favors distinct sites; the record-matching framing also permits counting conservation records. Clarify repeated records per site. This is a counting-unit choice, not ties.	BB--
6c0edf283c archeology	Match the conservation records to their sites by site code, and drop any with no matching site. How much does each site health and safety status (healtheval) account for in the conservation data? Treat a missing status as its own group, and give a count and that count's share of the sum. Then take only the statuses with an above-average share, and return one summary row with the average count, the average share, and the average amount by which the share exceeds the average share. Don't round.	**A:** Each status's count is the number of conservation records. **B:** Each status's count is the number of distinct sites that have conservation records.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	A count in the conservation data could count conservation records or represented sites; no counting entity is specified. This is not a tie-handling choice.	SS--
78923d5b9e archeology	Matching conservation records (scanconservation) to their sites on zoneref = sites.zoneregistry and dropping any without a site, which are the top 10 site reference and insurance status (insurstat) combinations by share of conservation records? Treat a missing status as its own value. Show the site reference, the insurance status, the record count, and the count as a fraction of all matched conservation records (unrounded), sorted by that fraction from highest to lowest with missing values last.	**A:** List the 10 site/insurance-status groups with the largest share of conservation records. Return exactly 10, breaking ties at the cutoff arbitrarily. **B:** List the site/insurance-status groups with the largest share of conservation records. Return all rows whose share reaches the 10th-highest value, including every row tied at the cutoff.	A: 10 rows, B: 53 rows; 43 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified. For B, use the tenth ordered row as the cutoff, not the tenth distinct share.	SN--
9b7e954a5a archeology	For each project (arcref) in the QC data, count its QC records. Work out its share as that count divided by the total over all projects. Compute the deviation as the share minus the average share across projects. Show the top 10 projects by deviation, highest first (missing last), with the project, the count, the share and the deviation, and no rounding.	**A:** Rank projects by how far their share of QC records is above the average share, and return exactly 10 rows, cutting through any tie at the cutoff. **B:** Rank projects by how far their share of QC records is above the average share, and return every project whose deviation is at least the 10th-highest, including ties.	A: 10 rows, B: 51 rows; 41 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	-N--
e14588a994 archeology	Looking at spatial scan data (scanspatial) linked to projects (projects.arcregistry = scanspatial.arcref), which funding sources (projects.fundflux, treating a missing source as its own group) are bigger than average? For each funding source, work out the average ratio of bounding box Y (boxy) to bounding box Z (boxz), ignoring missing values, and how big it is. For the sources above the average size across funding sources, give one row with the average of their ratios, their combined size, and how far on average they exceed the average size. No rounding.	**A:** The size of each funding source is measured by how many spatial records it has. **B:** The size of each funding source is measured by how many distinct projects it has spatial data for.	A: 1 row, B: 1 row; first row differs; largest value change: 10.5 (A) vs 9.75 (B), 7.1%	Size of a funding source is undefined: spatial-record volume and distinct represented projects are both plausible. This is a counting-unit choice, not ties.	NBB-
ea536393b9 archeology	Matching each scan to its site by site code (dropping scans with no matching site), group the scans by site code (scans.zoneref) and risk evaluation status (riskeval). For each group give the site code, the risk evaluation status, the average compression ratio (pressratio, ignoring missing values), the number of scans, and that number as a share of all matched scans. Show the top 10 by share, largest first. Don't round.	**A:** Rank the site-and-risk-status groups by share of scans and return exactly 10. Ties at the cut are broken arbitrarily. **B:** Rank the site-and-risk-status groups by share of scans and return every group whose share reaches the 10th-ranked share, including all ties.	A: 10 rows, B: 53 rows; 43 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	SN--

Set	Open request	Readings	Results	Review	Models
`fea75e8f25 archeology`	Among equipment whose average relative humidity (`scanenvironment.humepct`, ignoring missing readings) is above the average across all equipment, which are the top 10 by how far they sit above the mean of that above-average group? Show the equipment reference, its average humidity, how much it exceeds the all-equipment average, and its difference from the group mean. Return nothing if either reference average is missing, and sort from highest to lowest with missing values last.	**A:** List the above-average-humidity equipment ranked by how far it sits above the mean of that group, and return exactly 10 rows, cutting arbitrarily if there are ties at 10th place. **B:** List the same ranking, but keep all equipment that ties with the 10th row, so more than 10 rows may come back.	A: 10 rows, B: 11 rows; 1 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	`-S--`
`0fb0f319ce chinook_1`	Work out each invoice's average line `UnitPrice` (with its customer's country) and how far that is from the average across all invoices. Which are the top 10 invoices whose deviation sits furthest above their country's average deviation? Show the invoice id, country, average unit price, deviation from the overall average, and the deviation minus the country's average deviation, sorted highest first (blanks last) and unrounded.	**A:** Rank invoices by how far their deviation from the overall average invoice unit price exceeds their country's average deviation, and return exactly 10 rows, cutting ties arbitrarily. **B:** Rank invoices the same way, and return every invoice ranked in the top 10, including all invoices tied at 10th place.	A: 10 rows, B: 11 rows; 1 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	`SSS-`
`12d93c67a7 crypto`	Join fees to their orders (`fees.orderslink` = `orders.recordvault`; drop fees with no matching order) and group by time in force (`orders.timespan`). For each group, find the average `rebrate` / `feerate` ratio (ignoring missing values) and how many order matches it has. Keep only the groups whose match count is strictly above the average match count across all groups. For those groups, give me one row: the average of their average ratios, the total of their match counts, and the average amount by which their match count exceeds that cross-group average. No rounding.	**A:** Each time-in-force group's size is the number of joined fee-order records in it. Groups are compared to the average of these sizes, and the sizes are summed. **B:** Each time-in-force group's size is the number of distinct orders in it. Groups are compared to the average of these sizes, and the sizes are summed.	A: 1 row, B: 1 row; first row differs; largest value change: 14 (A) vs 14.25 (B), 1.8%	Order matches can mean joined fee-order rows or distinct matched orders. No uniqueness rule is specified. This is a counting-unit choice, not ties.	`SSSS`
`4aa1861890 crypto`	Matching risk/margin records to their orders (`ordervault` = `recordvault`) and orders to their users (`userlink` = `userstamp`), and ignoring anything that doesn't match: for each account scope (`users.acctscope`), how many risk/margin-matched orders are there? Show the scope, that count, its share of the total across all scopes, and the count divided by the average count per scope, without rounding.	**A:** For each account scope, count the risk/margin records that link through orders to users, so an order with several risk/margin records counts once per record. Report the count, its share of the total across scopes, and its ratio to the average count across scopes. **B:** For each account scope, count the distinct orders that have risk/margin records and link to users, so each order counts once. Report the count, its share of the total across scopes, and its ratio to the average count across scopes.	A: 4 rows, B: 4 rows; first row differs; largest value change: 224 (A) vs 223 (B), 0.45%	Matched orders favors distinct orders, while the record-join framing supports counting matches. Specify whether an order with multiple risk records counts once. This is a counting-unit choice, not ties.	`BB--`
`7b6fd5bf32 crypto`	Matching risk/margin records to their orders (`ordervault` = `recordvault`) and ignoring unmatched ones, how many risk/margin-matched orders does each order type (`ordertune`) have? For the types above the average count across types, give one row with the average of their counts and the average amount by which they exceed that average, without rounding.	**A:** Count the joined risk/margin records for each order type. For the types above the average count, return the average of their counts and their average excess over that average. **B:** Count the distinct orders with risk/margin records for each order type. For the types above the average count, return the average of their counts and their average excess over that average.	A: 1 row, B: 1 row; first row differs; largest value change: 6.5 (A) vs 6.75 (B), 3.7%	Matched orders favors distinct orders, while the record-join framing supports counting matches. Specify whether an order with multiple risk records counts once. This is a counting-unit choice, not ties.	`SBNB`
`9e932237ed crypto`	Take risk/margin records matched to their orders (`riskandmargin.ordervault` = `orders.recordvault`), skipping any with no matching order, and get each time-in-force (`orders.timespan`) bucket's count. Orders with no time in force are their own bucket. Work out each bucket's share of the total of all bucket counts, and keep the buckets whose share is above the average share. Give me one summary row with three values: the total count across those buckets, their average share, and how far above the average share they are on average. Leave everything unrounded.	**A:** A time-in-force bucket's count is the number of joined risk/margin records, so orders with several risk/margin records count once per record. **B:** A time-in-force bucket's count is the number of distinct orders that have at least one risk/margin record.	A: 1 row, B: 1 row; first row differs; largest value change: 0.014 (A) vs 0.01426 (B), 1.9%	Bucket count does not explicitly distinguish matched risk records from distinct orders. The opening record-level framing favors A. This is a counting-unit choice, not ties.	`SS--`
`d52cbe42aa crypto`	For each order reference (`orderexecutions.ordersmark`, with executions lacking a reference counted together as one group), count its execution records and compute its share of all execution records. What are the top 10 order references by how far their share is above the average share across order references? Show the order reference, execution count, share and share minus the average share, without rounding, from largest difference to smallest (missing values last).	**A:** Return exactly 10 rows; any rows tied with the 10th are dropped arbitrarily. **B:** Return every row whose rank is within the top 10, including all rows tied at the 10th place.	A: 10 rows, B: 999 rows; 989 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	`-S--`
`f7b50ba44b crypto`	Link each order execution to its order (where `recordvault` matches `ordersmark`), and ignore executions without a matching order. Find how many executed orders each order type (`ordertune`) has, with a missing order type treated as its own group. Then work out each type's share of the total as a fraction. Looking only at the order types whose share is strictly above the average share across all order types, return a single unrounded row with: the sum of their counts, their average share, and the average amount by which their share exceeds that overall average share.	**A:** Count execution records per order type, so each fill counts separately, and then compute shares and the above-average summary from those counts. **B:** Count distinct executed orders per order type, so each order counts once however many fills it has, and then compute shares and the above-average summary from those counts.	A: 1 row, B: 1 row; first row differs; largest value change: 0.0065 (A) vs 0.006757 (B), 3.8%	Executed orders favors distinct orders, while starting from execution records supports a fill-level tally. Clarify whether partial fills count separately. This is not about ties.	`NB--`
`ff5d7b5c34 crypto`	Using fees joined to their orders (`fees.orderslink` = `orders.recordvault`) and those orders' users (`orders.userlink` = `users.userstamp`), drop any fee without a match. Compute the average fee rate (`feerate`, ignoring missing values) for each pair of orderslink and account scope (`users.acctscope`). Keep only the pairs whose average is strictly above the mean of all the pair averages, and work out how far each one exceeds that mean. Within each account scope, rank the pairs by that excess from largest to smallest and give me the top 3 pairs in each scope. Show orderslink, account scope, average fee rate, excess, and the pair's rank within its scope, all unrounded.	**A:** Within each account scope, tied pairs share a rank and the next rank skips. All pairs ranked 3 or better are returned, including every tied pair. **B:** Within each account scope, the pairs are numbered consecutively with ties broken by orderslink. Exactly the first three are returned.	A: 28 rows, B: 12 rows; 16 row keys in only one answer; largest value change: 1 (A) vs 3 (B), 67%	Top 3 and rank do not define tie handling. Shared ranks with gaps and consecutive positions are both plausible; B adds an unstated but reasonable deterministic tiebreaker.	`SBBB`

Set	Open request	Readings	Results	Review	Models
`10d659166c customers_and_ invoices`	Going from invoices to their orders and on to the customers (skip anything that doesn't link up), how many invoiced orders does each customer gender (`Customers.gender`) have? Turn that into each gender's share of the overall total, and list the genders whose share is above the average share across genders: give the gender, the count, the share, and how far the share sits above the average. Don't round anything.	**A:** Each gender's volume is its number of invoices. Every invoice counts, including several invoices on the same order. **B:** Each gender's volume is its number of distinct orders that have at least one invoice.	A: 1 row, B: 1 row; first row differs; largest value change: 10 (A) vs 6 (B), 40%	Invoiced orders favors distinct orders, while starting from invoices can suggest invoice-level volume. Clarify whether multiple invoices for one order count repeatedly. This is a counting-unit choice, not ties.	`BBBN`
`8e740d9161 customers_and_ invoices`	Using invoices linked to their orders and the customers who placed them, how many invoiced orders does each customer gender have (a missing gender counts as its own group)? Take the plain average of these per-gender figures. For the genders above that average, give one row with their average figure and the average amount by which they exceed the per-gender average.	**A:** Per customer gender, count the invoices (invoice–order–customer joined rows). Average these counts across genders. For the genders above that average, return the mean count and the mean excess over the average. **B:** Per customer gender, count the distinct orders that have at least one invoice. Average these counts across genders. For the genders above that average, return the mean count and the mean excess over the average.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Invoiced orders favors distinct orders, while starting from invoices can suggest invoice-level volume. Clarify whether multiple invoices for one order count repeatedly. This is a counting-unit choice, not ties.	`SBSS`
`db855429b2 customers_and_ invoices`	Which are the top 10 customers by share of all orders? For each, show the customer id, how many orders they placed, and their share (their order count divided by the total number of orders), sorted from highest to lowest share with missing values last. Don't round anything.	**A:** Cut the list at exactly 10 rows. Which of the tied customers make the cut is arbitrary. **B:** Keep everyone whose share reaches the 10th-highest value, so customers tied at the boundary all appear.	A: 10 rows, B: 11 rows; 1 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	`-SS-`
`31caaab8c7 financial`	For each branch district (`account.district_id`), count its accounts and work out its share of all accounts (its count divided by the total number of accounts). Show the district id, the account count, the share, and how far the share is above the average share across districts (share minus that average). Which are the top 10 districts by that difference, from highest to lowest? Don't round anything.	**A:** Cut the sorted list at exactly 10 rows and drop rows tied with the 10th. **B:** Cut the sorted list at the 10th row's value and keep every row tied with it.	A: 10 rows, B: 12 rows; 2 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	`SNSS`
`53afca1d49 financial`	What are the top 10 branch districts (`account.district_id`) by share of all accounts? For each, show the district id, its number of accounts and its share of all accounts, without rounding, listed from the highest share down.	**A:** Rank districts by their share of accounts and return exactly the first 10 rows, dropping tied districts arbitrarily at the cut. **B:** Rank districts by their share of accounts and return every district whose share reaches the 10th-highest share, including all ties at the cut.	A: 10 rows, B: 12 rows; 2 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	`SNSS`
`6c4f216d01 financial`	Using each order's account and that account's branch district: which accounts rank in the top 10 for how far their number of orders is above the average number of orders per account (averaged over accounts that have orders)? For each one, show the account id, the region of its branch, its order count and that difference (order count minus the average), largest difference first, unrounded.	**A:** Return exactly 10 accounts. Any tie at 10th place is broken arbitrarily. **B:** Return every account whose difference is at least the 10th-place value, so all accounts tied at 10th place are included.	A: 10 rows, B: 62 rows; 52 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	`--SS`
`9d0419e2cc financial`	Count the accounts in each branch district (with its region), counting only accounts whose district is on file, and find the districts whose count is above their region's average district count. For those districts, show the district id, region, account count, how far the count is above the region average, and the count divided by the average count across all these above-average districts. Give me the top 10 by that ratio, highest first, unrounded.	**A:** Return exactly ten rows, breaking any tie at the tenth position arbitrarily. **B:** Return the top ten plus every further row whose ratio equals the tenth row's ratio.	A: 10 rows, B: 11 rows; 1 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	`SNSS`
`841fa977bf formula_1`	Which circuits have hosted an above-average share of races? For each circuit (`races.circuitId`), give its number of races, its share of all races (its count divided by the total across all circuits), and how far that share is above the average circuit's share. Show the top 10 by that margin, largest first.	**A:** Rank the above-average circuits by how far their race share exceeds the average share and return exactly 10 rows, cutting through any tie at 10th place arbitrarily. **B:** Rank the above-average circuits by how far their race share exceeds the average share and return all circuits ranked 10th or better, including every circuit tied with the 10th one.	A: 10 rows, B: 11 rows; 1 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	`SSSS`
`c1a8ce4d18 formula_1`	Which circuit cities have an above-average share of qualifying? For each, show the city, its qualifying count, that count as a share of the total across all cities, and how far its share is above the average city share (all unrounded).	**A:** Each city's count is its number of driver qualifying entries, and its share is its share of all qualifying entries. **B:** Each city's count is its number of distinct races with qualifying data (qualifying sessions), and its share is its share of all such sessions.	A: 16 rows, B: 17 rows; first row differs; 1 row keys in only one answer; largest value change: 222 (A) vs 10 (B), 95%	Qualifying count can mean driver entries or qualifying sessions represented by distinct races. The request specifies neither unit. This is not about ties.	`-S--`
`dabca577ba formula_1`	Which are the top 10 races by how far their number of constructor result entries exceeds the average number of entries per race? Look only at races above that average. For each one, give the race id, the circuit's country, the total constructor points scored in that race (ignoring missing points), the number of entries, and the difference from the average. List them from the biggest difference down.	**A:** Take the above-average races ranked by how far their entry count exceeds the average, and return exactly the first 10 rows. When races tie at 10th place, the cut falls inside the tied group. **B:** Take the above-average races ranked by how far their entry count exceeds the average, and return every race whose entry count reaches the 10th-highest entry count. All races tied at the cutoff are included.	A: 10 rows, B: 16 rows; 6 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	`-S--`

Set	Open request	Readings	Results	Review	Models
c078054b47 insider	For each transaction in sentimentandfundamentals (transref), take the average analyst coverage count (anlycount, ignoring missing values). Keep only the transactions whose average is above the overall average of those per-transaction averages. For each one, show transref, its average, how far it is above that overall average, and how far it is above the average among just the kept transactions. Give me the top 10 by that last difference, highest first (missing last). Don't round.	**A:** Return exactly ten rows. Ties at the cut are broken arbitrarily. **B:** Return the top ten plus every row tied with the tenth.	A: 10 rows, B: 31 rows; 21 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	SSS-
156f4616e3 sakila_1	For each city in the address table, show its city ID, how many addresses it has, that count relative to the average number of addresses per city, and its share of all addresses. Which are the top 10 cities by that share, highest first? Don't round anything.	**A:** Return exactly 10 cities. If cities are tied at the 10th place, keep an arbitrary subset of them. **B:** Return every city ranked in the top 10, including all cities tied with the 10th, so the result may have more than 10 rows.	A: 10 rows, B: 599 rows; 589 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	-N--
46093f9250 sakila_1	Looking at what has actually been rented (each rental traced through its inventory copy to the film, skipping any that don't link up), what share of the rented films does each film rating account for, out of the total across all ratings? For the ratings whose share is above the average share across ratings, give one row with: their combined number of rented films, their average share, and how far on average their share sits above that average share. No rounding.	**A:** Each rating's share is its number of rental transactions as a fraction of all rental transactions. **B:** Each rating's share is its number of distinct rented films as a fraction of the total distinct rented films summed over ratings.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Rented films can refer to rental transactions or distinct rented titles. Both fit the wording; the counting unit is unspecified. This is not a tie-handling choice.	NB--
5fd7aa9aff sakila_1	Looking at the films we have in stock by rating, which ratings have more films in inventory than the average rating does? Give me a single row with the total films in stock across those above-average ratings and, on average, how far above the across-rating average they are (no rounding).	**A:** For each rating, count the physical inventory copies; keep the ratings above the average count, then return the sum of their counts and their average amount above the average. **B:** For each rating, count the distinct film titles that have at least one copy in inventory; keep the ratings above the average count, then return the sum of their counts and their average amount above the average.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Films in stock can mean physical copies or distinct available titles; the counting unit is unspecified. This is a counting-unit choice, not ties.	BBSN
ab1cedcf3f sakila_1	Which are the top 10 cities by share of addresses? For each city_id, show how many addresses it has and what fraction of all addresses that is (its count divided by the total number of addresses, as an unrounded fraction rather than a percentage), sorted from the highest fraction to the lowest.	**A:** Per-city address counts and their share of all addresses, sorted by share in descending order. Return exactly 10 rows and break ties at the 10th place arbitrarily. **B:** Per-city address counts and their share of all addresses, sorted by share in descending order. Return the top 10 plus every city whose count ties with the 10th-ranked city's count.	A: 10 rows, B: 599 rows; 589 row keys in only one answer; no shared value changes	Top 10 can mean ten rows or all cities tied at the tenth row; no tie policy is given.	-N--
b77a784ee8 sakila_1	What are the top 10 films by number of copies in our inventory? For each, show the film_id, how many copies we have, and what share of all inventory copies that is (copies divided by the total number of copies, unrounded), with the highest share first.	**A:** List the top 10 films by share of inventory copies, returning exactly 10 rows and choosing arbitrarily among films tied at the cut. **B:** List the top films by share of inventory copies, keeping every film whose copy count reaches the 10th-highest count, so ties at the cut are all included.	A: 10 rows, B: 72 rows; 62 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	-N--
5c048dd871 student_club	Which are the top 10 members by share of total event attendance? Match attendance to members by member id, count attendance per member id and position, and show the member id, position, attendance count and share of all attendance, from highest share to lowest.	**A:** Rank member/position groups by their share of all attendance and return exactly 10 rows, cutting arbitrarily through any tie at 10th place. **B:** Rank member/position groups by their share of all attendance and return the top 10 plus every row tied with the 10th row's share.	A: 10 rows, B: 12 rows; 2 row keys in only one answer; no shared value changes	Top 10 leaves the cutoff-tie policy unspecified: ten rows and all rows tied with the tenth row are both reasonable. Use the tenth ordered row, not the tenth distinct value, for B.	SNSS
6b346d7be2 student_club	Taking attendance entries that match an event (others left out), how many attendees does each event status have, and what share of the total across all statuses is that? Show the status, its number of attendees, and the share, without rounding.	**A:** A status's attendance is the total number of member-event attendance entries for events with that status. **B:** A status's attendance is the number of different members who attended any event with that status.	A: 3 rows, B: 3 rows; first row differs; largest value change: 297 (A) vs 30 (B), 90%	Attendees can mean attendance occurrences or unique participating members. Neither repeated attendance nor deduplication is specified. This is a counting-unit choice, not ties.	SBSS
d3a162e8c3 student_club	What are the top 10 events by share of total attendance? For each one, give the event id, its number of attendance records, and its share of all attendance records (its count divided by the total, unrounded), from the largest share down, and return nothing if there's no total to divide by.	**A:** List the 10 events with the largest share of attendance records and stop at exactly 10 rows, breaking ties at the cut arbitrarily. **B:** List the events with the largest share of attendance records down to the 10th place, including every event tied with the 10th.	A: 10 rows, B: 11 rows; 1 row keys in only one answer; no shared value changes	Top 10 does not state whether cutoff ties should be included. A is the usual fixed-length interpretation.	SSNS
de3cf55570 student_club	Using event attendance and the city of each attending member's zip code (member.zip = zip_code.zip_code), find the number of event attendees from each city and each city's share of the total across all cities. For the cities whose share is above the average city share, give in one row the average number of attendees, the average share, and the average amount by which the share is above that average.	**A:** The per-city count is the number of attendance records (one per member per event). **B:** The per-city count is the number of distinct members who attended at least one event.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	Attendees can mean attendance occurrences or unique participating members. Neither repeated attendance nor deduplication is specified. This is a counting-unit choice, not ties.	NBBN
b6c13c0329 thrombosis_prediction	Looking at the examination data by patient sex (blank and missing sexes as their own groups), how many were examined in each sex, and what share of the overall total is that? For the sexes above the per-sex average, give one row with their combined total, their average share, and on average how far above the per-sex average they are.	**A:** For each sex, count the examination records. **B:** For each sex, count the distinct patients who have at least one examination.	A: 1 row, B: 1 row; first row differs; 2 row keys in only one answer; no shared value changes	How many were examined can mean examinations or distinct patients examined. B is more natural for people; the data framing also supports A. This is not about ties.	SBNN